50.039 Theory and Practice of Deep Learning
# W1-S1 Introduction and Machine Learning Reminders

Matthieu De Mari

A quick word about instructors

Matthieu (**Matt**) De Mari

- Lecturer at SUTD (Python, Deep Learning, Prog. Language Concepts, Capstone, etc.)

- Information Systems Technology and Design (ISTD) and Computer Science and Design (CSD) pillar/faculty

- PhD from CentraleSupelec (France)

- Email: matthieu_demari@sutd.edu.sg

- Office @ SUTD (in SUTD Academy/AI cluster): 2.401.07.

Objectives of 50.039 DL

- Introduce the technical aspects related to the **implementation of Deep Neural Networks from scratch**.

- Teach the students about the **most popular framework** for implementing Deep Neural Networks (as of Jan 2023): **PyTorch.**

- Discuss the **mathematical foundations and intuitions** behind the Deep Learning framework, layers and some techniques.

- Give the students a **global overview of the advanced techniques** related to Deep Learning and Deep Neural Networks.

- Describe **examples of practical applications** of Deep Learning, showing how some key AIs were implemented using Deep Learning.

Objectives of 50.039 DL

- Eventually, **assemble all concepts and tricks of the trade** discussed in this course to produce big architecture models that are close to the current state-of-the-art ones (e.g. ChatGPT, Dall-E, etc.).

- Discuss **ethical aspects** related to some of the concepts discussed.

- Give students **an understanding of the current state of research in the Deep Learning community**, and more specifically <u>**pointers**</u> as to where the most up-to-date information can be found.

- This includes **tracking big names** in the field, identifying **key scientific papers** that have changed the field, but also **newsletters**, **websites** or **magazines** that disclose information in Layman terms.

Skills needed for this course

- **Must-have CS:** Python, Numpy, Matplotlib.

- **Must-have Math:** Linear Algebra, multiple variables calculus, derivatives, optimization probability and statistics.

- **Good-to-have CS:** Machine Learning, Data Processing, Scipy, Sklearn.

- **Good-to-have Math:** Graph Theory, Game theory.

Need to revise? Time to check your 10.014 CTD, 10.013 M&A, 10.018 MS&S, 10.022 MU, 10.020 DDW, and 50.007 ML!

The way I teach things

An expert is not someone who

- Undertook a 6h-long online course,

- Implements Neural Networks by randomly connecting basic layers (typically Linear and Conv),

- Using a random framework he/she does not understand (or autoML!),

- and is happy to get an 80% accuracy on a simple dataset like MNIST or CIFAR-10.

An expert is someone who

- Understands how Neural Networks operate and the mathematical intuition behind typical (advanced) operations used in layers,

- Understands how the frameworks have been built and what they do behind the scenes,

- Knows AI is a fast-evolving field and knows how to stay up to date when it comes to AI.

The way I teach things

- **One topic per week** (or so) and hopefully, we will cover all the important concepts and important directions of DL these days.

- **Lectures slides**, along with **supporting notebooks**.

- **Mathematical aspects** discussed in class (to show what is happening behind the scenes of Neural Networks/Deep Learning).

- **Made a choice:** providing lots of code to demonstrate concepts, but **need you to play with them autonomously** to explore concepts!

- **Extra reading**, for curiosity (supporting papers, articles, etc.) and suggestions of pointers to follow to stay up to date.

- Suggestions for **continuing your learning**, after this course.

The way I teach things

I will upload materials on both **eDimension** and **Google Slides** (everything referenced on eDimension).

- **eDimension slides:** Read-only, reference slides from Prof, contains ground truth about things (lol).

- **Google slides:** shared and can be modified/commented by everyone who has link to slides.

- **Feeling like there is a notion rather obscure?**
  Comment/ask your questions on the Google Slides and other students/Prof can answer you!

The way I teach things

A slide with random and obscure content

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.

Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat.

Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur.

Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Matthieu De Mari
12:01 Today

What the hell is this? A latin prayer to summon a demon or something?

Matthieu De Mari
12:02 Today

It is actually a lorem ipsum, more about this here:
https://en.wikipedia.org /wiki/Lorem_ipsum

Reply or add others with @

to add speaker notes

The way I teach things

I will upload materials on both **eDimension** and **Google Slides** (everything referenced on eDimension).

- **eDimension slides:** Read-only, reference slides from Prof, contains materials to be presented on physical lectures.

- **Google slides:** Shared on eDimension and can be modified/commented by everyone who has link to the slides.

- **Feeling like there is a notion rather obscure?** Comment/ask your questions on the Google Slides and other students/Prof can answer you!

- (Will also greatly help me identify parts/materials that need clarifications/improvements for future generations!)

Syllabus

- **Week 1:** Introduction to course, some ML jargon reminders, linear and polynomial regression, generalization, ridge regression and regularization, overfitting/underfitting, logistic regression, neurons objects and how they relate to biology, our first shallow neural network, gradient descent reminders, training and testing procedure.

- **Week 2:** Introduction to PyTorch framework, tensors and dataloaders, implementing a shallow neural network in PyTorch, backpropagation in PyTorch with AutoGrad, advanced optimizers, multi-label classification with shallow neural networks, moving from shallow to deep neural networks.

Syllabus

- **Week 3:** Guided project and good practices for Deep Learning projects (train/test/dev, bias/variance, advanced regularization, dropout, normalizing inputs/outputs/layers, trainer functions, savers/loader functions for reproducibility and transfer learning).

- **Week 4:** The image data type, image processing techniques and typical computer vision operations, the convolution operation and layers, Convolutional Neural Networks, advanced CNNs and SotA. Preparing transition to the 50.035 Computer Vision course.

Syllabus

- **Week 5:** Adversarial machine learning, attacking a Neural Network with basic gradient-based attacks, fundamental limits of Neural Networks, defence mechanisms and state-of-the-art of some advanced attacks techniques.

- **Week 6:** Sequential data (times series, text, etc.), vanilla Recurrent Neural Networks, Gated Recurrent Units, Long-Short Term Memory cells, advanced RNN networks, mixing models for advanced architectures. MidTerm exam (based on W1-5).

Syllabus

- **Week 8:** The embedding problem, more advanced concepts on RNNs, introduction to Natural Language Processing (NLP) and Word Embeddings for NLP, brief state-of-the-art on NLP, attention and transformers architectures. Preparing transition to the 50.040 Natural Language Processing course.

- **Week 9:** Quick introduction to Graph Theory and typical graph datasets and problems, basics of Graph Convolutional Networks, brief state-of-the-art of advanced Graph Convolutional Networks.

Syllabus

- **Week 10:** Generative Models, Autoencoders and Variational Autoencoders, Generative Adversarial Networks (GANs), Advanced concepts on Generative Adversarial Networks, why is it so difficult to train GANs in practice. Basic introduction to diffusion models (more details to be covered in the 50.035 Computer Vision course). If time allows, practice on GANs.

- **Week 11:** Brief introduction to reinforcement learning, and state-action-rewards systems, multi-armed bandit problem and the exploration/exploitation trade-off, Q-learning and Deep Q-Learning. Brief state-of-the-art discussion about further works in Reinforcement Learning.

Syllabus

- **Week 12:** Explainability/Interpretability and open questions in research about Neural Networks.
  What will be the next revolution in AI? (a quick word on ChatGPT and other LLM models).
  Closing and future directions for studying Deep Learning.

- **Week 13:** Recap. Project presentations and guest conferences (TBA).

- **Week 14:** Final exam (probably W1-13).

Supporting Textbooks

Supporting textbooks, for your curiosity (not needed to understand this course).

- Michael A. Nielsen, "Neural networks and deep learning", 2015. (http://neuralnetworksanddeeplearning.com/)

- Ian Goodfellow, Yoshua Bengio and Aaron Courville, "Deep learning", 2016. (https://www.deeplearningbook.org/)

- Stevens et al., "Deep Learning with PyTorch", 2020. (https://www.manning.com/books/deep-learning-with-pytorch)

Neural Networks and Deep Learning

Neural Networks and Deep Learning is a free online book. The book will teach you about:

- Neural networks, a beautiful biologically-inspired programming paradigm which enables a computer to learn from observational data
- Deep learning, a powerful set of techniques for learning in neural networks

Neural networks and deep learning currently provide the best solutions to many problems in image recognition, speech recognition, and natural language processing. This book will teach you many of the core concepts behind neural networks and deep learning.

For more details about the approach taken in the book, see here. Or you can jump directly to Chapter 1 and get started.

Evaluation and grading

- **Homeworks (19%):** Given on Weeks 2, 5, 8, 10.
  Usually come in the form of a Jupyter Notebook, containing explanations, code snippets and questions.
  Submissions on eDimension, two weeks later or so, as a small PDF report containing code, figures and answers to questions.
  When time allows, debrief of homeworks in class.

- **MidTerm Exam (20%):** Given on Week 6, March 1st 2023.
  Theoretical, paper exam, notions of Week 1-5 to be tested.
  More details about venue and exam details to be announced closer to the exam date.

Evaluation and grading

- **Final Exam (30%):** Given on Week 14, April 26th 2023. Theoretical, paper exam, notions of Week 1-13 to be tested. More details about venue and exam details to be announced.

- **Project (29%):** Groups of 2-3 students. Submission for project (code, report, presentation) expected on Week 13. Problem statement to be freely decided by students, as long as it matches a list of given requirements. We will have a guided demo project on Week 3, more details about the project will be given on Week 5.

- **Student Feedback Survey (2%):** the usual.

Technical pre-requisites

- Install **Python 3**, if you have not done so already (this course does not cover C++/Java).

- **Libraries needed (maybe more based on projects/homeworks):** numpy, matplotlib, scipy, sklearn, networkx, pillow, gym.

- **Jupyter notebooks** for demos of code, along with the slides.

- In doubt, you can always use **Google Colab** to run the codes.

Technical pre-requisites

- Framework of choice will be **PyTorch!** (not Tensorflow, not Keras, not MXNet, etc.)

- Increasing popularity and preferred to Google's Tensorflow these days for many reasons.

- Learn more, if curious: https://www.assemblyai.com/blog/pytorch-vs-tensorflow-in-2022/

Fraction of Papers Using PyTorch vs. TensorFlow

Installing PyTorch and CUDA

- Install PyTorch, by getting the right version from https://pytorch.org/get-started/locally/

PyTorch Build	Stable (2.0.1)		Preview (Nightly)	
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python		C++ / Java	
Compute Platform	CUDA 11.7	CUDA 11.8	ROCm 5.4.2	CPU
Run this Command:	pip3 install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cu117			

Installing PyTorch and CUDA

- Check if your GPU is in the list of acceptable GPUs.
  https://developer.nvidia.com/cuda-gpus

- If so, install CUDA (check that you CUDA version number matchesthe one of your PyTorch install!)
  https://developer.nvidia.com/cuda-downloads

Installing PyTorch and CUDA

Question: How to use CUDA with MacOS?

Efforts are being made to make GPU computing compatible with AMD and MacOS GPUs, but still far from it (afaik).

Select Target Platform

Click on the green buttons that describe your target platform. Only supported platforms will be shown. By dowr terms and conditions of the CUDA EULA.

Operating System	Linux	Windows			
Architecture	x86_64				
Version	10	11	Server 2016	Server 2019	Server 2022
Installer Type	exe (local)	exe (network)			

Installing PyTorch and CUDA

- Your GPU not in the list of CUDA-enabled GPUs?

- You bought a MacOS laptop?

- **Most notebooks can still run on CPU (but they might take significantly longer to run).**

- Always the option of using **Google Colab**, or create an education account on AWS.

CPU vs GPU Training Time Comparison in Seconds

Runtime

Colab (GPU) ▮ 717

Colab (CPU) ▮▮▮▮▮▮▮▮▮▮▮▮ 22.120

0K 2K 4K 6K 8K 10K 12K 14K 16K 18K 20K 22K 24K

Training Time

Google Colaboratory

aws educate

Checking your PyTorch and CUDA install

- **Hello World for PyTorch:** to check you have PyTorch installed correctly. The code below should run and display a tensor.

```python
import torch
x = torch.rand(5, 3)
print(x)

tensor([[0.3380, 0.3845, 0.3217],
        [0.8337, 0.9050, 0.2650],
        [0.2979, 0.7141, 0.9069],
        [0.1449, 0.1132, 0.1375],
        [0.4675, 0.3947, 0.1426]])
```

- **Hello World for CUDA:** to check you have correctly installed CUDA on top of PyTorch. The code below should print *True,* as the output of *torch.cuda.is_available()*

```python
import torch
torch.cuda.is_available()
```

A quick word about PyTorch 2.0

PyTorch v2.0 is the current (latest version released in March 2023).

- Should be fully retro-compatible.

- Revising and testing notebooks (they were designed last year and tested on v1.13).

- Kindly let me know if you have any issues!

- Will let you know about new v2.0 features on W13.

About this week (Week 1)

1. What are the **typical concepts of Machine Learning** to be used as a starting point for this course?

2. What are the **different families of problems** in Deep Learning?

3. What is the **typical structure of a Deep Learning problem**?

4. What is **linear regression** and how to implement it?

5. What is the **gradient descent algorithm** and how is it used to **train Machine Learning models**?

6. What is **polynomial regression** and how to implement it?

7. What is **regularization** and how to implement it in **Ridge regression**?

About this week (Week 1)

8. What is **overfitting** and why is it bad?

9. What is **underfitting** and why is it bad?

10. What is **generalization** and how to evaluate it?

11. What is a **train-test split** and why is it related to **generalization**?

12. What is a **sigmoid** function? What is a **logistic** function?

13. How to perform **binary classification** using a **logistic regressor** and how is it related to linear regression?

About this week (Week 1)

14. What are **Neural Networks** and how do they relate to the **biology of a human brain**?

15. What is a **Neuron** in a Neural Network and how does it relate to linear/logistic regression?

16. What is the **difference** between a **shallow** and a **deep neural network**?

17. How to **implement a shallow Neural Network** manually and define a **forward propagation** method for it?

18. How to **train a shallow Neural Network** using **backpropagation**? How to define **backward propagation** and **trainer** functions?

What is AI/ML/NN/DL?

Definition (Artificial Intelligence):

In Computer Science, **Artificial Intelligence (AI)** refers to the theory and development of computer systems capable to **replicate/emulate the human brain**, more specifically **perform cognitive tasks** that normally require human intelligence, such as visual perception, speech recognition, decision-making, and translation between languages.

Artificial intelligence

Machine Learning

Neural Networks

Deep Learning

Artificial Intelligence

Machine Learning

Neural Networks

Deep Learning

What is AI/ML/NN/DL?

Definition (Machine Learning):

In Computer Science, **Machine Learning (ML)** refers to the field of study that describes **techniques** and **algorithms** that give computers the **ability to learn without being explicitly programmed**.

Some implementations of machine learning may rely on **data** and **neural networks**.

Conventional programming

Input (Given) → Function (Created) → Output (Obtained)

vs.

Machine Learning

Input (Available) → Function (?) → Output (Available)

What is AI/ML/NN/DL?

What we have done in programming so far was to design functions,

- which would do **specific operations,**
- and return **outputs,**
- for any **input** we could give it.

But sometimes, we can encounter problems where

- we can easily find **inputs** and **expected** outputs,
- but the **function** to be coded is **not simple** to figure out.
- **E.g., what animal in picture?**

| Input (Given) | → | Function (Created) | → | Output (Obtained) |

| Input (Available) | → | Function (?) | → | Output (Available) |

What is AI/ML/NN/DL?

E.g., what animal is in the picture?

Typical problem in **Computer Vision**, called **Image Recognition**.

Input x (available)	→	Function f (?)	→	Output y = f(x) (available)

It's a cat!

What is AI/ML/NN/DL?

E.g., what animal is in the picture?

Typical problem in **Computer Vision**, called **Image Recognition**.

Input x (available)	➡	Function f (?)	➡	Output y = f(x) (available)

Very easy for a human…
But, how would we explain the logic of it to a computer?

It's a cat!

What is AI/ML/NN/DL?

Other scenarios of difficult functions have to do with tasks where there is **no easy closed-form expression connecting inputs to outputs.**

- **E.g., what is a good selling price for my appartment?**

- Guessing the selling price of an appartment based on its parameters (size, location, etc.) and previous sales.

s$ 1,680,000   Negotiable

3 🛏   3 🛁   1184 sqft   S$ 1,418.92 psf

Details

Property Type	Floor Size
Executive Condominium For Sale	1184 sqft
Developer	PSF
Tampines EC Pte Ltd	S$ 1,418.92 psf
Furnishing	Floor Level
Partially Furnished	High Floor
Tenure	TOP
99-year Leasehold	January, 2016

What is AI/ML/NN/DL?

Definition (Neural Networks):

Neural Networks (**NNs**) are computing systems inspired by the biological neural networks that constitute animal brains.

NNs are based on a **collection of connected units (or nodes) called artificial neurons**, which are a (loose) model for the neurons in a biological brain.

What is AI/ML/NN/DL?

Definition (Neural Networks):

Neural Networks (**NNs**) are computing systems inspired by the biological neural networks that constitute animal brains.

NNs are based on a **collection of connected units (or nodes) called artificial neurons**, which are a (loose) model for the neurons in a biological brain.

Each connection, like the synapses in a biological brain, can transmit a signal to other neurons, therefore **processing any information** given as inputs to produce a final signal as output.

What is AI/ML/NN/DL?

Definition (Deep Learning):

**Deep Learning (DL)** is a subfield of machine learning, and <u>deep</u> **neural networks** make up the backbone of deep learning algorithms.

The **number of node layers,** or **depth**, of neural networks is what distinguishes a **shallow** **neural network** from a **deep** **neural network**, which must have more than three.

What is AI/ML/NN/DL?

# Supervised vs Unsupervised Learning

Definition (Supervised vs. Unsupervised Learning):

Supervised and **Unsupervised Learning** are the two techniques of ML.

The main difference is **the need for labelled training data**:

- Supervised machine learning relies on labelled input and output data to learn and make predictions,

- while **unsupervised learning does not require labelled data.**

Supervised Learning Examples

- Examples of **supervised learning:** spam detection, text classification, predicting the stock market, etc.

Figure 7. MAST stock price predictions using LSTM trained on AAPL

Regression vs. Classification

Definition (Regression):

Regression models are used to identify the relationships between the input and output variables.

Regression algorithms are used to **predict continuous values** such as price, salary, age, etc.

In regression, the outputs are often **continuous numerical values**.

S$ 1,680,000 Negotiable

3 🛏 3 🛁 1184 sqft S$ 1,418.92 psf

Details

Property Type	Floor Size
Executive Condominium For Sale	1184 sqft
Developer	PSF
Tampines EC Pte Ltd	S$ 1,418.92 psf
Furnishing	Floor Level
Partially Furnished	High Floor
Tenure	TOP
99-year Leasehold	January, 2016

Regression vs. Classification (Supervised)

Definition (Classification):

Classification models are used to divide the samples in the dataset into different **classes**.

Classification algorithms are then used to predict/classify **discrete values**, such as Male or Female, True or False, Spam or Not Spam, etc.

In Classification, the outputs are **discrete** or **categorical values**.

Supervised Learning Examples

Example: predicting the market.

Question: is it a **regression** or a **classification** task?

Figure 7. MAST stock price predictions using LSTM trained on AAPL

Supervised Learning Examples

Example: predicting the market.

Question: is it a **regression** or a **classification** task?

Depends.

If the output we are **predicting is the value of the stock in the future**, then probably **regression**.

Figure 7. MAST stock price predictions using LSTM trained on AAPL

Supervised Learning Examples

Example: predicting the market.

Question: is it a **regression** or a **classification** task?

Depends.

If the output we are **predicting is the value of the stock in the future**, then probably **regression**.

If the plan is to **predict whether we should buy, sell or hold**, then probably **classification**.

Figure 7. MAST stock price predictions using LSTM trained on AAPL

Supervised vs Unsupervised Learning

Definition (Supervised vs. Unsupervised Learning):

Supervised and **Unsupervised Learning** are the two techniques of ML.

The main difference is **the need for labelled training data**:

- Supervised machine learning relies on labelled input and output data to learn and make predictions,

- while **unsupervised learning does not require labelled data.**

Unsupervised Learning Examples

- Examples of **unsupervised learning:** clustering customers based on their behaviors, segmenting images, and identifying topics in a document.

TSNE Visualization of Book Embeddings

Unsupervised Learning Examples

- Examples of **unsupervised learning:** clustering customers based on their behaviors, segmenting images, and identifying topics in a document.

Identify topics that seem to appear in those papers and their prominence

TSNE Visualization of Book Embeddings

TYPES OF ML / NLP PAPERS

Clustering vs. Association (Unsupervised)

Definition (Clustering vs. Association):

Unsupervised learning can be used for two types of problems: **Clustering** and **Association**.

The main difference is that

- **clustering groups data based on similarity**,
- while **association looks for relationships between variables**.

Clustering vs. Association (Unsupervised)

Examples of clustering:

- Group customers into clusters based on the similarities in their purchasing behaviours, such as their spending habits or types of products they like.

- Group the catalogue of movies available on Netflix and identify what movies have in common, so you can later recommend a movie to a user based on what he/she has seen before.

TSNE Visualization of Book Embeddings

Clustering vs. Association (Unsupervised)

Example of association:

- Finding the relationship between customer buying behaviors and customer satisfaction.

- Association can be used to understand how changes in customer behavior impacts customer satisfaction, such as if buying a certain product increases customer satisfaction (or the opposite).

Deep Learning and Supervised/Unsupervised

While it is possible to use Deep Learning models and Neural Networks to perform both Supervised and Unsupervised tasks...

The vast majority of DL/NN tasks fall in the category of **supervised learning**.

(And we will therefore mostly focus on that)

Elements of a ML problem

Definition (Elements of a ML problem):

Machine learning problems should be well-defined, i.e. the following elements must be clearly identified:

1. **Task (T)** at hand,

2. **Inputs (I)** and **Outputs (O),**

3. **Model (M)** definition and its **trainable parameters (P),**

4. **Loss (L)** function to measure the performance of current model on said task and dataset.

s$1,680,000   Negotiable

3 🛏   3 🛁   1184 sqft   S$ 1,418.92 psf

Details

Property Type	Floor Size
Executive Condominium For Sale	1184 sqft

Developer	PSF
Tampines EC Pte Ltd	S$ 1,418.92 psf

Furnishing	Floor Level
Partially Furnished	High Floor

Tenure	TOP
99-year Leasehold	January, 2016

Elements of a ML problem

In the example on the right...

1. **Task (T)** at hand:

- Predict price of an apartment based on the apartment features.

- Using a dataset consisting of previously seen appartments to learn the selling prices of apartments.

- It is therefore a supervised regression task.

S$1,680,000 Negotiable

3 🛏   3 🛁   1184 sqft   S$ 1,418.92 psf

Details

Property Type		Floor Size	
Executive Condominium For Sale		1184 sqft	

Developer		PSF	
Tampines EC Pte Ltd		S$ 1,418.92 psf	

Furnishing		Floor Level	
Partially Furnished		High Floor	

Tenure		TOP	
99-year Leasehold		January, 2016	

Elements of a ML problem

In the example on the right…

2. Inputs (I) and **Outputs (O):**

- Inputs will be a list of apartment features (e.g. number of bathrooms, number of bedrooms, surface size, etc.).

- Output will be the selling price in SGD or surface price in SGD per sqft.

S$ 1,680,000 Negotiable

3 🛏 3 🛁 **1184** sqft S$ **1,418.92** psf

Details

Property Type	Floor Size
Executive Condominium For Sale	1184 sqft
Developer	PSF
Tampines EC Pte Ltd	S$ 1,418.92 psf
Furnishing	Floor Level
Partially Furnished	High Floor
Tenure	TOP
99-year Leasehold	January, 2016

Elements of a ML problem

In the example on the right...

3. **Model (M)** definition and its **trainable parameters (P):**

- Could be a linear or polynomial regression model (to be discussed this week).

- Or a neural network of some sort.

- Or even something else (check your ML course!)

s\$ **1,680,000**   Negotiable

3 🛏   3 🛁   **1184** sqft   S\$ **1,418.92** psf

Details

Property Type	Floor Size
Executive Condominium For Sale	1184 sqft
Developer	PSF
Tampines EC Pte Ltd	S\$ 1,418.92 psf
Furnishing	Floor Level
Partially Furnished	High Floor
Tenure	TOP
99-year Leasehold	January, 2016

Elements of a ML problem

In the example on the right...

4. **Loss (L)** function to measure the performance of current model on said task and dataset:

- To be discussed later, as this might depend on the model used.

s$1,680,000   Negotiable

3 🛏   3 🛁   1184 sqft   S$ 1,418.92 psf

Details

Property Type	Floor Size
Executive Condominium For Sale	1184 sqft
Developer	PSF
Tampines EC Pte Ltd	S$ 1,418.92 psf
Furnishing	Floor Level
Partially Furnished	High Floor
Tenure	TOP
99-year Leasehold	January, 2016

Our first toy example

Our first toy example: simplified version of the apartment price prediction.

1. **Task (T)** at hand:

- **Predict price of an apartment based on the apartment features.**

- We will generate a "fake" dataset consisting of apartments with features and selling prices.

- It is a **supervised regression task**.

S$ 1,680,000   Negotiable

3 🛏   3 🛁   1184 sqft   S$ 1,418.92 psf

Details

Property Type	Floor Size
Executive Condominium For Sale	1184 sqft
Developer	PSF
Tampines EC Pte Ltd	S$ 1,418.92 psf
Furnishing	Floor Level
Partially Furnished	High Floor
Tenure	TOP
99-year Leasehold	January, 2016

Our first toy example

Our first toy example: simplified version of the apartment price prediction.

2. **Inputs (I)** and **Outputs (O):**

- **Inputs** will be a **list of apartment features**, which here will only consist of the **surface of the apartment**, in square meters (sqm).

- **Output** will be the **selling price** in millions of SGD (mSGD).

s$1,680,000 Negotiable

3 🛏   3 🛁   1184 sqft   S$ 1,418.92 psf

Details

Property Type	Floor Size
Executive Condominium For Sale	1184 sqft
Developer	PSF
Tampines EC Pte Ltd	S$ 1,418.92 psf
Furnishing	Floor Level
Partially Furnished	High Floor
Tenure	TOP
99-year Leasehold	January, 2016

Our first toy example (mock dataset)

Have a look at Week 1, Notebook 1.

We will create a dataset, which uses surfaces as inputs, with values drawn randomly between 40 and 150 sqm.

We will assume that the price is randomly generated, using a uniform distribution, with:

- Average price being 100000 + 14373 times the surface in sqm,
- and a random "variation" of +/- 10% around the average value.

Our first toy example (mock dataset)

```python
# Two random generator functions to generate a mock dataset.
# 1. Surfaces randomly generated as uniform between min_surf and max_surf
def surface(min_surf, max_surf):
    return round(np.random.uniform(min_surf, max_surf), 2)
```

```python
# 2. Price is 100000 + 14373 times the surface in sqm, +/- 10%
# (randomly shifted to give the dataset some diversity).
def price(surface):
    # Note: this will return the price in millions of SGD.
    return round((100000 + 14373*surface)*(1 + np.random.uniform(-0.1, 0.1)))/1000000
```

```python
# 3. Generate dataset function
def generate_datasets(n_points, min_surf, max_surf):
    x = np.array([surface(min_surf, max_surf) for _ in range(n_points)])
    y = np.array([price(i) for i in x])
    return x, y
```

Our first toy example (mock dataset)

```python
1  # 4. Dataset generation (n_points points will be generated).
2  # Surfaces randomly generated as uniform between 40sqm and 150sqm.
3  # We will use a seed for reproducibility.
4  min_surf = 40
5  max_surf = 150
6  np.random.seed(27)
7  n_points = 50
8  inputs, outputs = generate_datasets(n_points, min_surf, max_surf)
9  print(inputs)
10 print(outputs)
```

```
[ 86.83 129.6   120.89 135.48   82.17 147.74 138.25   63.07 121.6   112.95
 137.55 134.38 122.42 135.72   60.54  75.81  81.02 127.31  56.62  58.69
  48.93  73.57 126.16  57.92  47.77 117.12  59.91 105.88  85.68  96.49
  64.27 119.81 133.44 142.18 120.95  92.42  94.22 105.4   48.36  52.92
 146.31 104.17  50.17  41.5  132.06 140.63 117.5   82.57 117.63 141.56]
[1.290893 2.034977 1.84501  2.201767 1.389632 2.218678 2.127228 0.959054
 2.029469 1.623609 2.111638 2.09194  2.012386 1.89553  1.004256 1.250228
 1.368325 1.830127 1.000719 0.906513 0.867629 1.065907 1.979544 1.001403
 0.796199 1.771816 0.867878 1.587176 1.25434  1.40047  1.002361 1.9972
 1.894479 2.293443 1.823577 1.340533 1.358613 1.602167 0.750759 0.863093
 2.30035  1.741468 0.7448   0.763732 2.093772 1.986868 1.90702  1.289541
 1.86578  2.231851]
```

```
1   # Display dataset and see that there is a rather clear linear trend.
2   plt.figure(figsize = (12, 8))
3   plt.scatter(inputs, outputs)
4   plt.xlabel("Surface (in sqm)")
5   plt.ylabel("Price (in millions SGD)")
6   plt.show()
```

Our first toy example

Our first toy example: simplified version of the apartment price prediction.

3. **Model (M)** definition and its **trainable parameters (P):**

- Here, our model will consist of a **linear regression model**.

s$ 1,680,000   Negotiable

3 🛏   3 🛁   1184 sqft   S$ 1,418.92 psf

Details

Property Type	Floor Size
Executive Condominium For Sale	1184 sqft
Developer	PSF
Tampines EC Pte Ltd	S$ 1,418.92 psf
Furnishing	Floor Level
Partially Furnished	High Floor
Tenure	TOP
99-year Leasehold	January, 2016

Our first model, the linear regression

Definition (Linear Regression):

Linear Regression is a model, which **assumes that there is a linear relationship between inputs and outputs.**

It therefore consists of two trainable parameters (a, b), to be freely chosen.

These will connect any input x_i to its respective output y_i, with the following equation:

$$y_i \approx a \, x_i + b$$

Our first model, the linear regression

Definition (Linear Regression):

Linear Regression is a model, which **assumes that there is a linear relationship between inputs and outputs.**

It therefore consists of **two trainable parameters** (a, b), to be freely chosen.

These will connect any **input** x_i to its respective **output** y_i, with the following equation:

$$y_i \approx a\,x_i + b$$

```python
1  # Linear regression has two trainable parameters (a and b).
2  # Other parameters, like min_surf, max_surf, n_points will
3  # help get points for the upcoming matplotlib displays.
4  def linreg_matplotlib(a, b, min_surf, max_surf, n_points = 50):
5      x = np.linspace(min_surf, max_surf, n_points)
6      y = a*x + b
7      return x, y
```

Let us try different values of (a, b)

Model 1

Model 2

Model 3

```python
1   # Display dataset
2   plt.figure(figsize = (12, 8))
3   plt.scatter(inputs, outputs)
4   plt.xlabel("Surface (in sqm)")
5   plt.ylabel("Price (in millions SGD)")
6
7   # Add some linreg (model 1)
8   a1 = 25000/1000000
9   b1 = 0
10  linreg_dataset1_inputs, linreg_dataset1_outputs = linreg_matplotlib(a1, b1, min_surf, max_surf, n_points)
11  plt.plot(linreg_dataset1_inputs, linreg_dataset1_outputs, 'r', label = "Model 1")
12
13  # Another linreg (model 2)
14  a2 = 17000/1000000
15  b2 = 50000/1000000
16  linreg_dataset2_inputs, linreg_dataset2_outputs = linreg_matplotlib(a2, b2, min_surf, max_surf, n_points)
17  plt.plot(linreg_dataset2_inputs, linreg_dataset2_outputs, 'g', label = "Model 2")
18
19  # A final linreg (model 3)
20  a3 = 14500/1000000
21  b3 = 95000/1000000
22  linreg_dataset3_inputs, linreg_dataset3_outputs = linreg_matplotlib(a3, b3, min_surf, max_surf, n_points)
23  plt.plot(linreg_dataset3_inputs, linreg_dataset3_outputs, 'k', label = "Model 3")
24
25  # Display
26  plt.legend(loc = 'best')
27  plt.show()
```

Which model is the best fit for this task in your opinion? (red, green or black?)

Why is one better than another?

The need for a loss function

Which model is the best fit for this task in your opinion? (red, green or black?)

Why is one better than another?

Telling which model is best can prove difficult, especially for close values of (a, b).

Good news, there is one last element we have not used.

4. **Loss (L)** function to measure the performance of the current model and its current parameters on said task and dataset.

The need for a loss function

Definition (Loss function):

A **loss function** (also known as a **cost function**) is a mathematical function that measures the difference between the predicted output of a model and the true output we should be matching.

It is used to

- **evaluate the performance of the model,**

- And more specifically **evaluate how good our choice of trainable parameters** (here, the parameters are a and b) are at fitting the data.

It is used to guide the optimization process during **training**.

The need for a loss function

Definition (The Mean Square Error loss function):

In our linear model, a good loss function we could use consists of the **Mean Square Error (MSE) loss function**.

The MSE is calculated by calculating the square difference between:

- The **prediction** $p_i(x_i, a, b) = ax_i + b$ formed by the model for some given inputs x_i,

- The **true value** y_i that we should be matching.

We then repeat this operation for every sample i, and average those errors together, i.e.

$$L(a, b, x, y) = \frac{1}{N} \sum_{i=1}^{N} (p(x_i, a, b) - y_i)^2 = \frac{1}{N} \sum_{i=1}^{N} (ax_i + b - y_i)^2$$

The need for a loss function

Recall the definition of the MSE loss function

$$L(a, b, x, y) = \frac{1}{N} \sum_{i=1}^{N} (ax_i + b - y_i)^2$$

In [9]:
```python
1  # Mean square error as a loss function
2  # Displaying loss using exponential notation (XXXe-YYY)
3  def loss_mse(a, b, x, y):
4      val =  np.sum((y - (a*x + b))**2)/x.shape[0]
5      return '{:.2e}'.format(val)
```

In [10]:
```python
1  # The lower the loss function, the better the linear regression values (a, b) fit the dataset.
2  loss1 = loss_mse(a1, b1, inputs, outputs)
3  loss2 = loss_mse(a2, b2, inputs, outputs)
4  loss3 = loss_mse(a3, b3, inputs, outputs)
5  print(loss1, loss2, loss3)
```

1.01e+00 5.11e-02 7.85e-03

Legend:
- Lin. Reg. model 1, using a = 0.025, b = 0.0, with loss = 1.01e+00
- Lin. Reg. model 2, using a = 0.017, b = 0.05, with loss = 5.11e-02
- Lin. Reg. model 3, using a = 0.0145, b = 0.095, with loss = 7.85e-03

Which model is the best fit for this task in your opinion? (red, green or black?)

Why is one better than another?

Answer: Black model, because it has a lower loss value than the green and red ones!

Training a model

Definition (Training a model):

Training a model consists of **finding the best values for the trainable parameters of the chosen model**, for the given task and dataset.

By best values for the trainable parameters, we mean the values that **minimize the chosen loss function**. It is therefore an optimization problem of some sort!

Once a minimal value is obtained on the loss function, then we can claim that the "optimal" trainable parameters have been found.

Or in other words, that the model has been trained.

Training a linear regression model

In the case of linear regression, training the model consists of solving the following optimization problem.

$$(a^*, b^*) = \underset{a,b}{\operatorname{argmin}} \left[L(a, b, x, y) \right]$$

That is, using the explicit formula of the loss function,

$$(a^*, b^*) = \underset{a,b}{\operatorname{argmin}} \left[\frac{1}{N} \sum_{i=1}^{N} (ax_i + b - y_i)^2 \right]$$

Training a linear regression model

Definition (The normal equation for linear regression):

Solving this optimization problem can be done analytically, as it can be proven that the optimal choice of parameters $W^* = \begin{pmatrix} b^* \\ a^* \end{pmatrix}$ follows the **normal equation**, below.

$$W^* = (X^T X)^{-1} X^T Y$$

With

$$X = \begin{pmatrix} 1 & x_1 \\ \dots & \dots \\ 1 & x_N \end{pmatrix}$$

And

$$Y = \begin{pmatrix} y_1 \\ \dots \\ y_N \end{pmatrix}$$

Finally, X^T denotes the transposed matrix of X and $(X)^{-1}$ its inverse.

Proof of the Normal Equation

Using the notations

$$X = \begin{pmatrix} 1 & x_1 \\ \dots & \dots \\ 1 & x_N \end{pmatrix},$$

$$Y = \begin{pmatrix} y_1 \\ \dots \\ y_N \end{pmatrix},$$

$$W^* = \begin{pmatrix} b^* \\ a^* \end{pmatrix}$$

We have, in matrix form

$$XW^* = \begin{pmatrix} a^*x_1 + b^* \\ \dots \\ a^*x_N + b^* \end{pmatrix}$$

And, then it follows that

$$XW^* - Y = \begin{pmatrix} a^*x_1 + b^* - y_1 \\ \dots \\ a^*x_N + b^* - y_N \end{pmatrix}$$

Proof of the Normal Equation

Our loss function, the MSE L, is therefore defined, in matrix format as:

$$L = \frac{1}{N}(XW^* - Y)^T(XW^* - Y)$$

This is equivalent to

$$L = \frac{1}{N}\left((XW^*)^T - Y^T\right)(XW^* - Y)$$

Proof of the Normal Equation

We then get

$$L = \frac{1}{N}\left(W^{*^T}X^T - Y^T\right)(XW^* - Y)$$

Expanding

$$L = \frac{1}{N}\left(W^{*^T}X^TXW^* - Y^TXW^* - W^{*^T}X^TY + Y^TY\right)$$

Proof of the Normal Equation

Recall that we are trying to find the solution to optimization problem

$$(a^*, b^*) = \arg \min_{a,b}[L(a, b, x, y)]$$

With

$$L = \frac{1}{N}(W^{*T}X^TXW^* - Y^TXW^* - W^{*T}X^TY + Y^TY)$$

Recognizing that $Y^TXW^* = W^{*T}X^TY$, differentiating our function L with respect to W and equating it to zero gives:

$$\frac{\partial L}{\partial W} = \frac{1}{N}(2X^TXW - 2X^TY) = 0$$

Proof of the Normal Equation

This is equivalent to

$$(X^T X W^* - X^T Y) = 0$$

Or,

$$X^T X W^* = X^T Y$$

Assuming $X^T X$ is invertible (and that might not always be the case based on your dataset!), this finally gives:

$$W^* = (X^T X)^{-1} X^T Y$$

Training a linear regression model

Normal Equation: $W^* = (X^T X)^{-1} X^T Y$

```python
1  # Defining the X matrix, following the notation above, as a numpy array.
2  X = np.array([[1, x_i] for x_i in inputs])
3  print(X)
```

```python
1  # While we are at it, let us define the transposed version of X.
2  XT = np.transpose(X)
3  print(XT)
```

```python
1  # Defining the Y matrix, following the notation above, as a numpy array.
2  Y = np.array([[y_i] for y_i in outputs])
3  print(Y)
```

```python
1  # Defining W_star according to our formula.
2  W_star = np.matmul(np.linalg.inv(np.matmul(XT,X)), np.matmul(XT,Y))
3  print(W_star)
4  # Extracting a_star and b_star values from W_star.
5  b_star, a_star = W_star[0, 0], W_star[1, 0]
6  print("Optimal a_star value: ", a_star)
7  print("The value we used for a in the mock dataset generation: ", 14373/1000000)
8  print("Optimal b_star value: ", b_star)
9  print("The value we used for b in the mock dataset generation: ", 100000/1000000)
```

Legend:
- Lin. Reg. model 1, using a = 0.025, b = 0.0, with loss = 1.01e+00
- Lin. Reg. model 2, using a = 0.017, b = 0.05, with loss = 5.11e-02
- Lin. Reg. model 3, using a = 0.0145, b = 0.095, with loss = 7.85e-03
- Optimal Lin. Reg. model, using a* = 0.014844435285058159, b* = 0.06941742189819247, with loss = 7.65e-03

Price (in millions SGD)

Surface (in sqm)

Purple line is the optimal set of parameters for our model on the given task and dataset! Black was pretty close!

The problem with the normal equation

The normal equation $W^* = (X^T X)^{-1} X^T Y$ immediately gives the optimal set of parameters (a^*, b^*) to use for linear regression.

- However, it can become **computationally expensive when the number of samples (or features) in X is large**.

- It might even be **impossible to use when the matrix product of feature variables $(X^T X)^{-1}$ is not invertible**.

Problem: In these cases, the normal equation can be slow to compute, or it may not even have a solution.

Another problem: More sophisticated models than linear regression will NOT have a normal equation, anyway.

Gradient Descent, to the rescue!

Definition (Gradient Descent):

Gradient Descent (GD) is an iterative algorithm used to solve optimization problems.

- It starts with an initial non-optimal set of parameters for our model.

- It then repeatedly updates the parameters in the direction of the negative gradient of the given cost function, until it converges to a local minimum.

Forgot about GD?
Have a look at your notes from 10.013 M&A!

Gradient Descent, to the rescue!

Definition (Gradient Descent):

Gradient Descent (GD) is an iterative algorithm used to solve optimization problems.

- It starts with an initial non-optimal set of parameters for our model.

- It then repeatedly updates the parameters in the direction of the negative gradient of the given cost function, until it converges to a local minimum.

- The normal equation has many problems...

- But GD, on the other hand can be used to find the optimal solution even when the normal equation is not applicable.

- While sub-optimal, the main advantage of gradient descent is that it can handle very large datasets and it can also be used for non-linear models.

GD in Linear Regression

In the case of Linear Regression, with the MSE loss defined earlier, we have

$$(a^*, b^*) = \underset{a,b}{\operatorname{argmin}} \left[L(a, b, x, y) = \frac{1}{N} \sum_{i=1}^{N} (ax_i + b - y_i)^2 \right]$$

GD in Linear Regression

We have the following derivatives with respect to a and b:

$$D_a = \frac{\partial L}{\partial a} = \frac{-2}{N} \sum_{i=1}^{N} x_i \left(y_i - (a\,x_i + b) \right)$$

$$D_b = \frac{\partial L}{\partial b} = \frac{-2}{N} \sum_{i}^{N} y_i - (a\,x_i + b)$$

GD in Linear Regression

The gradient descent **update rules** are therefore defined as

$$a \leftarrow a - \alpha D_a$$

$$a \leftarrow a + \frac{2\alpha}{N} \sum_{i=1}^{N} x_i \left(y_i - (a\, x_i + b) \right)$$

And

$$b \leftarrow b - \alpha D_b$$

$$b \leftarrow b + \frac{2\alpha}{N} \sum_{i}^{N} y_i - (a\, x_i + b)$$

With parameter α being the **learning rate** for the gradient descent, a parameter to be decided manually, later.

GD in Linear Regression

Gradient Descent Linear Regression procedure:

- We initialize a and b with some values (could be zero, random, or something else).

- For a given number of iterations, we will apply the two update rules defined earlier.

- (Optionally, we might decide to stop iterating, if we realize that the values of a and b are no longer changing.
  This is called **early stopping**, and is typically implemented by tracking the changes on each iteration.
  We would then decide to break the GD loop if the changes on a and b are less than a threshold δ.)

```python
def gradient_descent_linreg(a_0, b_0, x, y, alpha = 1e-5, delta = 1e-5, max_count = 1000):
    # Define the initial values of a and b as a_0 and b_0
    a, b = a_0, b_0
    # Define N as the number of elements in the dataset
    N = len(x)
    # Keep track of how much a and b changed on each iteration
    change = float("Inf")
    # Counter as safety to prevent infinite looping
    counter = 0
    # List of losses, to be used for display later
    losses = []
    while change > delta:
        # Helper to visualize iterations of while loop
        print("-----")
        # Use gradident descent update rules for a and b
        D_a = -2/N*(sum([x_i*(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        D_b = -2/N*(sum([(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        a = a - alpha*D_a
        b = b - alpha*D_b
        print("Gradients: ", D_a, D_b)
        print("New values for (a, b): ", a, b)
        # Compute change
        change = max(abs(alpha*D_a), abs(alpha*D_b))
        print("Change: ", change)
        # COmpute and display current loss
        loss = loss_mse(a, b, x, y)
        losses.append(float(loss))
        print("Loss: ", loss)
        # Counter update, will break if iterations number exceeds max_count,
        # to prevent gradient descent from going indefinitely.
        # (Just a safety measure, for good practice, we would definitely prefer to see
        # the while loop break "naturally", because change eventually fell under the threshold stop.)
        counter += 1
        if(counter > max_count):
            print("Maximal number of iterations reached.")
            break
    return a, b, losses
```

Initialize a and b as you please.

```python
def gradient_descent_linreg(a_0, b_0, x, y, alpha = 1e-5, delta = 1e-5, max_count = 1000):
    # Define the initial values of a and b as a_0 and b_0
    a, b = a_0, b_0
    # Define N as the number of elements in the dataset
    N = len(x)
    # Keep track of how much a and b changed on each iteration
    change = float("Inf")
    # Counter as safety to prevent infinite looping
    counter = 0
    # List of losses, to be used for display later
    losses = []
    while change > delta:
        # Helper to visualize iterations of while loop
        print("-----")
        # Use gradident descent update rules for a and b
        D_a = -2/N*(sum([x_i*(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        D_b = -2/N*(sum([(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        a = a - alpha*D_a
        b = b - alpha*D_b
        print("Gradients: ", D_a, D_b)
        print("New values for (a, b): ", a, b)
        # Compute change
        change = max(abs(alpha*D_a), abs(alpha*D_b))
        print("Change: ", change)
        # COmpute and display current loss
        loss = loss_mse(a, b, x, y)
        losses.append(float(loss))
        print("Loss: ", loss)
        # Counter update, will break if iterations number exceeds max_count,
        # to prevent gradient descent from going indefinitely.
        # (Just a safety measure, for good practice, we would definitely prefer to see
        # the while loop break "naturally", because change eventually fell under the threshold stop.)
        counter += 1
        if(counter > max_count):
            print("Maximal number of iterations reached.")
            break
    return a, b, losses
```

Number of samples in dataset.

```python
def gradient_descent_linreg(a_0, b_0, x, y, alpha = 1e-5, delta = 1e-5, max_count = 1000):
    # Define the initial values of a and b as a_0 and b_0
    a, b = a_0, b_0
    # Define N as the number of elements in the dataset
    N = len(x)
    # Keep track of how much a and b changed on each iteration
    change = float("Inf")
    # Counter as safety to prevent infinite looping
    counter = 0
    # List of losses, to be used for display later
    losses = []
    while change > delta:
        # Helper to visualize iterations of while loop
        print("-----")
        # Use gradident descent update rules for a and b
        D_a = -2/N*(sum([x_i*(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        D_b = -2/N*(sum([(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        a = a - alpha*D_a
        b = b - alpha*D_b
        print("Gradients: ", D_a, D_b)
        print("New values for (a, b): ", a, b)
        # Compute change
        change = max(abs(alpha*D_a), abs(alpha*D_b))
        print("Change: ", change)
        # COmpute and display current loss
        loss = loss_mse(a, b, x, y)
        losses.append(float(loss))
        print("Loss: ", loss)
        # Counter update, will break if iterations number exceeds max_count,
        # to prevent gradient descent from going indefinitely.
        # (Just a safety measure, for good practice, we would definitely prefer to see
        # the while loop break "naturally", because change eventually fell under the threshold stop.)
        counter += 1
        if(counter > max_count):
            print("Maximal number of iterations reached.")
            break
    return a, b, losses
```

Two possible stopping conditions, change < delta or counter > max_count.

```python
def gradient_descent_linreg(a_0, b_0, x, y, alpha = 1e-5, delta = 1e-5, max_count = 1000):
    # Define the initial values of a and b as a_0 and b_0
    a, b = a_0, b_0
    # Define N as the number of elements in the dataset
    N = len(x)
    # Keep track of how much a and b changed on each iteration
    change = float("Inf")
    # Counter as safety to prevent infinite looping
    counter = 0
    # List of losses, to be used for display later
    losses = []
    while change > delta:
        # Helper to visualize iterations of while loop
        print("-----")
        # Use gradient descent update rules for a and b
        D_a = -2/N*(sum([x_i*(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        D_b = -2/N*(sum([(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        a = a - alpha*D_a
        b = b - alpha*D_b
        print("Gradients: ", D_a, D_b)
        print("New values for (a, b): ", a, b)
        # Compute change
        change = max(abs(alpha*D_a), abs(alpha*D_b))
        print("Change: ", change)
        # COmpute and display current loss
        loss = loss_mse(a, b, x, y)
        losses.append(float(loss))
        print("Loss: ", loss)
        # Counter update, will break if iterations number exceeds max_count,
        # to prevent gradient descent from going indefinitely.
        # (Just a safety measure, for good practice, we would definitely prefer to see
        # the while loop break "naturally", because change eventually fell under the threshold stop.)
        counter += 1
        if(counter > max_count):
            print("Maximal number of iterations reached.")
            break
    return a, b, losses
```

Update using our GD update rules from earlier.

```python
def gradient_descent_linreg(a_0, b_0, x, y, alpha = 1e-5, delta = 1e-5, max_count = 1000):
    # Define the initial values of a and b as a_0 and b_0
    a, b = a_0, b_0
    # Define N as the number of elements in the dataset
    N = len(x)
    # Keep track of how much a and b changed on each iteration
    change = float("Inf")
    # Counter as safety to prevent infinite looping
    counter = 0
    # List of losses, to be used for display later
    losses = []
    while change > delta:
        # Helper to visualize iterations of while loop
        print("-----")
        # Use gradident descent update rules for a and b
        D_a = -2/N*(sum([x_i*(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        D_b = -2/N*(sum([(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        a = a - alpha*D_a
        b = b - alpha*D_b
        print("Gradients: ", D_a, D_b)
        print("New values for (a, b): ", a, b)
        # Compute change
        change = max(abs(alpha*D_a), abs(alpha*D_b))
        print("Change: ", change)
        # COmpute and display current loss
        loss = loss_mse(a, b, x, y)
        losses.append(float(loss))
        print("Loss: ", loss)
        # Counter update, will break if iterations number exceeds max_count,
        # to prevent gradient descent from going indefinitely.
        # (Just a safety measure, for good practice, we would definitely prefer to see
        # the while loop break "naturally", because change eventually fell under the threshold stop.)
        counter += 1
        if(counter > max_count):
            print("Maximal number of iterations reached.")
            break
    return a, b, losses
```

Compute change on this iteration (to decide on early stopping if changes become insignificant).

```python
def gradient_descent_linreg(a_0, b_0, x, y, alpha = 1e-5, delta = 1e-5, max_count = 1000):
    # Define the initial values of a and b as a_0 and b_0
    a, b = a_0, b_0
    # Define N as the number of elements in the dataset
    N = len(x)
    # Keep track of how much a and b changed on each iteration
    change = float("Inf")
    # Counter as safety to prevent infinite looping
    counter = 0
    # List of losses, to be used for display later
    losses = []
    while change > delta:
        # Helper to visualize iterations of while loop
        print("-----")
        # Use gradident descent update rules for a and b
        D_a = -2/N*(sum([x_i*(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        D_b = -2/N*(sum([(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        a = a - alpha*D_a
        b = b - alpha*D_b
        print("Gradients: ", D_a, D_b)
        print("New values for (a, b): ", a, b)
        # Compute change
        change = max(abs(alpha*D_a), abs(alpha*D_b))
        print("Change: ", change)
        # COmpute and display current loss
        loss = loss_mse(a, b, x, y)
        losses.append(float(loss))
        print("Loss: ", loss)
        # Counter update, will break if iterations number exceeds max_count,
        # to prevent gradient descent from going indefinitely.
        # (Just a safety measure, for good practice, we would definitely prefer to see
        # the while loop break "naturally", because change eventually fell under the threshold stop.)
        counter += 1
        if(counter > max_count):
            print("Maximal number of iterations reached.")
            break
    return a, b, losses
```

Calculate new MSE value using the new parameters of a and b. We also keep track of these losses by adding them to a list, used for display later.

```python
def gradient_descent_linreg(a_0, b_0, x, y, alpha = 1e-5, delta = 1e-5, max_count = 1000):
    # Define the initial values of a and b as a_0 and b_0
    a, b = a_0, b_0
    # Define N as the number of elements in the dataset
    N = len(x)
    # Keep track of how much a and b changed on each iteration
    change = float("Inf")
    # Counter as safety to prevent infinite looping
    counter = 0
    # List of losses, to be used for display later
    losses = []
    while change > delta:
        # Helper to visualize iterations of while loop
        print("-----")
        # Use gradident descent update rules for a and b
        D_a = -2/N*(sum([x_i*(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        D_b = -2/N*(sum([(y_i - (a*x_i + b)) for x_i, y_i in zip(x, y)]))
        a = a - alpha*D_a
        b = b - alpha*D_b
        print("Gradients: ", D_a, D_b)
        print("New values for (a, b): ", a, b)
        # Compute change
        change = max(abs(alpha*D_a), abs(alpha*D_b))
        print("Change: ", change)
        # COmpute and display current loss
        loss = loss_mse(a, b, x, y)
        losses.append(float(loss))
        print("Loss: ", loss)
        # Counter update, will break if iterations number exceeds max_count,
        # to prevent gradient descent from going indefinitely.
        # (Just a safety measure, for good practice, we would definitely prefer to see
        # the while loop break "naturally", because change eventually fell under the threshold stop.)
        counter += 1
        if(counter > max_count):
            print("Maximal number of iterations reached.")
            break
    return a, b, losses
```

**Return trained parameters and losses evolution on each iteration.**

```
1  a_gd, b_gd, losses = gradient_descent_linreg(a_0 = 0, b_0 = 0, x = inputs, y = outputs, alpha = 1e-5, delta = 1e-6)
```

```
-----
Gradients:  -342.3996226384001 3.105429920000001
New values for (a, b):  0.0034239962263840013 -3.105429920000001e-05
Change:  0.0034239962263840013
Loss:  1.61e+00
-----
Gradients:  -266.62829357559235 2.4212214483389536
New values for (a, b):  0.006090279162139925 -5.526651368338955e-05
Change:  0.0026662829357559236
Loss:  9.82e-01
-----
Gradients:  -207.62478780001163 1.8884249597020164
New values for (a, b):  0.008166527040140042 -7.415076328040972e-05
Change:  0.0020762478780001164
Loss:  5.99e-01
-----
Gradients:  -161.6784683033283 1.4735337252735503
New values for (a, b):  0.009783311723173324 -8.888610053314522e-05
Change:  0.001616784683033283
```

**Observation: Loss is decreasing over iterations
(which is good news for us!)**

```
1  # Display dataset
2  plt.figure(figsize = (10, 5))
3  plt.plot(losses, label = "Model losses over gradient descent iterations")
4
5  # Display
6  plt.legend(loc = 'best')
7  plt.show()
```


Observation: Confirmed with a plot.
Loss is decreasing over iterations
(which is good news for us!)

Checking for optimal parameters

We have generated the dataset ourselves, so we know the optimal values for a and b!

```
1  print("Optimal a_star value: ", a_star)
2  print("Value for a_star, found by gradient descent: ", a_gd)
3  print("We used 14373/1000000 in the mock dataset generation, which is: ", 14373/1000000)
4  print("Optimal b_star value: ", b_star)
5  print("Value for b_star, found by gradient descent: ", b_gd)
6  print("The value we used in the mock dataset generation: ", 100000/1000000)
```

```
Optimal a_star value:  0.0148444435285058159
Value for a_star, found by gradient descent:  0.015466847033579204
We used 14373/1000000 in the mock dataset generation, which is:  0.014373
Optimal b_star value:  0.06941742189819247
Value for b_star, found by gradient descent:  0.00014426903274128027
The value we used in the mock dataset generation:  0.1
```

Observation: Close match, good enough for us. (Not perfect due to randomness in the dataset generation!)

Optimal Lin. Reg. model, using a* = 0.014844435285058159, b* = 0.06941742189819247, with loss = 7.65e-03

Gradient Descent Lin. Reg. model, using a* = 0.015466847033579204, b* = 0.00014426903274128027, with loss = 8.12e-03

Using Sklearn for Linear regression

In practice, we never implement the linear regression model ourselves (but it is a good practice to try it at least once!)

- It is often better/faster to rely on the **sklearn** library and use the **LinearRegression** object!

```
1  # Creating a sklearn Linear Regression model.
2  # It uses the same analytical formula from earlier, i.e. W^* = (X^T X)^{-1} X^T Y.
3  reg = LinearRegression().fit(sk_inputs, sk_outputs)
4  # The coefficients for a* and b* are found using coeff_ and intercept_ respectively.
5  a_sk = reg.coef_[0]
6  b_sk = reg.intercept_
7  print("Optimal a_star value: ", a_star)
8  print("Value for a_star, found by sklearn: ", a_sk)
9  print("Optimal b_star value: ", b_star)
10 print("Value for b_star, found by sklearn: ", b_sk)
```

```
Optimal a_star value:  0.014844435285058159
Value for a_star, found by sklearn:  0.014844435285058166
Optimal b_star value:  0.06941742189819247
Value for b_star, found by sklearn:  0.069417421898189556
```

Predicting using Sklearn Linear regression

After training, it is also good practice to check if the predictor makes sense, by **testing** it and asking it to predict the price of an apartment it has never seen before.

• Confirm the pertinence of predicted values manually, if possible.

```
1  # We can later use this Linear Regression model, to predict the price of
2  # a new appartment with surface 105 sqm (price in millions SGD).
3  new_surf = 105
4  pred_price = reg.predict(np.array([[new_surf]]))[0]
5  print(pred_price)
```

1.628083126829297

```
1  avg_price = 14373*new_surf + 100000
2  min_val = 0.9*avg_price
3  max_val = 1.1*avg_price
4  print("Min, max, avg prices: ", [min_val, max_val, avg_price])
```

Min, max, avg prices: [1448248.5, 1770081.5000000002, 1609165]