

50.039 Theory and Practice of Deep Learning

W1-S3 Introduction and Machine Learning Reminders

Matthieu De Mari



SINGAPORE UNIVERSITY OF
TECHNOLOGY AND DESIGN

About this week (Week 1)

1. What are the **typical concepts of Machine Learning** to be used as a starting point for this course?
2. What are the **different families of problems** in Deep Learning?
3. What is the **typical structure of a Deep Learning problem**?
4. What is **linear regression** and how to implement it?
5. What is the **gradient descent algorithm** and how is it used to **train Machine Learning models**?
6. What is **polynomial regression** and how to implement it?
7. What is **regularization** and how to implement it in **Ridge regression**?

About this week (Week 1)

- 8. What is **overfitting** and why is it bad?
- 9. What is **underfitting** and why is it bad?
- 10. What is **generalization** and how to evaluate it?
- 11. What is a **train-test split** and why is it related to **generalization**?
- 12. What is a **sigmoid** function? What is a **logistic** function?
- 13. How to perform **binary classification** using a **logistic regressor** and how is it related to linear regression?

About this week (Week 1)

14. What are **Neural Networks** and how do they relate to the **biology of a human brain**?
15. What is a **Neuron** in a Neural Network and how does it relate to linear/logistic regression?
16. What is the **difference** between a **shallow** and a **deep neural network**?
17. How to **implement a shallow Neural Network** manually and define a **forward propagation** method for it?
18. (How to **train a shallow Neural Network** using **backpropagation**?
How to define **backward propagation** and **trainer** functions?)

Introducing the sigmoid function

Definition (the sigmoid function):

The **sigmoid function** is an important function used in Machine Learning. It is defined as

$$s(x) = \frac{1}{1 + \exp(-x)}$$

Or, equivalently

$$s(x) = \frac{\exp(x)}{1 + \exp(x)}$$

Note: Sometimes, the notation $\sigma(x)$ is used instead of $s(x)$.


The sigmoid function has the following properties (to be proven as an extra exercise!):

$$\forall x, \quad 0 < s(x) < 1$$

$$\lim_{x \rightarrow -\infty} s(x) = 0$$

$$\lim_{x \rightarrow \infty} s(x) = 1$$

$$s(0) = 0.5$$



Also introducing two logistic functions

Definition (the logistic functions):

Similarly, let us introduce two **logistic functions**, below:

$$l(x) = \ln(x)$$

$$l_2(x) = \ln(1 - x)$$

These two functions also have interesting properties, shown on the right.
(*Prove them as an extra practice?*)


$$\lim_{x \rightarrow 0} l(x) = -\infty$$

$$\lim_{x \rightarrow 1} l(x) = 0$$

$$\lim_{x \rightarrow 0} l_2(x) = 0$$

$$\lim_{x \rightarrow 1} l_2(x) = -\infty$$

$$l(0.5) = l_2(0.5) = -\ln(2)$$



To recap: functions and implementations

Definition (the **sigmoid function):**

$$s(x) = \frac{1}{1 + \exp(-x)}$$

Or, equivalently

$$s(x) = \frac{\exp(x)}{1 + \exp(x)}$$

Definition (the **logistic functions):**

$$l(x) = \ln(x)$$

And,

$$l_2(x) = \ln(1 - x)$$

```
1 def s(x):  
2     return 1/(1 + np.exp(-x))  
3 def l(x):  
4     return np.log(x)  
5 def l2(x):  
6     return np.log(1 - x)
```

Results on mixes for our two logistic functions

In addition, the following properties hold (some of them require using L'Hospital rule to prove them, try it out?).

$$\lim_{x \rightarrow 0} x \ln(x) = 0$$

$$\lim_{x \rightarrow 1} x \ln(x) = 0$$

$$\lim_{x \rightarrow 0} (1 - x) \ln(1 - x) = 0$$

$$\lim_{x \rightarrow 1} (1 - x) \ln(1 - x) = 0$$

Results on mixes for our two logistic functions


In addition, the following properties hold (some of them require using L'Hospital rule to prove them, try it out?).

$$\lim_{x \rightarrow 0} (1 - x) \ln(x) = -\infty$$

$$\lim_{x \rightarrow 1} (1 - x) \ln(x) = 0$$

$$\lim_{x \rightarrow 0} x \ln(1 - x) = 0$$


$$\lim_{x \rightarrow 1} x \ln(1 - x) = -\infty$$



New mock dataset generation

In notebook 6, we will prepare a dataset for a **single-variable binary classification**.

- The **inputs**, x , will consist of chest sizes for males between 94cm and 114cm, randomly generated.
- The **outputs**, y , will consists of **two classes**, corresponding to the shirt size to use, with two possible values:
 - **M** (class with value **0**, if chest size is below **104 cm**) size
 - and **L** size (class with value **1**, if chest size is above **104 cm**).



```

1 # All helper functions
2 def chest_size(min_size, max_size):
3     return round(np.random.uniform(min_size, max_size), 2)
4 def shirt_size(size, threshold):
5     return int(size >= threshold)
6 def generate_datasets(n_points, min_size, max_size, threshold):
7     x = np.array([chest_size(min_size, max_size) for _ in range(n_points)])
8     y = np.array([shirt_size(i, threshold) for i in x])
9     return x, y

```

```


1 # Dataset generation (n_points points will be generated).
2 # We will use a seed for reproducibility.
3 min_size = 94
4 max_size = 114
5 threshold = 104
6 np.random.seed(27)
7 n_points = 50
8 inputs, outputs = generate_datasets(n_points, min_size, max_size, threshold)
9 print(inputs)
10 print(outputs)

```

```

[102.51 110.29 108.71 111.36 101.67 113.59 111.86  98.19 108.84 107.26
 111.74 111.16 108.99 111.4   97.74 100.51 101.46 109.87  97.02  97.4
  95.62 100.1   109.67  97.26  95.41 108.02  97.62 105.98 102.31 104.27
  98.41 108.51 110.99 112.58 108.72 103.53 103.86 105.89  95.52  96.35
 113.33 105.67  95.85  94.27 110.74 112.3   108.09 101.74 108.11 112.47]
[0 1 1 1 0 1 1 0 1 1 1 1 1 1 0 0 0 1 0 0 0 0 1 0 0 1 0 1 0 1 0 1 1 1 1 0 0
 1 0 0 1 1 0 0 1 1 1 0 1 1]

```




A quick word about expert systems

Definition (**expert systems**):

An **expert system** is a computer program that mimics the decision-making abilities of a human expert in a specific domain or field.

It uses knowledge represented in a symbolic form, such as rules or ontologies, to make inferences and decisions based on the data it receives.

Expert systems are “the dream” and provide a high level of performance and accuracy in solving problems or making decisions. But these are difficult to design as they require a strong, and not necessarily obvious, human expertise.



A quick word about expert systems

Definition (**expert systems**):

An **expert system** is a computer program that mimics the decision-making abilities of a human expert in a specific domain or field.

It uses knowledge represented in a symbolic form, such as rules or ontologies, to make inferences and decisions based on the data it receives.

Important note: in our case, we could define an expert system as shown below, but most of the time the datasets will not have a clear logic and you will have to rely on ML algorithms instead, basically *“hoping the AI will figure out a logic that works”*.

```
1 # A simple expert system
2 def expert_system(inputs):
3     return [int(size > 104) for size in inputs]
```

```
1 # Try it and compare results to see we have 100% accuracy here!
2 pred = expert_system(inputs)
3 print((pred == outputs).all())
```

From Linear Regression to Logistic Regression

Definition(**Logistic Regression**):

The **logistic regression** model assumes that the classes y_i for every sample x_i are connected via the formula below.

$$\forall i, y_i = \begin{cases} 1 & \text{if } p(x_i) > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

This value $p(x)$ is often referred to as the **probability** of sample with value x being of class 1 (that is a L size shirt).

Respectively, $1 - p(x)$ is the **probability** of sample with value x being of class 0 (that is M size).

The **probability** function is then simply defined as

$$p(x) = s(ax + b)$$

The function p always has values between 0 and 1, thanks to the **sigmoid** function s .

An adjusted Logistic Regression

Definition(Logistic Regression):

The **logistic regression** model assumes that the classes y_i for every sample x_i are connected via the formula below.

$$\forall i, y_i = \begin{cases} 1 & \text{if } s(ax_i + b) > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

This value $s(ax_i + b)$ is often referred to as the **probability** of sample with value x_i being of class 1 (that is a L size shirt).

This model has **two trainable parameters**, a and b , to be decided as in the linear regression.

In fact, we could see the logistic regression as the **combination of linear regression and sigmoid together**.

In addition, note that the function $f(x) = ax + b$ is also often referred to as the **logit** function with value x .

From Linear Regression to Logistic Regression

We can implement the logistic regression below and try it out with some values a and b !

```

1 def logistic_regression(x, a, b):
2     return s(a*x + b)
3 def predict_samples(inputs, a, b):
4     return [int(logistic_regression(x, a, b) >= 0.5) for x in inputs]
5 def logistic_regression_matplotlib(a, b, min_size, max_size):
6     x_plt = np.linspace(min_size, max_size, 50)
7     y_plt = np.array([s(a*x + b) for x in x_plt])
8     return x_plt, y_plt

```

```

1 # Trying to predict with our logistic regression
2 # model, for two given values of a and b.
3 a = 1.7
4 b = -172
5 pred_y = predict_samples(inputs, a, b)
6 print(pred_y)


```

```

[1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0,
1, 1, 0, 0, 1, 1, 1, 1, 1, 1]

```

Restricted



Restricted

Training a logistic regressor

As before, in linear regression, we will have to **train the model**, i.e. choose values for a and b , that fit the dataset.

- While **MSE** seemed to work as a **loss function** to train the **linear regression** model, it is not that good for logistic regression.
- We need something better!

In logistic regression, we prefer to **minimize** a loss function, called the **log-likelihood (cross-entropy)** function, instead.

- This function has to do with the two **logistic functions** we have introduced earlier.

The log-likelihood function and loss

Definition (the log-likelihood function):

The **log-likelihood function** is defined as

$$L = \frac{1}{N} \sum_i^N [y_i \ln(s(ax_i + b)) + (1 - y_i) \ln(1 - s(ax_i + b))]$$

This is a good function to use for our problem for the two reasons:

- **When $p(x_i)$ is close to the ground truth value y_i , both terms of the loss function, will have values close to 0.** This follows from the properties we defined earlier for the mixes of logistic functions.
- **On the opposite, the log-likelihood loss function will have a different behaviour when $p(x_i)$ and the ground truth value y_i are different, producing large non-zero negative values instead.**

The log-likelihood function and loss

Definition (using the **log-likelihood function as a loss):**

This log-likelihood loss function is therefore a good choice of a performance/loss metric to measure the performance of our models.

Indeed, maximizing this log-likelihood loss function and bringing it to zero would be equivalent to finding the best choice of parameters a and b , and therefore the best fit.

In practice however, we **prefer to minimize loss functions**. This requires a simple adjustment, **multiplying the log-likelihood loss function by -1**.

$$a^*, b^* = \arg \min_{a, b} \left[\frac{-1}{N} \sum_i^N [y_i \ln(s(ax_i + b)) + (1 - y_i) \ln(1 - s(ax_i + b))] \right]$$

Implementing the log-likelihood loss

We can simply implement the log likelihood loss as shown below.

- Note: We could look for the “**normal equation**” for the logistic regression, but this simply will not work. It is simply impossible to solve the logistic regression problem analytically.
- (Not convinced? Try solving it analytically!).

```
1 def log_likelihood_loss(a, b, x, y):  
2     # Using Numpy broadcasting  
3     return -1*np.mean(y*np.log(logistic_regression(x, a, b)) \  
4                        + (1 - y)*np.log(1 - logistic_regression(x, a, b)))
```

Training a logistic regressor

Instead, it is often preferable to rely on the gradient descent method, or the Newton method.

- Typically, this is what the **LogisticRegression()** object in sklearn uses for training!
- (Newton method is not covered but feel free to look it up/try it out?)

```

1 # Reshaping inputs as a 2D array
2 # (As before, this is a requirement for sklearn models)
3 inputs_re = inputs.reshape(-1, 1)
4
5 # Create a logistic regression model, use the Newton method
6 # to find the best parameters a and b
7 logreg_model = LogisticRegression(solver = 'newton-cg')
8 logreg_model.fit(inputs_re, outputs)
9
10 # Test your model and inputs
11 pred_outputs = logreg_model.predict(inputs_re)
12 print(pred_outputs)
13 print(outputs)
14
15 # Retrieving coefficients a and b
16 a_sk = logreg_model.coef_[0, 0]
17 b_sk = logreg_model.intercept_[0]
18 print(a_sk, b_sk)


```

```

[0 1 1 1 0 1 1 0 1 1 1 1 1 1 0 0 0 1 0 0 0 0 1 0 0 1 0 1 0 1 0 1 1 1 0 0
 1 0 0 1 1 0 0 1 1 1 0 1 1]
[0 1 1 1 0 1 1 0 1 1 1 1 1 1 0 0 0 1 0 0 0 0 1 0 0 1 0 1 0 1 0 1 1 1 1 0 0
 1 0 0 1 1 0 0 1 1 1 0 1 1]
1.413470868291155 -147.3822278881603

```

Restricted



Restricted

Evaluating the logistic regression model

When attempting to evaluate a classification model, it is often a good idea to first display a confusion matrix for the predicted values of the dataset samples (test samples, preferably).

- We can simply calculate it by using the `confusion_matrix()` function from `sklearn`. The non-diagonal elements have zero values, which is the sign that our model classifies perfectly.
- Learn more, here: https://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html.

```
1 # Using a confusion matrix.
2 print(confusion_matrix(outputs, pred_outputs))
```

```
[[22  0]
 [ 0 28]]
```

```
1 # Using a classification report.
2 print(classification_report(outputs, pred_outputs))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	22
1	1.00	1.00	1.00	28
accuracy			1.00	50
macro avg	1.00	1.00	1.00	50
weighted avg	1.00	1.00	1.00	50

Evaluating the logistic regression model

We can also ask for a classification report for the trained model.

- This shows the precision, recall and F1 scores for each class, as well as the overall accuracy of the model.
- For each one of these interesting metrics, the best value is 1.0. Our model achieves a 1.0 score, suggesting that the model classifies perfectly.
- Learn more, here: https://scikit-learn.org/stable/modules/generated/sklearn.metrics.classification_report.html#sklearn.metrics.classification_report.

```
1 # Using a confusion matrix.
2 print(confusion_matrix(outputs, pred_outputs))
```

```
[[22  0]
 [ 0 28]]
```

```
1 # Using a classification report.
2 print(classification_report(outputs, pred_outputs))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	22
1	1.00	1.00	1.00	28
accuracy			1.00	50
macro avg	1.00	1.00	1.00	50
weighted avg	1.00	1.00	1.00	50


Note: the details of these metrics and what they represented is not discussed here (covered in previous ML classes).

A bit of biology

Many tasks we have seen earlier (e.g. appt prices) are “difficult” to solve, or do not have an obvious “explainable” logic.

Our brain, however, seems very capable to solve these tasks.

Million dollar question: Could we replace the linear/polynomial/logistic models with a model that mimics the brain?




Very easy for a human...
But, how would we explain the logic of it to a computer?

It's a cat!

A bit of biology

Question: In neural networks, we will implement neurons, but are they really close to the ones in the human brain?


- Neural networks are inspired by the neural structure and processes of the human brain.
- Neural networks are comprised of interconnected nodes, which are called **artificial neurons** or **perceptrons**.
- Neural networks are related to biology in that they are modelled after biological neural systems.



A bit of biology

In general,

- A **neuron will receive stimuli** through its dendrites, in the form of electrical signals playing the role of **inputs**,
- It will **process the inputs** in a certain way in its nucleus (and other parts of the neuron),
- If the integrated signals reach a threshold, **the neuron generates an action potential**, a brief electrical impulse. This is an output coming from the axon terminal, in the form of a signal (1) or no signal (0) being produced.




In other words,

Neuron
=
Logistic Regression?!

A bit of biology


- But a human brain is not just one neuron, it is many of them arranged into a network of some sort!
- Eventually, the output signal produced by a neuron will be passed to other neurons in the brain.
- In fact, the brain could be seen as a set of several neurons processing information, **both in parallel and in sequence.**
- This assembly, which resembles the human brain, is commonly referred to as a **neural network.**



A bit of biology

Research has shown that artificial neural networks (NN) could be used to solve some of the problems that biological NN can solve.

- For example, neural networks can be used to identify objects in images, just as biological neural networks can.
- The training process of CS neural networks is apparently very similar to the biological neural networks.
- This is actually an ongoing argument/debate in the community [Quanta2021] and [MITNews2022].



Starting with a mock dataset, again

As before, we will try to come up with a model capable of predicting the price of an apartment based on some of its features.

The inputs will this time consist of two parameters:

- the **surface** of the apartment, in sqm, just like before,
- and the **distance** between the apartment and the closest MRT station in meters.

We will generate a mock dataset by:

- randomly drawing **surfaces** between 40 sqm and 200 sqm,
- randomly drawing **distances** between 50 and 1000m,
- generate prices by assuming that the average price is simply defined as
$$avgprice = 100000 + 14373 \times surface + (1000 - distance) \times 1286$$

Finally, we will randomly apply a -/+ 10% variation on the average price to create some variance, and call that the output for a given sample.

Starting with a mock dataset, again

```
1 # All helper functions
2 min_surf = 40
3 max_surf = 200
4 def surface(min_surf, max_surf):
5     return round(np.random.uniform(min_surf, max_surf), 2)
6 min_dist = 50
7 max_dist = 1000
8 def distance(min_dist, max_dist):
9     return round(np.random.uniform(min_dist, max_dist), 2)
10 def price(surface, distance):
11     return round((100000 + 14373*surface + (1000 - distance)*1286)*(1 + np.random.uniform(-0.1, 0.1)))/1000000
12 n_points = 100
13 def create_dataset(n_points, min_surf, max_surf, min_dist, max_dist):
14     surfaces_list = np.array([surface(min_surf, max_surf) for _ in range(n_points)])
15     distances_list = np.array([distance(min_dist, max_dist) for _ in range(n_points)])
16     inputs = np.array([[s, d] for s, d in zip(surfaces_list, distances_list)])
17     outputs = np.array([price(s, d) for s, d in zip(surfaces_list, distances_list)]).reshape(n_points, 1)
18     return surfaces_list, distances_list, inputs, outputs
```

```



1 # Generate dataset
2 np.random.seed(47)
3 surfaces_list, distances_list, inputs, outputs = create_dataset(n_points, \
4                                                                 min_surf, \
5                                                                 max_surf, \
6                                                                 min_dist, \
7                                                                 max_dist)
8 # Check a few entries of the dataset
9 print(surfaces_list.shape)
10 print(distances_list.shape)
11 print(inputs.shape)
12 print(outputs.shape)
13 print(inputs[0:10, :])
14 print(outputs[0:10])



```

```

(100,)
(100,)
(100, 2)
(100, 1)
[[ 58.16 572.97]
 [195.92 809.8 ]
 [156.6  349.04]
 [ 96.23  86.82]
 [153.22 817.92]
 [167.94 806.25]
 [143.29 315.92]
 [106.34 482.67]
 [152.96 427.77]
 [ 79.46 955.76]]
[[1.581913]
 [3.450274]
 [2.978769]
 [2.808258]
 [2.556398]
 [3.023983]
 [3.099523]
 [2.121069]
 [3.136544]
 [1.273443]]

```





$$avgprice = 100000 + 14373 \times surface + (1000 - distance) \times 1286$$

 is actually the equation of a 2D plane in a 3D ambient space!

Coding a Neuron and a Neural Network


Definition (**Neuron** in Neural Networks):

Each artificial neuron in a neural network might for instance be represented as a linear regression (or later on, a logistic regression) model, which receives input, applies a transformation on that input, and produces an output.

And similar to the way that biological neurons are connected and communicate with each other, artificial neurons are connected to each other, in parallel and in sequence, through a network. This neural network can be thought of as pathways for information to flow.

(Later on, we will see more advanced ways to represent neurons than linear or logistic regression.)

Coding our first Neuron



```
1 # Let us define matrix W first, with two parameters w1 and w2.
2 # And the matrix b as well.
3 # Are these values randomly defined?
4 W = np.array([[14373/1000000], [-1286/1000000]]) # 14373/1000000 and -1286/1000000
5 b = np.ones(shape = (1, 1))*1.386 # 100000/1000000 + 1000*1286/1000000
```

```
1 # We can then implement the multiplication operation
2 # we defined above by using the matmul operation.
3 Z = np.matmul(inputs, W)
4 pred = Z + b
5 # We can then check the shape of the matrices we obtained
6 print(inputs.shape)
7 print(W.shape)
8 print(b.shape)
9 print(Z.shape)
10 print(pred.shape)
11 print(outputs.shape)
```

(100, 2)

(2, 1)

(1, 1)

(100, 1)

(100, 1)

(100, 1)

Defining a minimal Neural Network

Definition (**defining a minimal Neural Network**):

In order to define a minimal neural network with only one neuron, we need two things:

- First, an **__init__ method**, listing the **trainable parameters** for the model. In our case that is a **weight vector** $W = (w_1, w_2)$ with 2 elements and a single scalar **bias** value b .
- Second, a **forward method**, which will be used to formulate predictions for any set of given inputs. The **prediction formula** is simply $y_i = w_1 s_i + w_2 d_i + b$.

(For now, this single neuron is equivalent to a Linear Regression!)

```

1 class SimpleNeuralNet():
2
3     def __init__(self, W, b):
4         # Weights and biases matrices
5         self.W = W
6         self.b = b
7
8     def forward(self, x):
9         # wx + b operation as above
10        Z = np.matmul(x, self.W)
11        pred = Z + self.b
12        return pred

```

Note: You might find implementations that are using $XW + b$ as a formula for the forward method, and some implementations that are using $WX^T + b$ instead. Sometimes even omitting the “transpose” and using $WX + b$. This is roughly equivalent. What matters is that you are aware of the correct sizes for the different matrices and how to combine correctly.

```

1 # Create a minimal neuron neural network using our class
2 simple_neural_net = SimpleNeuralNet(W = np.array([[14373/1000000], [-1286/1000000]]), \
3                                         b = np.ones(shape = (1, 1))*1.386)
4 print(simple_neural_net.__dict__)

```

```

{'W': array([[ 0.014373],
              [-0.001286]]), 'b': array([[1.386]])}

```

```

1 # Predict and show first five samples
2 pred = simple_neural_net.forward(inputs)
3 print(pred[:5])

```

```

[[1.48509426]
 [3.16055536]
 [3.18794636]
 [2.65746327]
 [2.53638594]]

```

Adding a loss

This takes care of the **first three elements** of a machine learning problem (task, dataset, model).

The **fourth** element we need to define is a **loss function**.

- As this is a regression task, we can reuse the **MSE loss** and add this MSE loss function as a method to our class.
- For convenience, we will compute the MSE loss in matrix form.
- We will also reuse the **forward()** method we just coded to make predictions and then compared said predictions with the expected outputs from our dataset.

Adding a loss

```
1 class SimpleNeuralNet():
2
3     def __init__(self, W, b):
4         # Weights and biases matrices
5         self.W = W
6         self.b = b
7         # Loss, initialized as infinity before first calculation is made
8         self.loss = float("Inf")
9
10    def forward(self, inputs):
11        Z = np.matmul(inputs, self.W)
12        pred = Z + self.b
13        return pred
14
15    def MSE_loss(self, inputs, outputs):
16        outputs_re = outputs.reshape(-1, 1)
17        pred = self.forward(inputs)
18        losses = (pred - outputs_re)**2
19        self.loss = np.sum(losses)/outputs.shape[0]
20        return self.loss
```


Scaling up with more parameters

Technically, our Neural Network could operate with any number of inputs N_x and any number of outputs N_y .

We would simply need to adjust the sizes of the matrices of parameters:

- W as a $N_x \times N_y$ matrix,
- B as a N_y vector.

Note: This would be equivalent to having N_y neurons, each receiving N_x inputs, and each producing one output to be assembled as the Z vector.




Scaling up with more layers

Definition (**Shallow Neural Network**):

Following the logic of our SimpleNeuralNet, we could implement a Shallow Neural Network. **It will include two processing layers of neurons in a row, instead of just one, and each layer will be implementing its own $WX + b$ operation of some sort.**


Scaling up with more layers



Scaling up with more layers

The first layer, will receive inputs with dimensionality n_x and produce outputs with dimensionality n_h .


- The matrix W for this layer will therefore be of size $n_x \times n_h$, and the matrix b will be a simple 1D vector with size n_h .



Scaling up with more layers

The second layer, also called **hidden layer**, will receive the outputs from the previous layer with dimensionality n_h , use them as inputs and will produce outputs matching the dimensionality of the outputs in our dataset, n_y .

- The matrix W for this layer will be 2D $n_h \times n_y$, and b will be 1D n_y .



Our NN class

Small changes:

- Weights and biases are now **randomly generated**, instead of being passed to the init method.
- We pass the **expected sizes** n_x, n_h, n_y instead.
- And initialize parameters matrices as normal random with zero mean and variance 0.1.

```

1 class ShallowNeuralNet():
2
3     def __init__(self, n_x, n_h, n_y):
4         # Network dimensions
5         self.n_x = n_x
6         self.n_h = n_h
7         self.n_y = n_y
8         # Weights and biases matrices
9         self.W1 = np.random.randn(n_x, n_h)*0.1
10        self.b1 = np.random.randn(1, n_h)*0.1
11        self.W2 = np.random.randn(n_h, n_y)*0.1
12        self.b2 = np.random.randn(1, n_y)*0.1
13        # Loss, initialized as infinity before first calculation is made
14        self.loss = float("Inf")
15
16    def forward(self, inputs):
17        # Wx + b operation for the first layer
18        Z1 = np.matmul(inputs, self.W1)
19        Z1_b = Z1 + self.b1
20        # Wx + b operation for the second layer
21        Z2 = np.matmul(Z1_b, self.W2)
22        Z2_b = Z2 + self.b2
23        return Z2_b
24
25    def MSE_loss(self, inputs, outputs):
26        # MSE loss function as before
27        outputs_re = outputs.reshape(-1, 1)
28        pred = self.forward(inputs)
29        losses = (pred - outputs_re)**2
30        self.loss = np.sum(losses)/outputs.shape[0]
31        return self.loss

```

```

1 # Define neural network structure
2 n_x = 2
3 n_h = 4
4 n_y = 1
5 shallow_neural_net = ShallowNeuralNet(n_x, n_h, n_y)
6 print(shallow_neural_net.__dict__)

```

```

{'n_x': 2, 'n_h': 4, 'n_y': 1, 'W1': array([[ -0.10476816,  0.18570216,  0.03204007, -0.10951262],
      [-0.13867874, -0.03539496, -0.02856421,  0.20592501]]), 'b1': array([[ 0.0232776 , -0.16122469,  0.00718537,  0.0666335
  1]]), 'W2': array([[ 0.03321156],
      [-0.0336505 ],
      [ 0.04977554],
      [-0.1794089 ]]), 'b2': array([[0.03460341]]), 'loss': inf}

```

```

1 pred = shallow_neural_net.forward(inputs)
2 print(pred.shape)
3 print(outputs.shape)
4 print(pred[0:5])
5 print(outputs[0:5])

```

```

(100, 1)
(100, 1)
[[-23.24055489]
 [-31.54945952]
 [-12.75105332]
 [-2.49026451]
 [-32.3803654 ]]
[[1.581913]
 [3.450274]
 [2.978769]
 [2.808258]
 [2.556398]]

```

1. Initialize model

2. Forward to predict

3. Compute loss to evaluate model
(it is bad, because trainable parameters have random values!)

```

1 loss = shallow_neural_net.MSE_loss(inputs, outputs)
2 print(loss)

```

677.625448852107

The need for a training procedure

At the moment, we have coded a model that can initialize trainable parameters randomly, formulate predictions and evaluate its own performance.

- It is great, but the current models are garbage in terms of loss values. We can only try a few different initialization and pray the RNG gods to give us trainable parameters with a good loss.
- **We need a training procedure!**
- **What could it be...?**

```
1 np.random.seed(963)
2 shallow_neural_net1 = ShallowNeuralNet(n_x, n_h, n_y)
3 loss1 = shallow_neural_net1.MSE_loss(inputs, outputs)
4 shallow_neural_net2 = ShallowNeuralNet(n_x, n_h, n_y)
5 loss2 = shallow_neural_net2.MSE_loss(inputs, outputs)
6 shallow_neural_net3 = ShallowNeuralNet(n_x, n_h, n_y)
7 loss3 = shallow_neural_net3.MSE_loss(inputs, outputs)
8 shallow_neural_net4 = ShallowNeuralNet(n_x, n_h, n_y)
9 loss4 = shallow_neural_net4.MSE_loss(inputs, outputs)
10 print(loss1, loss2, loss3, loss4)
```

21.318364917457647 58.190236579106184 4.770288142049728 0.2093294704434788

The backpropagation mechanism


Definition (**Backpropagation**):

Backpropagation is an **algorithm** used to **train** neural networks.

It is a type of **gradient descent** (what a surprise!) algorithm that allows the network to learn by **adjusting the weights** of the connections between the neurons in the network.

Introduced by **Hinton** and **Rumelhart** [Rumelhart1986].

(Or was it? [Medium2020])



The backpropagation mechanism

Definition (**Backpropagation**):

Backpropagation is an **algorithm** used to **train** neural networks.

It is a type of **gradient descent** (what a surprise!) algorithm that allows the network to learn by **adjusting the weights** of the connections between the neurons in the network.

Introduced by **Hinton** and **Rumelhart** [Rumelhart1986].

(Or was it? [Medium2020])

Backpropagation will tell our model **how to adjust** the values of the W and b matrices in each layer, to improve the neural network model loss and prediction capabilities.

It will be implemented in the **backward()** method of our model.

Backpropagation is heavy to manually calculate and implement...

Strongly encouraging to try and redo the math/implementation shown hereafter on your own!

That's the best way to learn and you can always refer to these slides for answers!

After math is correct, have a look at how to implement it in a backward method (also shown in slides, but probably will be shown rapidly in class!)

Remember, you have no homework this week, so let us call that your homework!

Backpropagation in our model

While our Neural Network seems to be complicated in its architecture, at the end of the day, training it consists of solving a basic **optimization problem**.

$$W_1^*, b_1^*, W_2^*, b_2^* = \arg \min_{W_1, b_1, W_2, b_2} \left[\frac{1}{M} \left(Y_{pred}(X, W_1, b_1, W_2, b_2) - Y \right)^2 \right]$$
$$W_1^*, b_1^*, W_2^*, b_2^* = \arg \min_{W_1, b_1, W_2, b_2} \left[\frac{1}{M} (W_2(W_1 X + b_1) + b_2 - Y)^2 \right]$$

With M being the number of samples in the dataset. We will use **gradient descent**, like we did for the Linear Regression, except that this time we have 4 parameters being matrices and a loss function that is a bit more complex.

Before we continue

Before we continue, a quick note about square matrices.

$$W_1^*, b_1^*, W_2^*, b_2^* = \arg \min_{W_1, b_1, W_2, b_2} \left[\frac{1}{M} \left(Y_{pred}(X, W_1, b_1, W_2, b_2) - Y \right)^2 \right]$$
$$W_1^*, b_1^*, W_2^*, b_2^* = \arg \min_{W_1, b_1, W_2, b_2} \left[\frac{1}{M} (W_2(W_1 X + b_1) + b_2 - Y)^2 \right]$$

Note: The notation $(A)^2$ is somewhat abusive when A is a matrix, the correct way would be to write $A \cdot A^T$.

But in the formula above, that would lead to a very large formula (too large for slides!)

Chain rule, to the rescue!

Let us denote the error term ϵ .

$$\epsilon = Y_{pred}(X, W_1, b_1, W_2, b_2) - Y = W_2(W_1X + b_1) + b_2 - Y.$$

Using the chain rule, we can simply compute $\frac{\partial L}{\partial W_2}$ as

$$\frac{\partial L}{\partial W_2} = \frac{\partial L}{\partial Y_{pred}} \frac{\partial Y_{pred}}{\partial W_2}$$

Chain rule, to the rescue!

Since we have

$$\frac{\partial L}{\partial Y_{pred}} = \frac{2}{M} Y_{pred}(X, W_1, b_1, W_2, b_2) - Y = \frac{2\epsilon}{M}$$

And

$$\frac{\partial Y_{pred}}{\partial W_2} = W_1 X + b_1$$

The gradient descent update rule for W_2 is then defined as:

$$W_2 \leftarrow W_2 - \frac{2\epsilon\alpha}{M} (W_1 X + b_1)$$

Chain rule, to the rescue!

We have just computed the formula for the gradient descent update rule for W_2 , which was defined as:

$$W_2 \leftarrow W_2 - \frac{2\epsilon\alpha}{M} (W_1 X + b_1)$$

Similarly, we can define the gradient descent update rule for b_2 as:

$$b_2 \leftarrow b_2 - \frac{2\epsilon\alpha}{M}$$

Chain rule, to the rescue!

Similarly, we have:

$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial Y_{pred}} \frac{\partial Y_{pred}}{\partial W_1} = \left(\frac{2\epsilon}{M} \right) (W_2 X)$$

$$\text{And } \frac{\partial L}{\partial b_1} = \frac{\partial L}{\partial Y_{pred}} \frac{\partial Y_{pred}}{\partial b_1} = \left(\frac{2\epsilon}{M} \right) (W_2)$$

And therefore, we have:

$$W_1 \leftarrow W_1 - \frac{2\epsilon\alpha}{M} W_2 X$$

$$b_1 \leftarrow b_1 - \frac{2\epsilon\alpha}{M} W_2$$

Now that we have all update rules...

We have figured the four update rules to use for backpropagation...

$$W_2 \leftarrow W_2 - \frac{2\epsilon\alpha}{M} (W_1 X + b_1)$$

$$b_2 \leftarrow b_2 - \frac{2\epsilon\alpha}{M}$$


$$W_1 \leftarrow W_1 - \frac{2\epsilon\alpha}{M} W_2 X$$

$$b_1 \leftarrow b_1 - \frac{2\epsilon\alpha}{M} W_2$$

Now it is about implementing two additional methods for our Neural Network class:

1. **Backward method:** It performs the calculation of gradients, in matrix form, and parameters adjustments using the gradient descent update rules shown on the left.
2. **Trainer method:** It repeats the backward method, reusing the gradient descent for-loop function from earlier in Notebook 2, until a maximal number of iterations is reached or convergence is seen.

Training procedure, in short.



Backpropagation in our model

Update rules for W_2 and b_2

$$\epsilon = W_2(W_1X + b_1) + b_2 - Y$$

$$W_2 \leftarrow W_2 - \frac{2\epsilon\alpha}{M}(W_1X + b_1)$$

$$b_2 \leftarrow b_2 - \frac{2\epsilon\alpha}{M}$$

```
def backward(self, inputs, outputs, alpha = 1e-5):
    # Get the number of samples in dataset
    m = inputs.shape[0]

    # Forward propagate
    Z1 = np.matmul(inputs, self.W1)
    Z1_b = Z1 + self.b1
    Z2 = np.matmul(Z1_b, self.W2)
    y_pred = Z2 + self.b2

    # Compute error term
    epsilon = y_pred - outputs

    # Compute the gradient for W2 and b2
    dL_dW2 = (2/m)*np.matmul(Z1_b.T, epsilon)
    dL_db2 = (2/m)*np.sum(epsilon, axis = 0, keepdims = True)

    # Compute the loss derivative with respect to the first layer
    dL_dZ1 = np.matmul(epsilon, self.W2.T)
```

Backpropagation in our model

Update rules for W_1 and b_1

$$W_1 \leftarrow W_1 - \frac{2\epsilon\alpha}{M} W_2 X$$

$$b_1 \leftarrow b_1 - \frac{2\epsilon\alpha}{M} W_2$$

```
# Compute the loss derivative with respect to the first layer  
dL_dZ1 = np.matmul(epsilon, self.W2.T)
```

```
# Compute the gradient for W1 and b1  
dL_dW1 = (2/m)*np.matmul(inputs.T, dL_dZ1)  
dL_db1 = (2/m)*np.sum(dL_dZ1, axis = 0, keepdims = True)
```

```
# Update the weights and biases using gradient descent  
self.W1 -= alpha*dL_dW1  
self.b1 -= alpha*dL_db1  
self.W2 -= alpha*dL_dW2  
self.b2 -= alpha*dL_db2
```

```
# Update Loss  
self.MSE_loss(inputs, outputs)
```

Backpropagation in our model

Effect of the backpropagation

- Every time we call the backward() method, our model will update its parameters W_1 , W_2 , b_1 and b_2 .
- After every call of backward(), our model seems to become better at the task in question, shown by improvements in loss values.
- This can be seen by looking at the loss values decreasing after each iteration of the backward() method.

```
1 # Define neural network structure
2 n_x = 2
3 n_h = 4
4 n_y = 1
5 np.random.seed(967)
6 shallow_neural_net = ShallowNeuralNet(n_x, n_h, n_y)
7 loss = shallow_neural_net.MSE_loss(inputs, outputs)
8 print(shallow_neural_net.loss)
```

13.1869714693353

```
1 shallow_neural_net.backward(inputs, outputs, 1e-5)
2 print(shallow_neural_net.loss)
```

8.576174190387835


```
1 shallow_neural_net.backward(inputs, outputs, 1e-5)
2 print(shallow_neural_net.loss)
```

5.818528741557871

```
1 shallow_neural_net.backward(inputs, outputs, 1e-5)
2 print(shallow_neural_net.loss)
```

4.14790610131927

Training procedure, in short (reminder).



Trainer function for our model

Train function implementation

- Iterate the backward() method for a number of iterations $N_{\max}$.
- Display new loss values and append them in a list for display.
- Also implemented an early stopping, which will break for loop if little to no change in parameters during an iteration of backward().
- Similar to GD in Notebook 2!

```
# Our trainer function
def train(shallow_neural_net, inputs, outputs, N_max = 1000, \
          alpha = 1e-5, delta = 1e-5, display = True):
    # List of losses, starts with the current loss
    losses_list = [shallow_neural_net.loss]
    # Repeat iterations
    for iteration_number in range(1, N_max + 1):
        # Backpropagate
        shallow_neural_net.backward(inputs, outputs, alpha)
        new_loss = shallow_neural_net.loss
        # Update losses list
        losses_list.append(new_loss)
        # Display
        if(display):
            print("Iteration {} - Loss = {}".format(iteration_number, new_loss))
        # Check for delta value and early stop criterion
        difference = abs(losses_list[-1] - losses_list[-2])
        if(difference < delta):
            if(display):
                print("Stopping early - loss evolution was less than delta.")
            break
    else:
        # Else on for loop will execute if break did not trigger
        if(display):
            print("Stopping - Maximal number of iterations reached.")
    return losses_list
```

Trainer function for our model


```
1 losses_list = train(shallow_neural_net, inputs, outputs, N_max = 10000, alpha = 1e-5, delta = 1e-6, display = True)
```

```
Iteration 1 - Loss = 4.639845655363769
Iteration 2 - Loss = 3.6055739235388646
Iteration 3 - Loss = 3.1219544752264055
Iteration 4 - Loss = 2.891370725402625
Iteration 5 - Loss = 2.778663798354082
Iteration 6 - Loss = 2.721333641596964
Iteration 7 - Loss = 2.6901646643755135
Iteration 8 - Loss = 2.671407441095785
Iteration 9 - Loss = 2.6585626629544303
Iteration 10 - Loss = 2.648548914291606
Iteration 11 - Loss = 2.639902127093024
Iteration 12 - Loss = 2.6319255476288577
Iteration 13 - Loss = 2.6242871305644138
Iteration 14 - Loss = 2.6168284145839156
Iteration 15 - Loss = 2.60947364178999
Iteration 16 - Loss = 2.6021864900028118
Iteration 17 - Loss = 2.5949494536811937
Iteration 18 - Loss = 2.58775401064343
Iteration 19 - Loss = 2.5805050707014652
```

Successful training as losses are decreasing over iterations?

Trainer function for our model

The curve below is called a **training curve**, it is useful to show how the training went (make an extra viz method to be used after training?).



Loss has decreased and
converged.
Confirmed training?

Finally, we will make the trainer function and the performance curves function into methods for our class.

```

66 def train(self, inputs, outputs, N_max = 1000, alpha = 1e-5, delta = 1e-5, display = True):
67     # List of losses, starts with the current loss
68     self.losses_list = [self.loss]
69     # Repeat iterations
70     for iteration_number in range(1, N_max + 1):
71         # Backpropagate
72         self.backward(inputs, outputs, alpha)
73         new_loss = self.loss
74         # Update losses list
75         self.losses_list.append(new_loss)
76         # Display
77         if(display):
78             print("Iteration {} - Loss = {}".format(iteration_number, new_loss))
79         # Check for delta value and early stop criterion
80         difference = abs(self.losses_list[-1] - self.losses_list[-2])
81         if(difference < delta):
82             if(display):
83                 print("Stopping early - loss evolution was less than delta.")
84             break
85         else:
86             # Else on for loop will execute if break did not trigger
87             if(display):
88                 print("Stopping - Maximal number of iterations reached.")
89
90 def show_losses_over_training(self):
91     # Initialize matplotlib
92     fig, axs = plt.subplots(1, 2, figsize = (15, 5))
93     axs[0].plot(list(range(len(self.losses_list))), self.losses_list)
94     axs[0].set_xlabel("Iteration number")
95     axs[0].set_ylabel("Loss")
96     axs[1].plot(list(range(len(self.losses_list))), self.losses_list)
97     axs[1].set_xlabel("Iteration number")
98     axs[1].set_ylabel("Loss (in logarithmic scale)")
99     axs[1].set_yscale("log")
100     # Display
101     plt.show()

```

Backpropagation is heavy to manually calculate and implement...
Strongly encouraging to try and redo the math/implementation shown on your own!

That's the best way to learn and you can always refer to these slides for answers!

After math is correct, have a look at how to implement it in a backward method (also shown in slides, but probably will be shown rapidly in class!)

Remember, you have no homework this week, so let us call that your homework!

Conclusion (Week 1)

- Reminders of Machine Learning
- Linear
- Using a normal equation to train a linear regression
- Gradient descent as a training procedure for linear regression
- Polynomial Regression
- Regularization in Ridge/Lasso Regression
- Train-test split
- Overfitting and underfitting
- Generalization
- Sigmoid and Logistic functions
- From linear regression to logistic regression
- Using logistic regression for binary classification

Conclusion (Week 1)

- Implementing a shallow neural network in Numpy
- Forward propagation method, to formulate predictions
- The backpropagation mechanism, as the gradient descent on Neural Network
- Backward method for training and trainer functions to iterate backward iterations
- Performance/training curves

Next week?

- Initializations to break symmetries
- Exploding and vanishing gradients
- Activation functions and non-linearities in Neural Networks
- Advanced optimizers
- Validation sets, early stopping, saver and loader functions

Learn more about these topics

Out of class, supporting papers, for those of you who are curious.

- [Géron2019] A. Géron , “How Neural Networks Work”, 2019.
- [Yamins2016] Yamins et al., “Deep neural networks are robust computational models of the human visual system”, 2016.
- [Kriegeskorte2013] Kriegeskorte and Kievit, “Neural Network Models of the Human Brain”, 2013.
- [Rumelhart1986] D.E. **Rumelhart**, G. **Hinton**, Williams, “Learning representations by back-propagating errors“, 1986.
<https://www.nature.com/articles/323533a0>

Learn more about these topics

Keep track of important names and follow their social media (Follow them on Scholar, Twitter, or whatever works for you!)

- **Geoffrey Hinton**: Emeritus Professor at **University of Toronto**, one of the three Godfathers of Deep Learning and **2018 Turing Award** winner (highest distinction in Computer Science), Nobel Prize Physics 2025. Formerly, Google.



<https://www.cs.toronto.edu/~hinton/>

<https://scholar.google.co.uk/citations?user=JicYPdAAAAAJ&hl=en>

- **David Rumelhart**: Former professor at **University of California**, credited for inventing **backpropagation** along with **Hinton**. Passed away in 2011.

<https://www.nytimes.com/2011/03/19/health/19rumelhart.html>


Twitter, the theater for AI/DL drama and big AI announcements



vibrant portrait painting of Salvador Dalí with a robotic half face



a shiba inu wearing a beret and black turtleneck



Learn more about these topics

Some extra (easy) reading and videos for those of you who are curious.

- [Quanta2021] “Artificial Neural Nets Finally Yield Clues to How Brains Learn”, 2021.

<https://www.quantamagazine.org/artificial-neural-nets-finally-yield-clues-to-how-brains-learn-20210218/>

- [MITNews2022] “Study urges caution when comparing neural networks to the brain”, 2022.

<https://news.mit.edu/2022/neural-networks-brain-function-1102>

- [Medium2020] “Who Invented Backpropagation? Hinton Says He Didn’t, but His Work Made It Popular”, 2020.

<https://medium.com/syncedreview/who-invented-backpropagation-hinton-says-he-didnt-but-his-work-made-it-popular-e0854504d6d1/>