50.039 Theory and Practice of Deep Learning
W2-S1 Neural Networks, Initializers, Optimizers and other Good Practices

Matthieu De Mari

About this week (Week 2)

1. What are the **typical initializers for trainable parameters** in Neural Networks?

2. What is **symmetry** in a Neural Network and why is it essential to **break it** with proper initializers?

3. What is the **exploding gradient** problem?

4. How to **spot and fix an exploding gradient** problem?

5. What is the **vanishing gradient** problem?

6. How to **spot and fix a vanishing gradient** problem?

7. Why are **activation functions** needed in Neural Networks?

8. What is the **universal approximation theorem**?

About this week (Week 2)

9. What are typical **examples** of **activation functions**? Which activations should I be using?

10. What are typical examples of **advanced optimizers**? (Adagrad, RMSProp, Adam, etc.).

11. How to **implement** said **optimizers** manually?

12. What is a **stochastic gradient descent** and what are its **benefits**?

13. How to implement the **stochastic gradient descent** manually?

14. What is a **mini-batch gradient descent** and what are its **benefits**?

15. How to implement the **mini-batch gradient descent** manually?

16. How to **combine all optimizers concepts** into a great optimizer?

About this week (Week 2)

17. What is the **no free lunch theorem**?

18. What is the **train-test-validation split**?
Why is it **good practice** to use an extra validation set?

19. What is the **early stopping** of optimizer concept?
Why is it **good practice** to use it?

20. What are **saver** and **loader** functions?
Why is it **good practice** to use them?

21. What are **other common good practices** when it comes to Neural Networks?

Our shallow NN with two processing layers
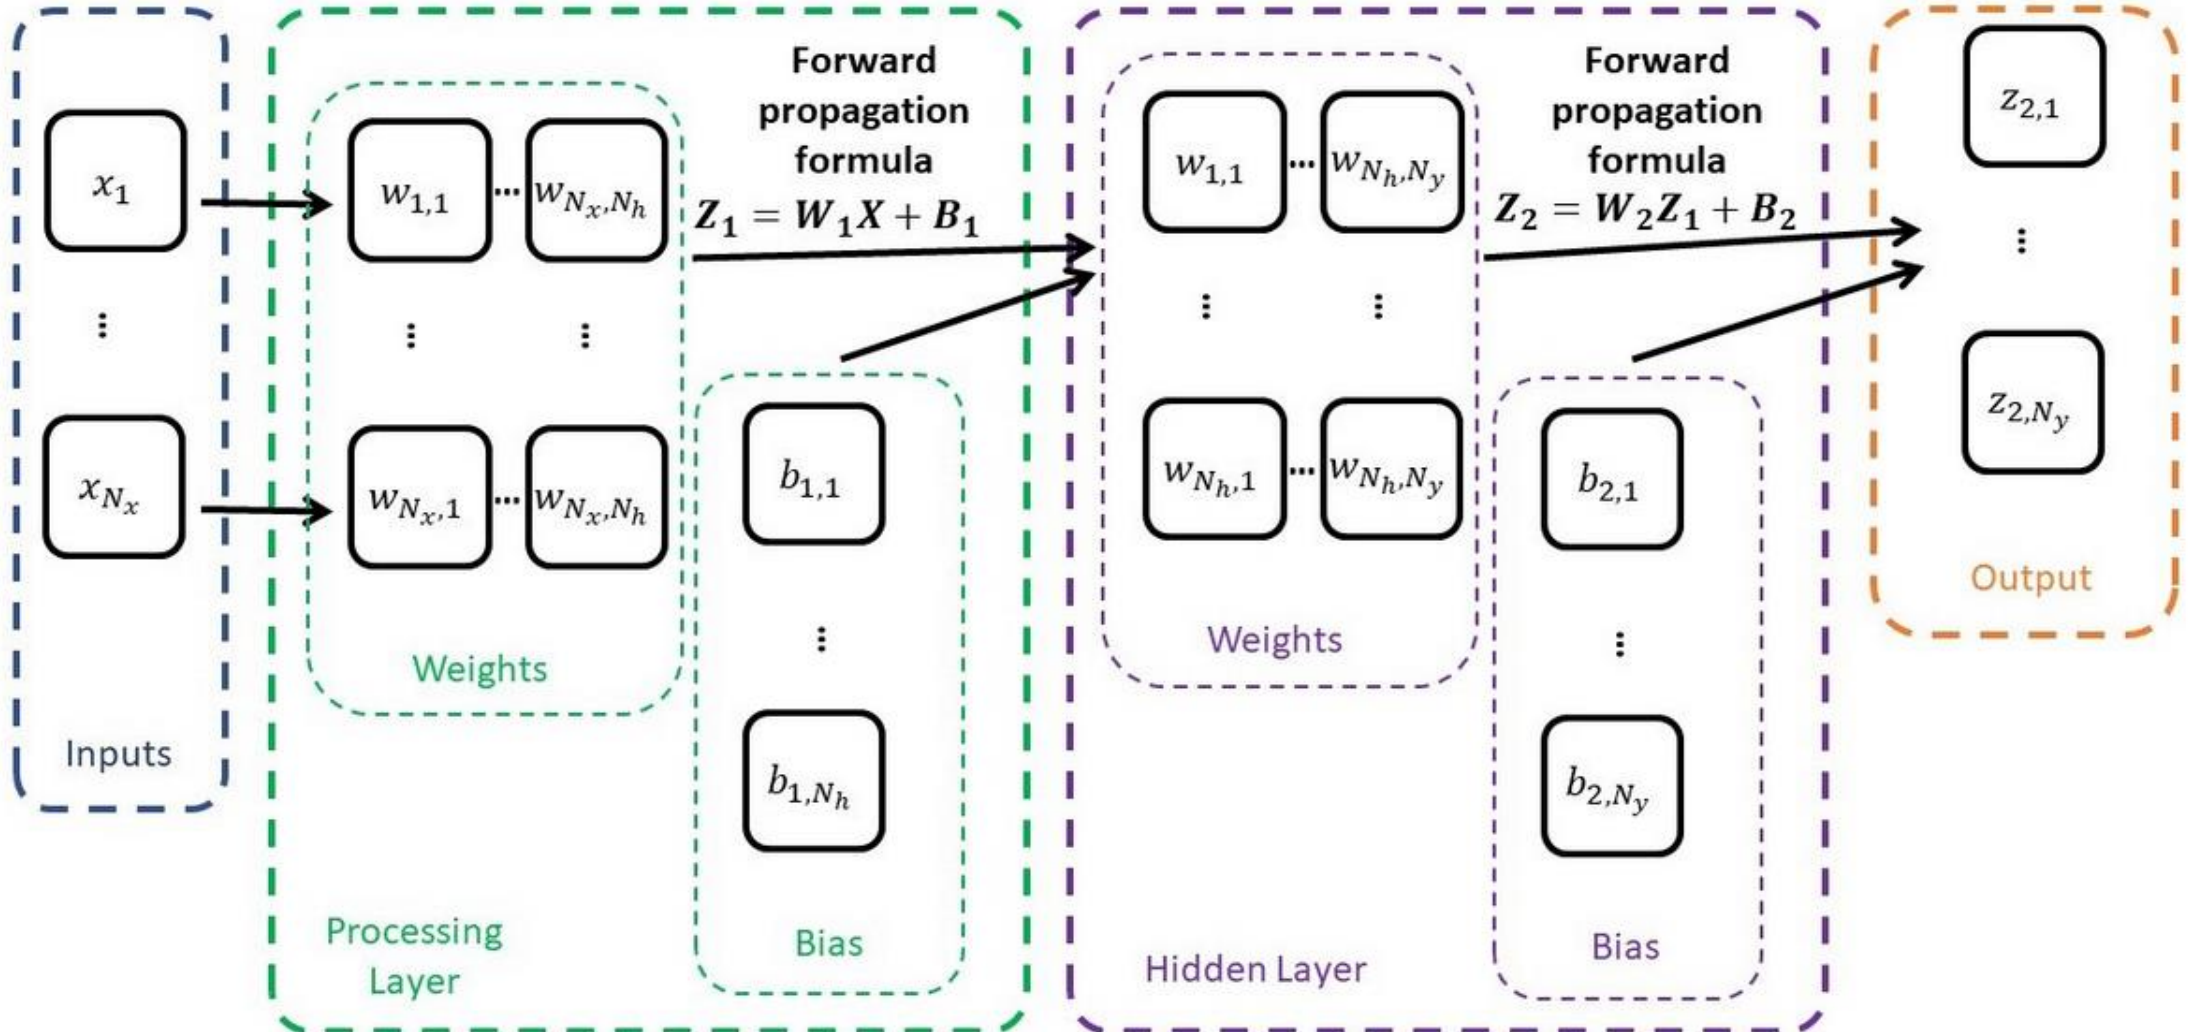
Our NN class

How it works:

- Weights and biases are now initialized as normal random with zero mean and variance 1.

- Forward method for making predictions using current parameters.

- Loss function for performance evaluation.

```python
class ShallowNeuralNet():

    def __init__(self, n_x, n_h, n_y):
        # Network dimensions
        self.n_x = n_x
        self.n_h = n_h
        self.n_y = n_y
        # Weights and biases matrices
        self.W1 = np.random.randn(n_x, n_h)*0.1
        self.b1 = np.random.randn(1, n_h)*0.1
        self.W2 = np.random.randn(n_h, n_y)*0.1
        self.b2 = np.random.randn(1, n_y)*0.1
        # Loss, initialized as infinity before first calculation is made
        self.loss = float("Inf")

    def forward(self, inputs):
        # Wx + b operation for the first layer
        Z1 = np.matmul(inputs, self.W1)
        Z1_b = Z1 + self.b1
        # Wx + b operation for the second layer
        Z2 = np.matmul(Z1_b, self.W2)
        Z2_b = Z2 + self.b2
        return Z2_b

    def MSE_loss(self, inputs, outputs):
        # MSE loss function as before
        outputs_re = outputs.reshape(-1, 1)
        pred = self.forward(inputs)
        losses = (pred - outputs_re)**2
        self.loss = np.sum(losses)/outputs.shape[0]
        return self.loss
```

```
1  # Define neural network structure
2  n_x = 2
3  n_h = 4
4  n_y = 1
5  shallow_neural_net = ShallowNeuralNet(n_x, n_h, n_y)
6  print(shallow_neural_net.__dict__)
```

```
{'n_x': 2, 'n_h': 4, 'n_y': 1, 'W1': array([[-0.10476816,  0.18570216,  0.03204007, -0.10951262],
       [-0.13867874, -0.03539496, -0.02856421,  0.20592501]]), 'b1': array([[ 0.0232776 , -0.16122469,  0.00718537,  0.0666335
1]]), 'W2': array([[ 0.03321156],
       [-0.0336505 ],
       [ 0.04977554],
       [-0.1794089 ]]), 'b2': array([[0.03460341]]), 'loss': inf}
```

```
1  pred = shallow_neural_net.forward(inputs)
2  print(pred.shape)
3  print(outputs.shape)
4  print(pred[0:5])
5  print(outputs[0:5])
```

```
(100, 1)
(100, 1)
[[-23.24055489]
 [-31.54945952]
 [-12.75105332]
 [ -2.49026451]
 [-32.3803654 ]]
[[1.581913]
 [3.450274]
 [2.978769]
 [2.808258]
 [2.556398]]
```

1. Initialize model
2. Forward to predict
3. Compute loss to evaluate model (it is bad, because trainable parameters have received random values!)

```
1  loss = shallow_neural_net.MSE_loss(inputs, outputs)
2  print(loss)
```

```
677.625448852107
```

Backpropagation in our model

Update rules for W_2 and b_2

(Obtained after using chain rule to compute derivatives wrt. parameters for loss)

$$W_2 \leftarrow W_2 - \frac{2\epsilon\alpha}{M}(W_1 X + b_1)$$

$$b_2 \leftarrow b_2 - \frac{2\epsilon\alpha}{M}$$

```python
def backward(self, inputs, outputs, alpha = 1e-5):
    # Get the number of samples in dataset
    m = inputs.shape[0]

    # Forward propagate
    Z1 = np.matmul(inputs, self.W1)
    Z1_b = Z1 + self.b1
    Z2 = np.matmul(Z1_b, self.W2)
    y_pred = Z2 + self.b2

    # Compute error term
    epsilon = y_pred - outputs

    # Compute the gradient for W2 and b2
    dL_dW2 = (2/m)*np.matmul(Z1_b.T, epsilon)
    dL_db2 = (2/m)*np.sum(epsilon, axis = 0, keepdims = True)

    # Compute the Loss derivative with respect to the first layer
    dL_dZ1 = np.matmul(epsilon, self.W2.T)
```

Backpropagation in our model

Update rules for W_1 and b_1

(Obtained after using chain rule to compute derivatives wrt. parameters for loss)

$$W_1 \leftarrow W_1 - \frac{2\epsilon\alpha}{M}W_2X$$

$$b_1 \leftarrow b_1 - \frac{2\epsilon\alpha}{M}W_2$$

```python
# Compute the Loss derivative with respect to the first layer
dL_dZ1 = np.matmul(epsilon, self.W2.T)

# Compute the gradient for W1 and b1
dL_dW1 = (2/m)*np.matmul(inputs.T, dL_dZ1)
dL_db1 = (2/m)*np.sum(dL_dZ1, axis = 0, keepdims = True)

# Update the weights and biases using gradient descent
self.W1 -= alpha*dL_dW1
self.b1 -= alpha*dL_db1
self.W2 -= alpha*dL_dW2
self.b2 -= alpha*dL_db2

# Update Loss
self.MSE_loss(inputs, outputs)
```

Training procedure, in short.

Forward pass on dataset samples

Compute errors and loss function

Use new loss value to compute gradients

Adjust parameters with backprop

Repeat until convergence or max number of iterations

Made that iterative process a trainer method for our class.

```python
    def train(self, inputs, outputs, N_max = 1000, alpha = 1e-5, delta = 1e-5, display = True):
        # List of losses, starts with the current loss
        self.losses_list = [self.loss]
        # Repeat iterations
        for iteration_number in range(1, N_max + 1):
            # Backpropagate
            self.backward(inputs, outputs, alpha)
            new_loss = self.loss
            # Update losses list
            self.losses_list.append(new_loss)
            # Display
            if(display):
                print("Iteration {} - Loss = {}".format(iteration_number, new_loss))
            # Check for delta value and early stop criterion
            difference = abs(self.losses_list[-1] - self.losses_list[-2])
            if(difference < delta):
                if(display):
                    print("Stopping early - loss evolution was less than delta.")
                break
        else:
            # Else on for loop will execute if break did not trigger
            if(display):
                print("Stopping - Maximal number of iterations reached.")

    def show_losses_over_training(self):
        # Initialize matplotlib
        fig, axs = plt.subplots(1, 2, figsize = (15, 5))
        axs[0].plot(list(range(len(self.losses_list))), self.losses_list)
        axs[0].set_xlabel("Iteration number")
        axs[0].set_ylabel("Loss")
        axs[1].plot(list(range(len(self.losses_list))), self.losses_list)
        axs[1].set_xlabel("Iteration number")
        axs[1].set_ylabel("Loss (in logarithmic scale)")
        axs[1].set_yscale("log")
        # Display
        plt.show()
```

Trainer function for our model

```
1  losses_list = train(shallow_neural_net, inputs, outputs, N_max = 10000, alpha = 1e-5, delta = 1e-6, display = True)
```

```
Iteration 1 - Loss = 4.639845655363769
Iteration 2 - Loss = 3.6055739235388646
Iteration 3 - Loss = 3.1219544752264055
Iteration 4 - Loss = 2.891370725402625
Iteration 5 - Loss = 2.778663798354082
Iteration 6 - Loss = 2.721333641596964
Iteration 7 - Loss = 2.6901646643755135
Iteration 8 - Loss = 2.671407441095785
Iteration 9 - Loss = 2.6585626629544303
Iteration 10 - Loss = 2.648548914291606
Iteration 11 - Loss = 2.6399021270930924
Iteration 12 - Loss = 2.6319255476288577
Iteration 13 - Loss = 2.6242871305644138
Iteration 14 - Loss = 2.6168284145839156
Iteration 15 - Loss = 2.60947364178999
Iteration 16 - Loss = 2.6021864900028118
Iteration 17 - Loss = 2.5949494536811937
Iteration 18 - Loss = 2.58775401064343
Iteration 19 - Loss = 2.5805950207914653
```

A bad initialization

In our previous version of the model, we initialized our parameters randomly.

- While it seemed unimportant, this is actually essential.

- **Initializing parameters as constants would in fact be a terrible idea.**

- **But we have to try it at least once to understand why.**

```python
class ShallowNeuralNet():

    def __init__(self, n_x, n_h, n_y):
        # Network dimensions
        self.n_x = n_x
        self.n_h = n_h
        self.n_y = n_y
        # Weights and biases matrices (randomly initialized)
        self.W1 = np.random.randn(n_x, n_h)*0.1
        self.b1 = np.random.randn(1, n_h)*0.1
        self.W2 = np.random.randn(n_h, n_y)*0.1
        self.b2 = np.random.randn(1, n_y)*0.1
        # Loss, initialized as infinity before first calculation is made
        self.loss = float("Inf")
```

```python
class ShallowNeuralNet_ConstantInit():

    def __init__(self, n_x, n_h, n_y, const_val):
        # Network dimensions
        self.n_x = n_x
        self.n_h = n_h
        self.n_y = n_y
        # Weights and biases matrices (all initialized as constant
        # matrices filled with 0.1 values)
        self.W1 = np.ones(shape = (n_x, n_h))*const_val
        self.b1 = np.ones(shape = (1, n_h))*const_val
        self.W2 = np.ones(shape = (n_h, n_y))*const_val
        self.b2 = np.ones(shape = (1, n_y))*const_val
        # Loss, initialized as infinity before first calculation is made
        self.loss = float("Inf")
```

Training the constant initialization model

```
1  # Define and train neural network structure
2  n_x = 2
3  n_h = 4
4  n_y = 1
5  np.random.seed(37)
6  shallow_neural_net = ShallowNeuralNet_ConstantInit(n_x, n_h, n_y, 0.1)
7  shallow_neural_net.train(inputs, outputs, N_max = 10000, alpha = 1e-6, delta = 1e-6, display = True)
```

```
Iteration 1 - Loss = 618.4538740049891
W1, b1, W2, b2:
[[0.09938446 0.09938446 0.09938446 0.09938446]
 [0.09676805 0.09676805 0.09676805 0.09676805]]
[[0.09999521 0.09999521 0.09999521 0.09999521]]
[[0.09614772]
 [0.09614772]
 [0.09614772]
 [0.09614772]]
[[0.09995206]]
Iteration 501 - Loss = 1.3726340052732087
W1, b1, W2, b2:
[[0.09064391 0.09064391 0.09064391 0.09064391]
 [0.04434246 0.04434246 0.04434246 0.04434246]]
[[0.09992629 0.09992629 0.09992629 0.09992629]]
[[0.0145902]
 [0.0145902]
 [0.0145902]
 [0.0145902]]
```

Training the constant initialization model

```python
1  # Define and train neural network structure
2  n_x = 2
3  n_h = 4
4  n_y = 1
5  np.random.seed(37)
6  shallow_neural_net = ShallowNeuralNet_ConstantInit(n_x, n_h, n_y, 0.1)
7  shallow_neural_net.train(inputs, outputs, N_max = 10000, alpha = 1e-6, delta = 1e-6, display = True)
```

```
Iteration 1 - Loss = 618.4538740049891
W1, b1, W2, b2:
[[0.09938446 0.09938446 0.09938446 0.09938446]
 [0.09676805 0.09676805 0.09676805 0.09676805]]
[[0.09999521 0.09999521 0.09999521 0.09999521]]
[[0.09614772]
 [0.09614772]
 [0.09614772]
 [0.09614772]]
[[0.09995206]]
Iteration 501 - Loss = 1.3726340052732087
W1, b1, W2, b2:
[[0.09064391 0.09064391 0.09064391 0.09064391]
 [0.04434246 0.04434246 0.04434246 0.04434246]]
[[0.09992629 0.09992629 0.09992629 0.09992629]]
[[0.0145902]
 [0.0145902]
 [0.0145902]
 [0.0145902]]
```

Same values on all groups of weights

Training the constant initialization model

```
1  # Define and train neural network structure
2  n_x = 2
3  n_h = 4
4  n_y = 1
5  np.random.seed(37)
6  shallow_neural_net = ShallowNeuralNet_ConstantInit(n_x, n_h, n_y, 0.1)
7  shallow_neural_net.train(inputs, outputs, N_max = 10000, alpha = 1e-6, delta = 1e-6, display = True)
```

```
Iteration 1 - Loss = 618.4538740049891
W1, b1, W2, b2:
[[0.09938446 0.09938446 0.09938446 0.09938446]
 [0.09676805 0.09676805 0.09676805 0.09676805]]
[[0.09999521 0.09999521 0.09999521 0.09999521]]
[[0.09614772]
 [0.09614772]
 [0.09614772]
 [0.09614772]]
[[0.09995206]]
Iteration 501 - Loss = 1.3726340052732087
W1, b1, W2, b2:
[[0.09064391 0.09064391 0.09064391 0.09064391]
 [0.04434246 0.04434246 0.04434246 0.04434246]]
[[0.09992629 0.09992629 0.09992629 0.09992629]]
[[0.0145902]
 [0.0145902]
 [0.0145902]
 [0.0145902]]
[[0.00000030]]
```

Same values on all groups of weights

Same thing for bias vector

Training the constant initialization model

```
1   # Define and train neural network structure
2   n_x = 2
3   n_h = 4
4   n_y = 1
5   np.random.seed(37)
6   shallow_neural_net = ShallowNeuralNet_ConstantInit(n_x, n_h, n_y, 0.1)
7   shallow_neural_net.train(inputs, outputs, N_max = 10000, alpha = 1e-6, delta = 1e-6, display = True)
```

```
Iteration 1 - Loss = 618.4538740049891
W1, b1, W2, b2:
[[0.09938446 0.09938446 0.09938446 0.09938446]
 [0.09676805 0.09676805 0.09676805 0.09676805]]
[[0.09999521 0.09999521 0.09999521 0.09999521]]
[[0.09614772]
 [0.09614772]
 [0.09614772]
 [0.09614772]]
[[0.09995206]]
Iteration 501 - Loss = 1.3726340052732087
W1, b1, W2, b2:
[[0.09064391 0.09064391 0.09064391 0.09064391]
 [0.04434246 0.04434246 0.04434246 0.04434246]]
[[0.09992629 0.09992629 0.09992629 0.09992629]]
[[0.0145902]
 [0.0145902]
 [0.0145902]
 [0.0145902]]
```

Same values on all groups of weights

Same thing for bias vector

Keeps happening without discontinuity no matter the number of iterations…

This means all neurons are doing the same thing! (a.k.a. **symmetry** in the neural network)

Symmetry in the Neural Network

Definition (Symmetry):

Symmetry in a neural network is the **tendency of all neurons to have the same weights and processes.**

As a **lack of diversity**, this can lead to a **lack of generalization** and can prevent the NN from learning.

Additionally, symmetrical neural networks can be vulnerable to **adversarial attacks** (on Week 6).

```
Iteration 1 - Loss = 618.4538740049891
W1, b1, W2, b2:
[[0.09938446 0.09938446 0.09938446 0.09938446]
 [0.09676805 0.09676805 0.09676805 0.09676805]]
[[0.09999521 0.09999521 0.09999521 0.09999521]]
[[0.09614772]
 [0.09614772]
 [0.09614772]
 [0.09614772]]
[[0.09995206]]
Iteration 501 - Loss = 1.3726340052732087
W1, b1, W2, b2:
[[0.09064391 0.09064391 0.09064391 0.09064391]
 [0.04434246 0.04434246 0.04434246 0.04434246]]
[[0.09992629 0.09992629 0.09992629 0.09992629]]
[[0.0145902]
 [0.0145902]
 [0.0145902]
 [0.0145902]]
```

Symmetry in the Neural Network

Definition (Symmetry):

Symmetry in a neural network is the **tendency of all neurons to have the same weights and processes.**

As a **lack of diversity**, this can lead to a **lack of generalization** and can prevent the NN from learning.

Additionally, symmetrical neural networks can be vulnerable to **adversarial attacks** (on Week 6).

- Essentially, this happened because **all weights and biases had the same starting point** (due to the constant initialization).

- The backward process then updated the parameters identically, and they end up keeping the same values over the course of training.

- **Need to enforce diversity on these parameters from the start! Hence, random initialization.**

The random normal initialization

Definition (the random normal initialization [LeCun1998]):

The **normal random initializer** will setup **parameters randomly,** by sampling from a **normal distribution**, with a **mean 0** and a **standard deviation of 1** (or lower).

This initializer is **straightforward and simple** but can lead to **slow convergence or even divergence as the network gets deeper**.

```python
28    def init_parameters_normal(self):
29        # Weights and biases matrices (randomly initialized)
30        self.W1 = np.random.randn(self.n_x, self.n_h)*0.1
31        self.b1 = np.random.randn(1, self.n_h)*0.1
32        self.W2 = np.random.randn(self.n_h, self.n_y)*0.1
33        self.b2 = np.random.randn(1, self.n_y)*0.1
```

Using a random initialization is just better

The Xavier initialization

Definition (the Xavier initialization [Glorot2010]):

The **Xavier initializer** is based on a **Gaussian distribution**, with parameters initialized to zero-mean and a **variance of** $\frac{2}{N_x+N_y}$, with N_x being the input size and N_y the output size.

It is the **most popular initializer** in deep learning and is **a good starting point for most architectures**.

```python
35        def init_parameters_xavier(self):
36            # Weights and biases matrices (Xavier initialized)
37            var = np.sqrt(2.0/(self.n_x + self.n_y))
38            self.W1 = np.random.randn(self.n_x, self.n_h)*var
39            self.b1 = np.random.randn(1, self.n_h)*var
40            self.W2 = np.random.randn(self.n_h, self.n_y)*var
41            self.b2 = np.random.randn(1, self.n_y)*var
```

The He initialization

Definition (the He initialization [LeCun1998]):

The **He initializer** is similar to the **Xavier initializer**, but with a **variance** of **variance of** $\dfrac{1}{N_{in}+N_{out}}$, where N_{in} (resp. N_{out}) is the size of each layer input (resp. output). Each layer is individually initialized.

It seems to **improve performance when working with deeper networks**.

```python
43      def init_parameters_he(self):
44          # Weights and biases matrices (He initialized)
45          range1 = np.sqrt(4/(self.n_x + self.n_h))
46          self.W1 = np.random.randn(self.n_x, self.n_h)*range1
47          self.b1 = np.random.randn(1, self.n_h)*range1
48          range2 = np.sqrt(4/(self.n_h + self.n_y))
49          self.W2 = np.random.randn(self.n_h, self.n_y)*range2
50          self.b2 = np.random.randn(1, self.n_y)*range2
```

The LeCun initialization

Definition (the LeCun initialization [LeCun1998bis]):

The **LeCun initializer** is **based on a random normal distribution**, with a variance of $\sqrt{\dfrac{1}{N_{in}}}$, with N_{in} being the number of inputs of each layer.

This initializer is particularly useful for **architectures with sigmoid and tanh activation functions**.

```
52      def init_parameters_lecun(self):
53          # Weights and biases matrices (LeCun initialized)
54          range1 = np.sqrt(1/self.n_x)
55          self.W1 = np.random.randn(self.n_x, self.n_h)*range1
56          self.b1 = np.zeros((1, self.n_h))
57          range2 = np.sqrt(1/self.n_h)
58          self.W2 = np.random.randn(self.n_h, self.n_y)*range2
59          self.b2 = np.zeros((1, self.n_y))
```

Initializations variations and more

Many more initialization formulas exist:

• Xavier has random uniform and normal variations (same for He),

• Glorot proposed more initializations,

• Orthogonal initializations are sometimes useful but rare,

• Variance scaling initialization also exists but rarely used,

• Etc.

PyTorch has listed a lot of them, ready to use:
https://pytorch.org/docs/stable/nn.init.html

No free lunch theorem

Very important note: What I listed earlier are empirical observations only, certainly not absolute rules.

- For this reason, it is a good idea to try a few different initialization techniques and see which one works best for your network and dataset.

- Some recent research has also suggested that using a combination of different initialization techniques (e.g., Xavier initialization for some layers and He initialization for others) can improve performance.

- **Also, careful:** Every now and then, we get a "smart*ss" claiming to have found a "ground-breaking new initializer beating all others".

- In reality, there is something that the machine learning community likes to call the **"No free lunch theorem".**

No free lunch theorem

Definition (the no free lunch (or NFL) theorem):

The **no free lunch (NFL) theorem** for supervised machine learning is a theorem that essentially implies that **no single machine learning algorithm or tool is universally the best-performing algorithm for all problems**.

(https://en.wikipedia.org/wiki/No_free_lunch_theorem).

FREE LUNCH!

What do you think?

Son, if a deal is too good to be true, then it usually is.

GUY DOWNES ©

officeguycartoons.com

When in doubt about a supposedly "ground-breaking new technique beating all other techniques", **try them out**.

No free lunch theorem – season 1

Definition (the no free lunch (or NFL) theorem – season 1):

The **no free lunch (NFL) theorem** also implies that there is **no such thing "one initializer that works for all models and problems".**

No choice: try some of them and see how it goes!

(A good read on this matter, also: https://www.deeplearning.ai/ai-notes/initialization/index.html)

FREE LUNCH!

What do you think?

Son, if a deal is too good to be true, then it usually is.

GUY DOWNES ©

officeguycartoons.com

When in doubt about a supposedly "ground-breaking new technique beating all other techniques", **try them out**.

A weird initializer and a new problem

Definition (our custom "weird" initializer):

Let us consider the **"weird" initializer** below. It uses a r**andom uniform distribution**, **scaling parameters based on the input and output sizes**, like in Xavier/He.

We could in fact see this as a **random uniform variation of Xavier/He**.

And yet, when we use it, something weird happens.

```python
68    def init_parameters_weird(self):
69        # Weights and biases matrices (Weird? initialized)
70        init_val = np.sqrt(6.0/(self.n_x + self.n_y))
71        self.W1 = np.random.uniform(-init_val, init_val, (self.n_x, self.n_h))
72        self.b1 = np.random.uniform(-init_val, init_val, (1, self.n_h))
73        self.W2 = np.random.uniform(-init_val, init_val, (self.n_h, self.n_y))
74        self.b2 = np.random.uniform(-init_val, init_val, (1, self.n_y))
```

```
1  # Define and train neural network structure (Weird initialization)
2  n_x = 2
3  n_h = 4
4  n_y = 1
5  init_type = "Weird"
6  np.random.seed(37)
7  shallow_neural_net_weird = ShallowNeuralNet(n_x, n_h, n_y, init_type)
8  # Train and show final loss
9  shallow_neural_net_weird.train(inputs, outputs, N_max = 100, alpha = 1e-6, delta = 1e-6, display = True)
```

```
- Gradients:
[[ 0.0347825  -0.28526197 -0.24439365  0.18661634]
 [ 0.17610429 -1.44428536 -1.23736846  0.94484113]]
[[ 0.00026903 -0.00220639 -0.00189029  0.0014434 ]]
[[-0.64676422]
 [-0.55320161]
 [ 1.42348809]
 [-0.81598256]]
[[-0.00168314]]
- Parameters:
[[ 1.25722625 -0.1015457  -0.86890687  0.23163369]
 [ 0.33964944  0.52106421 -1.12164798  0.69431032]]
[[-0.61665631  0.71679298  0.82789654  0.3603433 ]]
[[-0.15983769]
 [ 1.31087797]
 [ 1.12307381]
 [-0.85756698]]
[[0.27139079]]
Iteration 1 - Loss = 5638943.191688167
```

Wow, that's a big loss value!?

```
1  # Define and train neural network structure (Weird initialization)
2  n_x = 2
3  n_h = 4
4  n_y = 1
5  init_type = "Weird"
6  np.random.seed(37)
7  shallow_neural_net_weird = ShallowNeuralNet(n_x, n_h, n_y, init_type)
8  # Train and show final loss
9  shallow_neural_net_weird.train(inputs, outputs, N_max = 100, alpha = 1e-6, delta = 1e-6, display = True)
```

```
Iteration 26 - Loss = nan
- Gradients:
[[nan nan nan nan]
 [nan nan nan nan]]
[[nan nan nan nan]]
[[nan]
 [nan]
 [nan]
 [nan]]
[[nan]]
- Parameters:
[[nan nan nan nan]
 [nan nan nan nan]]
[[nan nan nan nan]]
[[nan]
 [nan]
 [nan]
 [nan]]
[[nan]]
```

And after iteration 26, we are getting NaNs values (Not a Number values) all over the place…?!
What is going on…?!

The exploding gradient problem

Definition (exploding gradients):

These NaN values are a typical symptom for a phenomenon called **exploding gradients**.

This typically occurs when the gradient descent rule has changes (in $\alpha \frac{\partial dL}{\partial W1}$ for instance) that are far greater than the current values in the matrices (e.g. W_1).

The values will explode to infinity and eventually become NaNs.

This typically happens with

- **unlucky initializations**,
- **when the learning rate α is too large for the given parameters**,
- and sometimes when **gradients are approximated with formulas that are a bit buggish**…

In order to avoid this problem, many techniques have been developed (something for later).

Unlucky initialization

Definition (Unlucky Initialization):

Unlucky Initialization is the term used to describe when a system's behaviour is impossible to predict due to the **randomness of its initial conditions**.

In some cases, the output of a system will be **drastically different** depending on the initial conditions, which can lead to **vastly different outcomes**.

In our previous "weird" initializer, we can resolve the NaN problem, by just using a different seed for the initialization…!

Yup, as simple as that!

(BTW, that could be a proof that the NFL theorem is indeed true!)

```
1  # Define and train neural network structure (Weird initialization)
2  n_x = 2
3  n_h = 4
4  n_y = 1
5  init_type = "Weird"
6  np.random.seed(17)
7  shallow_neural_net_weird2 = ShallowNeuralNet(n_x, n_h, n_y, init_type)
8  # Train and show final loss
9  shallow_neural_net_weird2.train(inputs, outputs, N_max = 500, alpha = 1e-6, delta = 1e-6, display = True)
```

```
- Gradients:
[[-9.86521735e-08 -7.36853902e-08 -8.11338502e-08  1.17729102e-07]
 [-3.95240319e-08 -2.95213335e-08 -3.25054864e-08  4.71670175e-08]]
[[-3.87197953e-08 -2.89206322e-08 -3.18440634e-08  4.62072611e-08]]
[[ 4.10683166e-07]
 [-1.03823853e-07]
 [ 3.91291750e-07]
 [ 7.42268171e-07]]
[[-1.96398395e-07]]
- Parameters:
[[-0.57336145  0.13297469 -1.0271704  -1.19171781]
 [ 0.65980609  0.57963967 -0.6216372   0.49082227]]
[[-1.30374448 -0.40183954  1.25928682 -1.24410091]]
[[ 0.19714925]
 [ 0.14725493]
 [ 0.16214014]
 [-0.23527311]]
[[0.14792844]]
Stopping early - loss evolution was less than beta on iteration 342.
```

Different seed…
Now that works?!

```
1  # Showing final loss
2  print(shallow_neural_net_weird2.loss)
```

```
0.12422436714449628
```

The learning rate tradeoff

Definition (Learning rate tradeoff):

The **learning rate** is a **hyperparameter** which controls the step size taken by the algorithm in order to minimize the loss function during training.

- If the **learning rate is too small**, the algorithm will **take longer to converge**, but it can also prevent overshooting the global minimum.

- If the **learning rate is too high**, the algorithm **may converge faster**, but **it may lead to exploding gradients**.

The learning rate tradeoff

Definition (Learning rate tradeoff):

There is therefore a **tradeoff** on the deep learning rate, and as in all things with tradeoffs, it is about finding the **optimal middle ground**.

High Learning rate value:

Trains faster

Risk of overshooting min

Risk of exploding gradient

Low learning value:

Trains slow, but safe

Risk of training so slow we barely make any changes

Triggering an exploding gradient on purpose

And, as a proof of that…

We can, in fact, **trigger an exploding gradient problem, on purpose, by setting a very high learning rate value!**

Let us try a much higher value for the learning rate (e.g. using $1e^{-2}$ instead of $1e^{-6}$).

```
 1  # Define and train neural network structure (random normal initialization)
 2  n_x = 2
 3  n_h = 4
 4  n_y = 1
 5  init_type = "Normal"
 6  np.random.seed(37)
 7  shallow_neural_net_normal1 = ShallowNeuralNet(n_x, n_h, n_y, init_type)
 8  # Train and show final loss
 9  shallow_neural_net_normal1.train(inputs, outputs, N_max = 1000, alpha = 1e-6, delta = 1e-6, display = True)
10  print(shallow_neural_net_normal1.loss)
```

```
- Gradients:
[[ 1.42356523e-05 -1.73577846e-06 -9.53259136e-07  1.24839272e-06]
 [ 1.32743203e-04 -1.61856153e-05 -8.88885652e-06  1.16408890e-05]]
[[ 1.27309997e-07 -1.55231348e-08 -8.52503385e-09  1.11644251e-08]]
[[ 3.86686371e-04]
 [ 2.71147622e-04]
 [ 8.05690592e-05]
 [-1.49929229e-04]]
[[2.45019007e-06]]
- Parameters:
[[-0.00544636  0.06743081  0.0346647  -0.13003462]
 [ 0.15185119  0.09898237  0.02776809 -0.04485894]]
[[ 0.09619662 -0.08275786  0.05346571  0.12283862]]
[[ 0.05195923]
 [-0.00633548]
 [-0.00347934]
 [ 0.00455655]]
[[0.14480251]]
Iteration 1 - Loss = 6.635227700991098
```

Normal initialization and low learning rate (1e-6)
This is fine...

```
 1  # Show loss after training
 2  print(shallow_neural_net_normal1.loss)
```

```
2.0907010654169245
```

```
 1  # Define and train neural network structure (random normal initialization)
 2  n_x = 2
 3  n_h = 4
 4  n_y = 1
 5  init_type = "Normal"
 6  np.random.seed(37)
 7  shallow_neural_net_normal2 = ShallowNeuralNet(n_x, n_h, n_y, init_type)
 8  # Train and show final loss
 9  shallow_neural_net_normal2.train(inputs, outputs, N_max = 1000, alpha = 1e-2, delta = 1e-6, display = True)
10  print(shallow_neural_net_normal2.loss)
```

```
- Gradients:
[[ 0.14235652 -0.01735778 -0.00953259  0.01248393]
 [ 1.32743203 -0.16185615 -0.08888857  0.11640889]]
[[ 1.27309997e-03 -1.55231348e-04 -8.52503385e-05  1.11644251e-04]]
[[ 3.86686371]
 [ 2.71147622]
 [ 0.80569059]
 [-1.49929229]]
[[0.0245019]]
- Parameters:
[[-0.00544636  0.06743081  0.0346647  -0.13003462]
 [ 0.15185119  0.09898237  0.02776809 -0.04485894]]
[[ 0.09619662 -0.08275786  0.05346571  0.12283862]]
[[ 0.05195923]
 [-0.00633548]
 [-0.00347934]
 [ 0.00455655]]
[[0.14480251]]
Iteration 1 - Loss = 4354741.971946698
```

Learning rate too high leads to exploding gradients

```
 1  # Show loss after training
 2  print(shallow_neural_net_normal2.loss)
```

```
nan
```

No free lunch theorem – season 2

Definition (the no free lunch (or NFL) theorem – season 2):

The **no free lunch (NFL) theorem** also implies that there is **no "one value that works for all models" for hyperparameters such as the learning rate.**

No choice (again!): try some values and see how it goes!

ONE DOES NOT SIMPLY

HAVE A VALUE THAT ALWAYS WORKS FOR THE LEARNING RATE

The learning rate tradeoff

Definition (Hyperparameters and Hyperparameter tuning):

As we have seen, the **learning rate** is a **hyperparameter** that needs to be carefully chosen.

Thus, an important part of deep learning is **hyperparameter tuning.**

It consists, for instance, of **finding the optimal learning rate** that will allow the algorithm to **converge quickly and reliably**.

For a given modeln there could be multiple hyperparameters.
For instance, the lambda value we used in regularization is also a hyperparameter, which requires hyperparameter tuning.

The learning rate tradeoff

Note: Sadly, that is sometimes all it takes to make the difference between a model that successfully trains and a model that does not!

Controlling gradients to prevent explosions

Adjusting and trying different values for the learning rate is often good enough to prevent exploding gradients.

There are other options, like:

- **Downscaling initial values:** Divide initial parameter values by a certain factor.

- **Gradient clipping:** Set a max value for change on a single iteration and truncate any change that goes above that (to be discussed later).

- **Time evolution on learning rate:** Make the learning rate value change over time (to be discussed later).

```
1  # Define and train neural network structure (Weird initialization)
2  n_x = 2
3  n_h = 4
4  n_y = 1
5  init_type = "Weird"
6  np.random.seed(37)
7  shallow_neural_net_weird = ShallowNeuralNet(n_x, n_h, n_y, init_type)
8  # Train and show final loss
9  shallow_neural_net_weird.train(inputs, outputs, N_max = 100, alpha = 1e-6, delta = 1e-6, display = True)
```

```
Iteration 26 - Loss = nan
- Gradients:
[[nan nan nan nan]
 [nan nan nan nan]]
[[nan nan nan nan]]
[[nan]
 [nan]
 [nan]
 [nan]]
[[nan]]
- Parameters:
[[nan nan nan nan]
 [nan nan nan nan]]
[[nan nan nan nan]]
[[nan]
 [nan]
 [nan]
 [nan]]
[[nan]]
```

Back to our first exploding gradient example

```
1   # Define and train neural network structure (Weird initialization = some sort of a Uniform Xavier)
2   n_x = 2
3   n_h = 4
4   n_y = 1
5   init_type = "Weird"
6   np.random.seed(37)
7   shallow_neural_net_xavier2 = ShallowNeuralNet(n_x, n_h, n_y, init_type)
8   # Divide initial values by 10!
9   shallow_neural_net_xavier2.W1 /= 100
10  shallow_neural_net_xavier2.b1 /= 100
11  shallow_neural_net_xavier2.W2 /= 100
12  shallow_neural_net_xavier2.b2 /= 100
13  # Train and show final loss
14  shallow_neural_net_xavier2.train(inputs, outputs, N_max = 1000, alpha = 1e-6, delta = 1e-6, display = True)
15  print(shallow_neural_net_xavier2.loss)
```

```
- Gradients:
[[ 1.08643744e-06 -8.91020716e-06 -7.63367797e-06  5.82899367e-06]
 [ 4.17432348e-06 -3.42349093e-05 -2.93302128e-05  2.23962323e-05]]
[[ 8.17657760e-09 -6.70586248e-08 -5.74514080e-08  4.38692717e-08]]
[[-1.73842740e-05]
 [-1.29545695e-05]
 [ 3.51566967e-05]
 [-1.97254965e-05]]
[[-5.1155505e-06]]
- Parameters:
[[ 0.01257226 -0.00101546 -0.00868907  0.00231634]
 [ 0.00339649  0.00521064 -0.01121648  0.0069431 ]]
[[-0.00616656  0.00716793  0.00827897  0.00360343]]
[[-0.00159838]
 [ 0.01310878]
 [ 0.01123074]
 [-0.00857567]]
[[0.00271391]]
Iteration 1 - Loss = 7.036477890046345
```

Downscaling parameters after initialization fixed the exploding issue, but it leads to slower training.

```
1   # Show loss after training
2   print(shallow_neural_net_xavier2.loss)
```

```
2.2689392753336914
```

The vanishing gradient problem

Definition (vanishing gradients):

We observed a phenomenon called **exploding gradients**. Its counterpart is called **vanishing gradients**.

This typically occurs when the gradient descent rule has changes (in $\alpha \frac{\partial dL}{\partial W1}$ for instance) that are far smaller than the current values in the matrices (e.g. W_1).

It might also happen for other reasons (to be discussed later).

- We can typically force the apparition of a vanishing gradient problem by forcing the parameters to be initialized as very small values (or even zeroes).

- We can also force the apparition of a vanishing gradient problem by using a very small learning rate α.

- We can then observe that most parameters remain zero during training as the changes are close, or even equal to, zero.

```
1  # Define and train neural network structure (random normal initialization)
2  n_x = 2
3  n_h = 4
4  n_y = 1
5  init_type = "Normal"
6  np.random.seed(37)
7  shallow_neural_net_normal = ShallowNeuralNet(n_x, n_h, n_y, init_type)
8  # Train and show final loss
9  shallow_neural_net_normal.train(inputs, outputs, N_max = 100, alpha = 1e-6, delta = 1e-6, display = True)
10 print(shallow_neural_net_normal.loss)
```

```
- Gradients:
[[ 1.42356523e-05 -1.73577846e-06 -9.53259136e-07  1.24839272e-06]
 [ 1.32743203e-04 -1.61856153e-05 -8.88885652e-06  1.16408890e-05]]
[[ 1.27309997e-07 -1.55231348e-08 -8.52503385e-09  1.11644251e-08]]
[[ 3.86686371e-04]
 [ 2.71147622e-04]
 [ 8.05690592e-05]
 [-1.49929229e-04]]
[[2.45019007e-06]]
- Parameters:
[[-0.00544636  0.06743081  0.0346647  -0.13003462]
 [ 0.15185119  0.09898237  0.02776809 -0.04485894]]
[[ 0.09619662 -0.08275786  0.05346571  0.12283862]]
[[ 0.05195923]
 [-0.00633548]
 [-0.00347934]
 [ 0.00455655]]
[[0.14480251]]
Iteration 1 - Loss = 6.635227700991098
```

This is fine.

```
1  # Show final loss
2  print(shallow_neural_net_normal.loss)
```

```
2.654115820074732
```

```
1   # Define and train neural network structure (Zero initialization)
2   n_x = 2
3   n_h = 4
4   n_y = 1
5   init_type = "Zero"
6   np.random.seed(37)
7   shallow_neural_net_zero = ShallowNeuralNet(n_x, n_h, n_y, init_type)
8   # Train and show final loss
9   shallow_neural_net_zero.train(inputs, outputs, N_max = 100, alpha = 1e-6, delta = 1e-6, display = True)
10  print(shallow_neural_net_zero.loss)
```

```
- Gradients:
[[0. 0. 0. 0.]
 [0. 0. 0. 0.]]
[[0. 0. 0. 0.]]
[[0.]
 [0.]
 [0.]
 [0.]]
[[-4.9531058e-06]]
- Parameters:
[[0. 0. 0. 0.]
 [0. 0. 0. 0.]]
[[0. 0. 0. 0.]]
[[0.]
 [0.]
 [0.]
 [0.]]
[[0.]]
Iteration 1 - Loss = 6.641131615309766
```

Initializing parameters as zeroes is seriously problematic and the model just does not train.

```
1   # Show final loss
2   print(shallow_neural_net_zero.loss)
```

```
6.638703310983079
```

```
1  # Define and train neural network structure
2  # (random normal initialization,
3  # but very low learning rate alpha)
4  n_x = 2
5  n_h = 4
6  n_y = 1
7  init_type = "Normal"
8  np.random.seed(37)
9  shallow_neural_net_normal_zerolr = ShallowNeuralNet(n_x, n_h, n_y, init_type)
10 # Train and show final loss
11 shallow_neural_net_normal_zerolr.train(inputs, outputs, N_max = 100, alpha = 5e-8, delta = 1e-6, display = True)
12 print(shallow_neural_net_normal_zerolr.loss)
```

```
- Gradients:
[[ 7.11782617e-07 -8.67889228e-08 -4.76629568e-08  6.24196361e-08]
 [ 6.63716016e-06 -8.09280764e-07 -4.44442826e-07  5.82044450e-07]]
[[ 6.36549987e-09 -7.76156741e-10 -4.26251693e-10  5.58221254e-10]]
[[ 1.93343185e-05]
 [ 1.35573811e-05]
 [ 4.02845296e-06]
 [-7.49646145e-06]]
[[1.22509503e-07]]
- Parameters:
[[-0.00544636  0.06743081  0.0346647  -0.13003462]
 [ 0.15185119  0.09898237  0.02776809 -0.04485894]]
[[ 0.09619662 -0.08275786  0.05346571  0.12283862]]
[[ 0.05195923]
 [-0.00633548]
 [-0.00347934]
 [ 0.00455655]]
[[0.14480251]]
Iteration 1 - Loss = 6.887598572349596
```

Setting up a learning rate that is too small might seriously hinder the training.
And sometimes even lead to a "convergence" to wrong minimum.

```
1  # Show final loss
2  print(shallow_neural_net_normal_zerolr.loss)
```

```
5.748091602996055
```

Remember, the learning rate tradeoff

**Definition (Hyperparameters and hyperparameter tuning):**

As we have seen, the **learning rate** is a **hyperparameter** that needs to be carefully chosen.

Thus, an important part of deep learning is **hyperparameter tuning.**

It consists, for instance, of **finding the optimal learning rate** that will allow the algorithm to **converge quickly and reliably**.

New dataset! (An abstract one this time)

New dataset! (An abstract one this time)

- Each sample consists of two randomly drawn input values (x_1, x_2), with
$$-1 < x_1, x_2 < 1.$$

- Black line serves as (non-linear) frontier between both classes, with equation.
$$f(x) = -\frac{1}{4} + \frac{3}{4}\sin(x\pi)$$

- Two classes (green and red):

Green (class 1) if $x_2 > f(x_1)$ and red (class 0) otherwise.

New dataset!

- We can generate some entries using the helper functions here.

- Basically, drawing pairs of values randomly and checking if above or below the sine line.

```python
# ALL helper functions
eps = 1e-5
min_val = -1 + eps
max_val = 1 - eps
def val(min_val, max_val):
    return round(np.random.uniform(min_val, max_val), 2)
def class_for_val(val1, val2):
    k = np.pi
    return int(val2 >= -1/4 + 3/4*np.sin(val1*k))
n_points = 100
def create_dataset(n_points, min_val, max_val):
    val1_list = np.array([val(min_val, max_val) for _ in range(n_points)])
    val2_list = np.array([val(min_val, max_val) for _ in range(n_points)])
    inputs = np.array([[v1, v2] for v1, v2 in zip(val1_list, val2_list)])
    outputs = np.array([class_for_val(v1, v2) for v1, v2 in zip(val1_list, val2_list)]).reshape(n_points, 1)
    return val1_list, val2_list, inputs, outputs
```

```python
# Generate dataset
np.random.seed(47)
val1_list, val2_list, inputs, outputs = create_dataset(n_points, min_val, max_val)
# Check a few entries of the dataset
print(val1_list.shape)
print(val2_list.shape)
print(inputs.shape)
print(outputs.shape)
print(inputs[0:5, :])
print(outputs[0:5])
```

```
(100,)
(100,)
(100, 2)
(100, 1)
[[-0.77  0.1 ]
 [ 0.95  0.6 ]
 [ 0.46 -0.37]
 [-0.3  -0.92]
 [ 0.42  0.62]]
[[1]
 [1]
 [0]
 [0]
 [1]]
```

Adjusting our shallow NN

We can adjust our shallow neural network to the task by:

- Using the same two layers and forward pass procedure as before, leaving our init and forward methods untouched.

- Replace MSE with our cross-entropy (another name for the log-likelihood loss).

- Backward and trainer methods are left untouched.

```python
class ShallowNeuralNet():

    def __init__(self, n_x, n_h, n_y):
        # Network dimensions
        self.n_x = n_x
        self.n_h = n_h
        self.n_y = n_y
        # Initialize parameters
        self.init_parameters_normal()
        # Loss, initialized as infinity before first calculation is made
        self.loss = float("Inf")

    def init_parameters_normal(self):
        # Weights and biases matrices (randomly initialized)
        self.W1 = np.random.randn(self.n_x, self.n_h)*0.1
        self.b1 = np.random.randn(1, self.n_h)*0.1
        self.W2 = np.random.randn(self.n_h, self.n_y)*0.1
        self.b2 = np.random.randn(1, self.n_y)*0.1

    def forward(self, inputs):
        # Wx + b operation for the first layer
        Z1 = np.matmul(inputs, self.W1)
        Z1_b = Z1 + self.b1
        # Wx + b operation for the second layer
        Z2 = np.matmul(Z1_b, self.W2)
        Z2_b = Z2 + self.b2
        #  Adding clipping to keep prediction values
        # between 0 and 1 (see loss function!)
        pred = np.clip(Z2_b, 0, 1)
        return pred

    def CE_loss(self, inputs, outputs):
        # MSE loss function as before
        outputs_re = outputs.reshape(-1, 1)
        pred = self.forward(inputs)
        eps = 1e-5
        losses = outputs*np.log(pred + eps) + (1 - outputs)*np.log(1 - pred + eps)
        self.loss = -np.sum(losses)/outputs.shape[0]
        return self.loss

    def backward(self, inputs, outputs, alpha = 1e-5):
        # Get the number of samples in dataset
```
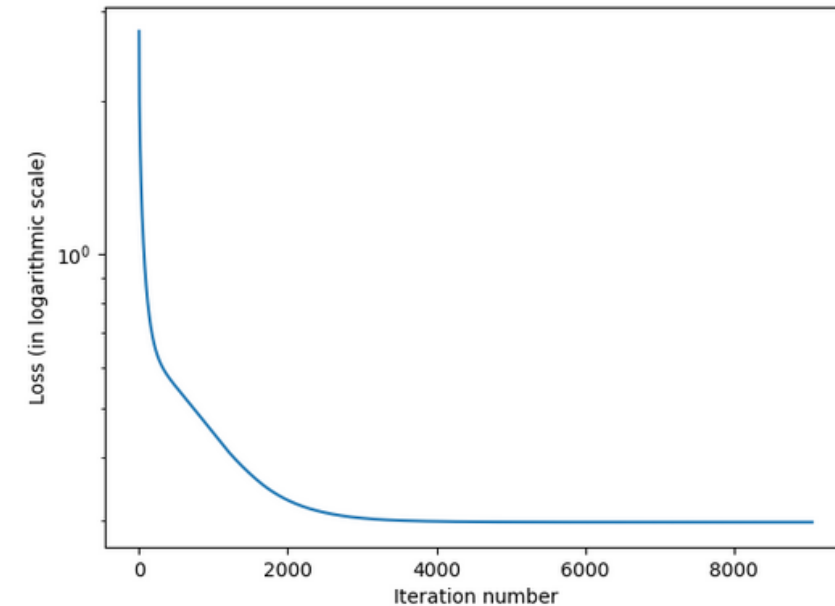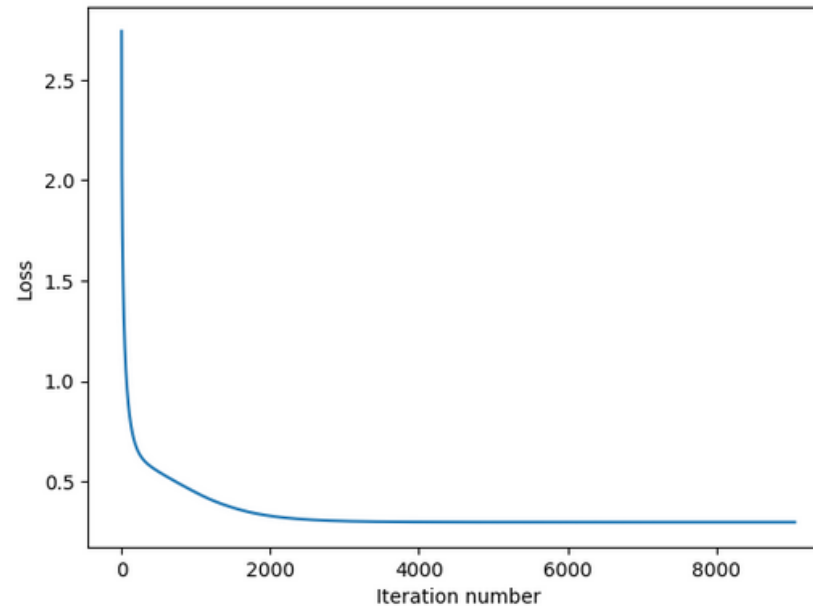
Trains just fine! (or does it?)

```
1  # Define and train neural network structure (no activation)
2  n_x = 2
3  n_h = 10
4  n_y = 1
5  np.random.seed(37)
6  shallow_neural_net = ShallowNeuralNet(n_x, n_h, n_y)
7  # Train and show final loss
8  shallow_neural_net.train(inputs, outputs, N_max = 10000, alpha = 5e-3, delta = 1e-8, display = False)
9  print(shallow_neural_net.loss)
```

0.29783823594717523

```
1  shallow_neural_net.show_losses_over_training()
```

Trains just fine! (or does it?)

NOPE

Linear layers necessarily lead to linear boundary

As shown in Notebook,

- Our model consists of two linear layers, with two WX+B operations in a row.

- The final result is necessarily linear also, as shown in the final y_{pred} equation.

- Boundary is therefore linear!

To demonstrate, let us consider that $n_x = 2$, $n_h = 4$ and $n_y = 1$. In this configuration, we have:

$$W_1 = \begin{pmatrix} w_{1,1}^{(1)}, & w_{1,2}^{(1)}, & w_{1,3}^{(1)}, & w_{1,4}^{(1)} \\ w_{2,1}^{(1)}, & w_{2,2}^{(1)}, & w_{2,3}^{(1)}, & w_{2,4}^{(1)} \end{pmatrix}$$

$$b_1 = \begin{pmatrix} b_{1,1}^{(1)}, & b_{1,2}^{(1)}, & b_{1,3}^{(1)}, & b_{1,4}^{(1)} \end{pmatrix}$$

$$W_2 = \begin{pmatrix} w_{1,1}^{(2)} \\ w_{1,2}^{(2)} \\ w_{1,3}^{(2)} \\ w_{1,4}^{(2)} \end{pmatrix}$$

$$b_2 = \begin{pmatrix} b_{1,1}^{(2)} \end{pmatrix}$$

After the first operation $Z_1 = XW_1 + b_1$, we have:

$$Z_1 = \begin{pmatrix} w_{1,1}^{(1)}x_1 + w_{2,1}^{(1)}x_2 + b_{1,1}^{(1)}, & w_{1,2}^{(1)}x_1 + w_{2,2}^{(1)}x_2 + b_{1,2}^{(1)}, & w_{1,3}^{(1)}x_1 + w_{2,3}^{(1)}x_2 + b_{1,3}^{(1)}, & w_{1,4}^{(1)}x_1 + w_{2,4}^{(1)}x_2 + b_{1,4}^{(1)} \end{pmatrix}$$

After the second operation, the predicted price y_{pred} is given by $y_{pred} = Z_1 W_2 + b_2$

$$y_{pred} = \left(\sum_{k=1}^{4} w_{1,k}^{(1)}w_{1,k}^{(2)} \right) x_1 + \left(\sum_{k=1}^{4} w_{2,k}^{(1)}w_{1,k}^{(2)} \right) x_2 + \left(\sum_{k=1}^{4} b_{1,k}^{(1)}w_{1,k}^{(2)} \right) + b_{1,1}^{(2)}$$

Trains just fine! (or does it?)

Well, that is going to be a problem, because the black frontier is definitely not linear…

Or, in other words, the data is **not linearly separable.**

Activation functions and introducing non-linearity in the Neural Networks

Definition (Activation functions and non-linearity):

Activation functions are an important component of neural networks because they **introduce non-linearity** to the model.

This is important because most real-world data is non-linear.

Simply done by **adding an extra operation** after the linear one, for instance **our sigmoid function**.

Layer input
X

Linear operation
Z = WX+B

Activation function
A = s(Z)

Layer output
A

Our new layer prototype, to be used in place of the Linear operation only

Adding sigmoid activations

We can update our model by:

- Adding a sigmoid method, implementing the sigmoid operation from earlier.

- Adding sigmoid operations after each linear operation in the forward method.

- Our cross-entropy loss method remains unchanged, as it will simply use the updated forward procedure.

```python
class ShallowNeuralNet_WithAct():

    def __init__(self, n_x, n_h, n_y):
        # Network dimensions
        self.n_x = n_x
        self.n_h = n_h
        self.n_y = n_y
        # Initialize parameters
        self.init_parameters_normal()
        # Loss, initialized as infinity before first calculation is made
        self.loss = float("Inf")

    def init_parameters_normal(self):
        # Weights and biases matrices (randomly initialized)
        self.W1 = np.random.randn(self.n_x, self.n_h)*0.1
        self.b1 = np.random.randn(1, self.n_h)*0.1
        self.W2 = np.random.randn(self.n_h, self.n_y)*0.1
        self.b2 = np.random.randn(1, self.n_y)*0.1

    def sigmoid(self, val):
        return 1/(1 + np.exp(-val))

    def forward(self, inputs):
        # Wx + b operation for the first layer
        Z1 = np.matmul(inputs, self.W1)
        Z1_b = Z1 + self.b1
        A1 = self.sigmoid(Z1_b)
        # Wx + b operation for the second layer
        Z2 = np.matmul(A1, self.W2)
        Z2_b = Z2 + self.b2
        y_pred = self.sigmoid(Z2_b)
        return y_pred

    def CF_loss(self, inputs, outputs):
```

Adding sigmoid activations

Careful now…

- **As the forward procedure has changed, we have to update the backward propagation.**

- Going back to the chain rule procedure and accounting for the new sigmoid operations gives us the new update rules to use in the backward method…

```python
class ShallowNeuralNet_WithAct():

    def __init__(self, n_x, n_h, n_y):
        # Network dimensions
        self.n_x = n_x
        self.n_h = n_h
        self.n_y = n_y
        # Initialize parameters
        self.init_parameters_normal()
        # Loss, initialized as infinity before first calculation is made
        self.loss = float("Inf")

    def init_parameters_normal(self):
        # Weights and biases matrices (randomly initialized)
        self.W1 = np.random.randn(self.n_x, self.n_h)*0.1
        self.b1 = np.random.randn(1, self.n_h)*0.1
        self.W2 = np.random.randn(self.n_h, self.n_y)*0.1
        self.b2 = np.random.randn(1, self.n_y)*0.1

    def sigmoid(self, val):
        return 1/(1 + np.exp(-val))

    def forward(self, inputs):
        # Wx + b operation for the first layer
        Z1 = np.matmul(inputs, self.W1)
        Z1_b = Z1 + self.b1
        A1 = self.sigmoid(Z1_b)
        # Wx + b operation for the second layer
        Z2 = np.matmul(A1, self.W2)
        Z2_b = Z2 + self.b2
        y_pred = self.sigmoid(Z2_b)
        return y_pred
```

Step 1: Drawing a computation graph

Inputs X

Linear #1
$Z_1 = W_1 X + b_1$

Sigmoid
$A_1 = s(Z_1)$

Linear #2
$Z_2 = W_2 X + b_2$

Sigmoid
$A_2 = s(Z_2)$

Cross Entropy Loss
$L(A_2, Y)$

Step 2: Forward equations

Our forward method gives:

$$Z_1 = W_1 X + b_1$$

$$A_1 = s(Z_1)$$

$$Z_2 = W_2 A_1 + b_2$$

$$A_2 = s(Z_2)$$

$$L = \frac{-1}{N} \sum_i^N Y\, ln(A_2) + (1 - Y)ln(1 - A_2)$$

Step 3: Use the chain rule to backpropagate

Retrieving the gradient descent update rules takes a few steps and requires some organizing…

- First, recall that our loss function is defined as

$$L = \frac{-1}{N} \sum_i^N Y\,ln(A_2) + (1 - Y)ln(1 - A_2)$$

Therefore, we have

$$\frac{\partial L}{\partial A_2} = -\frac{Y}{A_2} + \frac{1 - Y}{1 - A_2}$$

Step 3: Use the chain rule to backpropagate

- Then, recall that

$$A_2 = s(Z_2)$$

We can then prove that

$$s'(X) = s(X)\big(1 - s(X)\big)$$

Using the chain rule, we then obtain

$$\frac{\partial L}{\partial Z_2} = \frac{\partial L}{\partial A_2}\frac{\partial A_2}{\partial Z_2} = \frac{\partial L}{\partial A_2}A_2(1 - A_2)$$

Step 3: Use the chain rule to backpropagate

- Let us continue going backwards

$$Z_2 = W_2 A_1 + b_2$$

Using the chain rule, we then obtain

$$\frac{\partial L}{\partial W_2} = \frac{\partial L}{\partial Z_2}\frac{\partial Z_2}{\partial W_2} = \frac{\partial L}{\partial Z_2}A_1$$

$$\frac{\partial L}{\partial b_2} = \frac{\partial L}{\partial Z_2}\frac{\partial Z_2}{\partial b_2} = \frac{\partial L}{\partial Z_2}$$

Step 3: Use the chain rule to backpropagate

- Let us continue going backwards

$$Z_2 \;=\; W_2 A_1 \;+\; b_2$$
$$A_1 \;=\; s(Z_1)$$

Using the chain rule and our sigmoid derivative, we then obtain

$$\frac{\partial L}{\partial A_1} = \frac{\partial L}{\partial Z_2}\frac{\partial Z_2}{\partial A_1} = \frac{\partial L}{\partial Z_2} W_2$$

$$\frac{\partial L}{\partial Z_1} = \frac{\partial L}{\partial A_1}\frac{\partial A_1}{\partial Z_1} = \frac{\partial L}{\partial A_1} A_1 (1 - A_1)$$

Step 3: Use the chain rule to backpropagate

- Let us continue going backwards

$$Z_2 = W_2 A_1 + b_2$$
$$A_1 = s(Z_1)$$

Using the chain rule and our sigmoid derivative, we then obtain

$$\frac{\partial L}{\partial A_1} = \frac{\partial L}{\partial Z_2}\frac{\partial Z_2}{\partial A_1} = \frac{\partial L}{\partial Z_2}W_2$$

$$\frac{\partial L}{\partial Z_1} = \frac{\partial L}{\partial A_1}\frac{\partial A_1}{\partial Z_1} = \frac{\partial L}{\partial A_1}A_1(1 - A_1)$$

Step 3: Use the chain rule to backpropagate

- Let us continue going backwards

$$Z_1 \;=\; W_1 X \;+\; b_1$$

Using the chain rule and our sigmoid derivative, we then obtain

$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial Z_1}\frac{\partial Z_1}{\partial W_1} = \frac{\partial L}{\partial Z_1} X$$

$$\frac{\partial L}{\partial b_1} = \frac{\partial L}{\partial Z_1}\frac{\partial Z_1}{\partial b_1} = \frac{\partial L}{\partial Z_1}$$

Backward Method Update

Reusing our formulas carefully gives the following update for the backward method.

- Notice how we do not rush and compute each term, one operation at a time, to avoid making mistakes.

- Take it slow!

```python
def backward(self, inputs, outputs, alpha = 1e-5):
    # Get the number of samples in dataset
    m = inputs.shape[0]

    # Forward propagate
    Z1 = np.matmul(inputs, self.W1)
    Z1_b = Z1 + self.b1
    A1 = self.sigmoid(Z1_b)
    Z2 = np.matmul(A1, self.W2)
    Z2_b = Z2 + self.b2
    A2 = self.sigmoid(Z2_b)

    # Compute error term
    dL_dA2 = -outputs/A2 + (1 - outputs)/(1 - A2)
    dL_dZ2 = dL_dA2*A2*(1 - A2)
    dL_dA1 = np.dot(dL_dZ2, self.W2.T)
    dL_dZ1 = dL_dA1*A1*(1 - A1)

    # Gradient descent update rules
    self.W2 -= (1/m)*alpha*np.dot(A1.T, dL_dZ2)
    self.W1 -= (1/m)*alpha*np.dot(inputs.T, dL_dZ1)
    self.b2 -= (1/m)*alpha*np.sum(dL_dZ2, axis = 0, keepdims = True)
    self.b1 -= (1/m)*alpha*np.sum(dL_dZ1, axis = 0, keepdims = True)

    # Update Loss
    self.CE_loss(inputs, outputs)
```

Trains just fine!

```
1  # Define and train neural network structure (with activation)
2  n_x = 2
3  n_h = 10
4  n_y = 1
5  np.random.seed(37)
6  shallow_neural_net_act = ShallowNeuralNet_WithAct(n_x, n_h, n_y)
7  # Train and show final loss
8  shallow_neural_net_act.train(inputs, outputs, N_max = 100000, alpha = 5, delta = 1e-8, display = False)
9  print(shallow_neural_net_act.loss)
```

0.00022943092333062472

```
1  shallow_neural_net_act.show_losses_over_training()
```

Trains just fine!

Using the sigmoid activation functions helped the network create some non-linearity! (Looking good!)
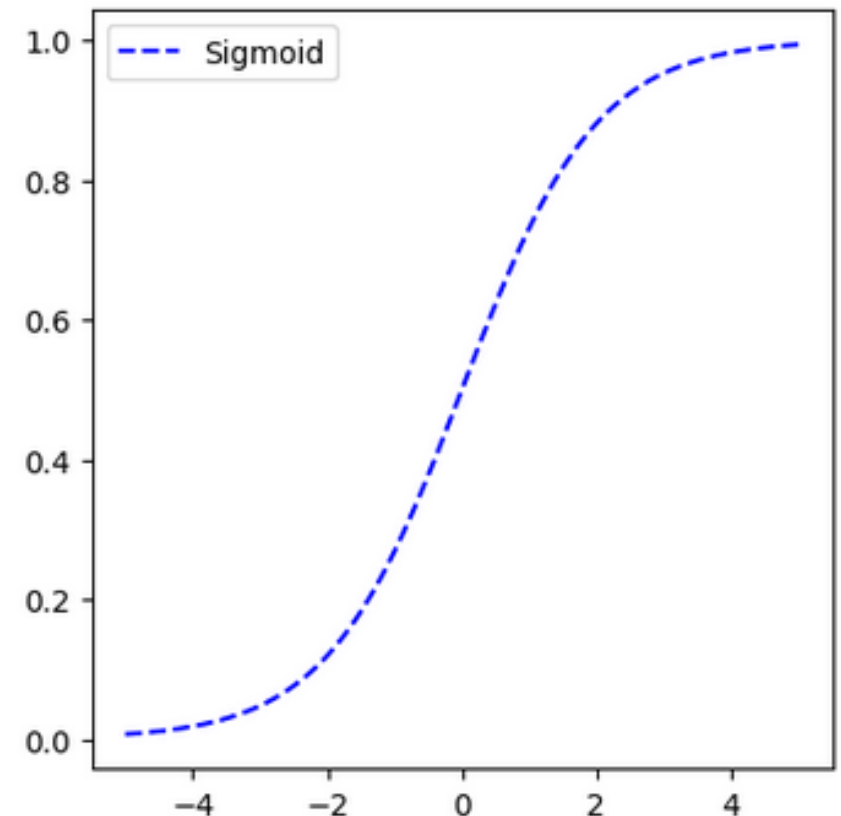
Examples of activations: sigmoid

Definition (sigmoid function):

$$s(x) = \frac{1}{1 + \exp(-x)}$$

The **sigmoid activation function** is often used in the **output layer of a binary classification model**, because the output will then have values between 0 and 1, and can therefore be interpreted as a **probability**.

```python
def sigmoid(val):
    return 1/(1 + np.exp(-val))
```

Examples of activations: Hyperbolic tangent

Definition (tanh function):

The **tanh (short for "hyperbolic tangent") activation function** is similar to the sigmoid function, but it maps values to a range between -1 and 1.

Like the sigmoid function, it is often used in the output layer of a classification model.

```python
1  def tanh(val):
2      return np.tanh(val)
```

Examples of activations: ReLU

Definition (ReLU function):

The **ReLU** (short for **"rectified linear unit"**) **activation function** is a simple function that maps any input value less than 0 to 0, and any input value greater than or equal to 0 to the input value itself.

It is widely used because it is computationally efficient and does not saturate (i.e., "die") like some other activation functions.

```python
def ReLU(val):
    return np.maximum(0, val)
```

Examples of activations: **Leaky ReLU**

Definition (Leaky ReLU function):

The **leaky ReLU activation** function is similar to the ReLU function, but it allows a **small negative slope** for input values less than 0.

This can help to alleviate the "dying ReLU" problem, where some neurons in the network "die" and no longer respond to input.

```
1  def leaky_ReLU(val, alpha = 0.1):
2      return np.where(val > 0, val, alpha*val)
```

Examples of activations: ELU

```python
def ELU(val, alpha = 0.1):
    return np.where(val > 0, val, alpha*(np.exp(val) - 1))
```

Definition (ELU function):

The **ELU** (short for "**exponential linear unit**") **activation function** is similar to the ReLU function, but it has a negative slope for input values less than 0.

Like the leaky ReLU function, it can help to alleviate the "dying ReLU" problem.

A note on activation functions

We have listed the most common activation functions, which are used 99% of the time.

Keep in mind that **many more activation functions exist**, some of them being very niche, but worth keeping an eye on. Have a look at the list here:
https://pytorch.org/docs/stable/nn.html#non-linear-activations-weighted-sum-nonlinearity

The universal approximation theorem

Definition (the Universal Approximation Theorem):

Mathematically speaking, any neural network architecture aims at finding the mathematical function $y = f(x)$ that can map some given inputs x to outputs y. And the function $f(x)$ can be arbitrarily complex, based on the dataset and task at hand.

The Universal Approximation Theorem tells us that **Neural Networks have a kind of universality property.**

This means that no matter what $f(x)$ might be, there is a network that can approximately approach the function $f(x)$!

As such, this is the most important theorem of Deep Learning!

Learn more about these topics

Out of class, supporting papers, for those of you who are curious.

- [LeCun1998] Y. **LeCun**, Y. **Bengio**, and G. **Hinton**, "Efficient backprop. In Neural networks: Tricks of the trade", 1998.

- [Glorot2010] X. **Glorot**, and Y. **Bengio**, "Understanding the difficulty of training deep feedforward neural networks", 2010.

- [He2015] K. **He**, X. Zhang, S. Ren, & J. Sun, "Delving deep into rectifiers: Surpassing human-level performance on imagenet classification.", 2015.

- [LeCun1998bis] Y. **LeCun**, L. **Bottou**, Y. **Bengio**, and P. Haffner "Gradient-based learning applied to document recognition.", 1998

Learn more about these topics

Tracking important names (Track their works and follow them on Scholar, Twitter, or whatever works for you!)

- **Yann LeCun: Chief AI Scientist** at **Facebook/Meta** and **Silver Professor** at the **New York University**, <u>**another one of the three Godfathers of Deep Learning**</u> and **2018 Turing Award** winner (highest distinction in Computer Science).
http://yann.lecun.com/
https://scholar.google.com/citations?user=WLN3QrAAAAAJ&hl=fr

- **Yoshua Bengio: Professor** at **University of Montreal**, <u>**last of the three Godfathers of Deep Learning**</u> and **2018 Turing Award** winner (highest distinction in Computer Science).
https://yoshuabengio.org
https://scholar.google.com/citations?user=kukA0LcAAAAJ&hl=fr

Learn more about these topics

Tracking important names (Track their works and follow them on Scholar, Twitter, or whatever works for you!)

- **Xavier Glorot:** Currently works at **DeepMind**.
  https://scholar.google.com/citations?user=_WnkXlkAAAAJ&hl=fr

- **Kaiming He: Researcher** at **Facebook AI Research (FAIR/Meta AI).**
  https://scholar.google.com/citations?user=DhtAFkwAAAAJ&hl=en

Learn more about these topics

Tracking important names (Track their works and follow them on Scholar, Twitter, or whatever works for you!)

- **Leon Bottou:** Works at **FAIR**/**Meta AI**.
https://leon.bottou.org/
https://scholar.google.fr/citations?user=kbN88gsAAAAJ&hl=fr

- **FAIR**/**Meta AI**: Facebook Artificial Intelligence Research, now Meta Artificial Intelligence is an academic research laboratory focused on generating knowledge for the AI community.
https://ai.facebook.com/
https://twitter.com/MetaAI?ref_src=twsrc%5Egoogle%7Ctwcamp%5Eserp%7Ctwgr%5Eauthor