50.039 Theory and Practice of Deep Learning
W2-S2 Neural Networks, Initializers, Optimizers and other Good Practices

Matthieu De Mari

SINGAPORE UNIVERSITY OF
TECHNOLOGY AND DESIGN

About this week (Week 2)

1. What are the **typical initializers for trainable parameters** in Neural Networks?

2. What is **symmetry** in a Neural Network and why is it essential to **break it** with proper initializers?

3. What is the **exploding gradient** problem?

4. How to **spot and fix an exploding gradient** problem?

5. What is the **vanishing gradient** problem?

6. How to **spot and fix a vanishing gradient** problem?

7. Why are **activation functions** needed in Neural Networks?

8. What is the **universal approximation theorem**?

About this week (Week 2)

9. What are typical **examples** of **activation functions**? Which activations should I be using?

10. What are typical examples of **advanced optimizers**? (Adagrad, RMSProp, Adam, etc.).

11. How to **implement** said **optimizers** manually?

12. What is a **stochastic gradient descent** and what are its **benefits**?

13. How to implement the **stochastic gradient descent** manually?

14. What is a **mini-batch gradient descent** and what are its **benefits**?

15. How to implement the **mini-batch gradient descent** manually?

16. How to **combine all optimizers concepts** into a great optimizer?

About this week (Week 2)

17. What is the **no free lunch theorem**?

18. What is the **train-test-validation split**?
Why is it **good practice** to use an extra validation set?

19. What is the **early stopping** of optimizer concept?
Why is it **good practice** to use it?

20. What are **saver** and **loader** functions?
Why is it **good practice** to use them?

21. What are **other common good practices** when it comes to Neural Networks?

The problem with gradient descent algorithm

Let us consider the following function:

$$f(x) = \frac{319}{8400} x^4 + \frac{43}{4200} x^3 - \frac{6799}{8400} x^2 - \frac{299}{840} x + 6$$

And let us assume we are trying to solve the optimization problem:

$$x^* = \arg\min_x f(x)$$

The problem with gradient descent algorithm

The derivative of f is:
$$f'(x) = \frac{1276}{8400}\, x^3 + \frac{129}{4200}\, x^2 - \frac{13598}{8400}\, x - \frac{299}{840}$$

And we could use gradient descent to solve this optimization by using the following update rule:

$$x \leftarrow x - \alpha f'(x)$$

The problem with gradient descent algorithm

Our vanilla gradient descent algorithm is implemented by defining the objective function f and its derivative f'.

We can then implement our vanilla gradient descent algorithm for this optimization problem as shown below.

```python
def obj_fun(val):
    a4 = 319/8400
    a3 = 43/4200
    a2 = -6799/8400
    a1 = -299/840
    a0 = 6
    return a4*val**4 + a3*val**3 + a2*val**2 + a1*val + a0
```

```python
def obj_fun_deriv(val):
    a4 = 319/8400
    a3 = 43/4200
    a2 = -6799/8400
    a1 = -299/840
    return 4*a4*val**3 + 3*a3*val**2 + 2*a2*val + a1
```

```python
def vanilla_gd(start_val, alpha = 0.1, n_iter = 100):
    val = start_val
    list_vals = [val]
    for iter_num in range(n_iter):
        # Compute gradient and update value
        val -= alpha*obj_fun_deriv(val)
        list_vals.append(val)
    return val, list_vals
```
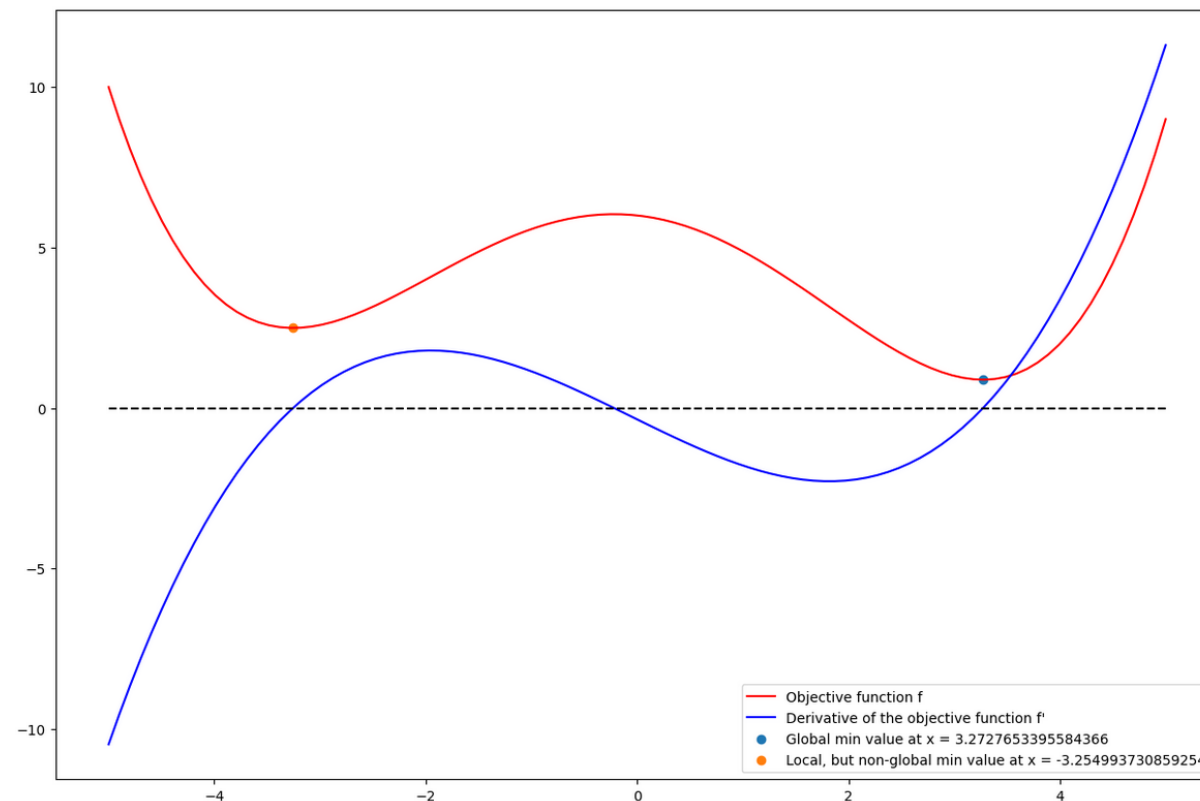
The problem with gradient descent algorithm

The function f admits a global minimum at $x \approx 3.27$ and a local (non-global) minimum at $x \approx -3.25$, as shown in the figure below.

Objective function f
Derivative of the objective function f'
Global min value at x = 3.2727653395584366
Local, but non-global min value at x = -3.254993730859254

The problem with gradient descent algorithm

When using $x_0 = 5$ as a starting point for our gradient descent, we do converge to the global minimum.

```python
opt_val_vanilla_gd1, val_list_vanilla_gd1 = vanilla_gd(start_val = 5, \
                                                        alpha = 0.1, \
                                                        n_iter = 100)
print("Optimal, found by vanilla gd: ", opt_val_vanilla_gd1)
print("Global min: ", approx_min_x)
print("Local, non-global min: ", approx_min_x2)
```

```
Optimal, found by vanilla gd:  3.2727653395584366
Global min:  3.2727653395584366
Local, non-global min:  -3.254993730859254
```

The problem with gradient descent algorithm

When using $x_0 = -5$ as a starting point for our gradient descent, however, we end up converging to the local (non-global) minimum.

Unfortunately, in the presence of local non-global minima, the GD algorithm might end up getting trapped and converging to the wrong value.

```
1  opt_val_vanilla_gd2, val_list_vanilla_gd2 = vanilla_gd(start_val = -5, \
2                                              alpha = 0.1, \
3                                              n_iter = 100)
4  print("Optimal, found by vanilla gd: ", opt_val_vanilla_gd2)
5  print("Global min: ", approx_min_x)
6  print("Local, non-global min: ", approx_min_x2)
```

```
Optimal, found by vanilla gd:  -3.254993730859254
Global min:  3.2727653395584366
Local, non-global min:  -3.254993730859254
```

The problem with gradient descent algorithm

When using $x_0 = -5$ as a starting point for our gradient descent, however, we end up converging to the local (non-global) minimum.

Unfortunately, in the presence of local non-global minima, the GD algorithm might end up getting trapped and converging to the wrong value.

In the case of our Neural Networks:

- this could lead to **suboptimal parameters W and b**, which are not going to minimize the loss functions used for training,

- and therefore lead to **a Neural Network that has not been optimally trained for its task**.

The problem with gradient descent algorithm

- **When the function is convex:** gradient descent is guaranteed to converge to the global minimum, no matter the starting point.

- **When the function is not convex (and our function f is not):** the starting point has an impact and might lead to a completely different outcome.

- Some solutions exist, but are not guaranteed to work.

Adding momentum to gradient descent

Definition (Gradient Descent with Added Momentum):

In an attempt to alleviate this phenomenon, we propose to use **momentum**.

It is a method that helps accelerate the training process by updating the weights of the model by reusing values from the past gradients.

The idea is that, if a direction in the parameters space consistently leads to good reduction in the loss function, then that direction should be reinforced.

Momentum can be implemented by adding a term to the parameter update, for instance:

$$x \leftarrow x - \alpha\, f'(x) + \mu\, f'_{prev}$$

As before,

- α is a hyperparameter that controls the step size/learning rate at each iteration,

- And the coefficient μ is another hyperparameter that controls the influence of the momentum, or previous gradient update f'_{prev}.

Adding momentum to gradient descent

Using momentum can help the model converge faster and can also **help escape local minima** and converge to global minima instead.

However, it can also cause the model to overshoot the global minimum and diverge, so – just like with the learning rate – **it is important to tune the momentum hyperparameter μ carefully (NFL theorem!)**.

Insufficient Momentum

Sufficient Momentum

Adding momentum to gradient descent

Definition (Gradient Descent with Added Momentum):

Momentum is a method that helps accelerate the training process by updating the weights of the model by reusing values from the past gradients.

Momentum can be implemented by adding a term to the parameter update, for instance:

$$x \leftarrow x - \alpha f'(x) + \mu f'_{prev}$$

```python
1  def momentum_gd(start_val, alpha = 0.1, mu = 0.1, n_iter = 100):
2      val = start_val
3      mom = 0
4      list_vals = [val]
5      for iter_num in range(n_iter):
6          # Compute gradient and momentum
7          grad = -alpha*obj_fun_deriv(val)
8          mom = mu*mom + grad
9          # Update value
10         val += mom
11         list_vals.append(val)
12     return val, list_vals
```

Adding momentum to gradient descent

As before, if the starting point is good and not too much momentum is used, we will converge to the correct minimum.

```python
opt_val_momentum_gd1, val_list_momentum_gd1 = momentum_gd(start_val = 5, \
                                                          alpha = 0.1, \
                                                          mu = 0.1, \
                                                          n_iter = 100)
print("Optimal, found by momentum gd: ", opt_val_momentum_gd1)
print("Global min: ", approx_min_x)
print("Local, non-global min: ", approx_min_x2)
```

```
Optimal, found by momentum gd:  3.272765339558436
Global min:  3.27276533955844366
Local, non-global min:  -3.254993730859254
```

Adding momentum to gradient descent

As before, if the starting point is bad and not too much momentum is used, we will end up converging to the incorrect minimum, like before.

```python
opt_val_momentum_gd2, val_list_momentum_gd2 = momentum_gd(start_val = -5, \
                                                          alpha = 0.1, \
                                                          mu = 0.1, \
                                                          n_iter = 100)
print("Optimal, found by momentum gd: ", opt_val_momentum_gd2)
print("Global min: ", approx_min_x)
print("Local, non-global min: ", approx_min_x2)
```

```
Optimal, found by momentum gd:  -3.254993730859254
Global min:  3.2727653395584366
Local, non-global min:  -3.254993730859254
```

Adding momentum to gradient descent

However, if the starting point is bad and enough momentum is used, we might be able to escape the local minimum and end up converging to the correct minimum.

```python
opt_val_momentum_gd3, val_list_momentum_gd3 = momentum_gd(start_val = -5, \
                                                          alpha = 0.1, \
                                                          mu = 0.9, \
                                                          n_iter = 100)
print("Optimal, found by momentum gd: ", opt_val_momentum_gd3)
print("Global min: ", approx_min_x)
print("Local, non-global min: ", approx_min_x2)
```

```
Optimal, found by momentum gd:  3.267681061691328
Global min:  3.27276533955584366
Local, non-global min:  -3.254993730859254
```

Adding momentum to gradient descent

A different momentum GD

Definition (Gradient Descent with Nesterov Momentum):

Nesterov **accelerated gradient descent (NAG)** is a variant of momentum gradient descent that can help accelerate convergence and improve the optimization of deep learning models.

It does this by incorporating the concept of momentum, although **in a slightly different way** than before.

This could help the optimization algorithm to continue moving in the same direction even if the gradients change, but with **less oscillations than observed with the previous momentum.**

In my opinion, this is neither better or worse, than the previous formula, just different (have I mentioned the NFL theorem yet?).

A different momentum GD

Definition (Gradient Descent with Nesterov Momentum):

Nesterov accelerated gradient descent (NAG) is a variant of momentum gradient descent that can help accelerate convergence and improve the optimization of deep learning models.

```python
def nesterov_gd(start_val, alpha = 0.1, mu = 0.1, n_iter = 100):
    val = start_val
    mom = 0
    list_vals = [val]
    for iter_num in range(n_iter):
        # Compute gradient and Nesterov momentum
        grad = -alpha*obj_fun_deriv(val)
        mom = mu*mom + grad
        # Update value
        val += mu*mom + grad
        list_vals.append(val)
    return val, list_vals
```

Comparing both momentum GD

Gradient Descent with Added **Momentum**

Nesterov accelerated gradient descent (NAG)

Adding a decay to hyperparameters

Definition (learning rate decay):

Learning rate decay is a technique used to **gradually reduce the learning rate** of gradient descent over iterations.

The idea is to **start with a relatively high learning rate, which can help the model to converge quickly in the early stages of training**, and then **gradually reduce the learning rate as training progresses**.

This can help the model to continue learning and improving even as it gets closer to a good solution, and can also help to prevent overfitting by preventing the model from making overly large updates to its parameters.

There are **several ways to implement learning rate decay**, for instance by reducing the learning rate by a fixed amount at regular intervals, updating the learning rate as $\alpha \leftarrow 0.99\alpha$, every N iterations.

Adding a decay to hyperparameters

A possible implementation of gradient descent with learning rate decay is shown below.

Notice how, on every iteration, the learning rate is updated and decayed as $\alpha \leftarrow K\alpha$, with K a decay factor between 0 and 1.

```python
def vanilla_gd_lrdecay(start_val, alpha = 0.1, alpha_decay_fac = 0.99, n_iter = 100):
    val = start_val
    list_vals = [val]
    for iter_num in range(n_iter):
        # Decay on LR
        alpha *= alpha_decay_fac
        # Update value
        val += -alpha*obj_fun_deriv(val)
        list_vals.append(val)
    return val, list_vals
```

# Adding a decay to hyperparameters

As before, the learning rate decay parameter $K$ is a hyperparameter that requires to be adjusted carefully:

- a high decay might help not overshoot a minimum,

- but decaying too fast might make $\alpha$ go to 0 too soon, leading to a wrong convergence point.

# Adding a decay to hyperparameters

Different decay methods can be implemented, with different effects on the gradient descent algorithm: worth exploring.

# Gradient-based learning rate control

**Definition (gradient-based learning rate control):**

Other options for **controlling the learning rate** suggest to **use the value of the gradients to adjust the learning rate** accordingly.

Gradients are often used in learning rate formulas because they provide a measure of how much the parameters are changing with each update.

The magnitude of the gradients give us an indication of how much "learning" is done at each step, and can be used to adjust the learning rate accordingly.

In notebook 5, we implement a possible learning rate control mechanism, using the gradients and the learning rate factor $\alpha_{cf}$, below.

$$\alpha = \frac{1}{(1 + \alpha_{cf} \sqrt{|f'(x)|})}$$

# Gradient-based learning rate control

Gradient-based learning control scheme:

$$\alpha = \frac{1}{(1 + \alpha_{cf} \sqrt{|f'(x)|})}$$

```python
def vanilla_gd_gradlr(start_val, alpha = 0.1, alpha_control_fac = 0.99, n_iter = 100):
    val = start_val
    list_vals = [val]
    for iter_num in range(n_iter):
        # Compute gradients
        grad = obj_fun_deriv(val)
        # Gradient-based decay on LR
        alpha *= 1/(1 + alpha_control_fac*np.sqrt(np.abs(grad)))
        # Update value
        val += -alpha*grad
        list_vals.append(val)
    return val, list_vals
```

Gradient-based learning rate control

If the gradients are large, it may indicate that we are is still making significant progress and are able to adjust parameters effectively at a high learning rate. In this case, we may want to keep the learning rate high to allow the algorithm to continue adjusting the parameters quickly.

On the other hand, **if the gradients are small**, it may indicate that we are close to converging and should only make small updates to the parameters. In this case, we may want to reduce the learning rate to continue improving, but at a slower pace to prevent overshooting the minimum.

By incorporating gradients into the learning rate adjustment formula, we dynamically adjust the learning rate based on the current learning progress. This helps improving effectively as we get closer to a good solution.

Gradient-based learning rate control

Do note however, that this can be a **risky strategy**.

And it is often preferable for learning rate control schemes to do the opposite, that is making the learning rate inversely proportional to the mean gradient value (or its mean squared value). By doing so,

- **if the gradients are small** (i.e. the model is making small updates to its weights), the learning rate will be high to prevent vanishing gradients.

- and **if the gradients are large** (i.e. the model is making large updates to its weights), the learning rate will be low to prevent exploding gradients.

# Gradient-based learning rate control

Gradient-based learning rate control

This can also help to prevent the model from making overly large updates, which can – in turn – help to improve generalization by preventing overfitting.

Most advanced gradient descent algorithms will typically combine

- **a gradient-based learning rate control,**

- **a learning rate decay,**

- **and some momentum formula**.

Going back to our classification dataset

So far, training using "vanilla" gradient descent

```
43    def backward(self, inputs, outputs, alpha = 1e-5):
44            # Get the number of samples in dataset
45            m = inputs.shape[0]
46
47            # Forward propagate
48            Z1 = np.matmul(inputs, self.W1)
49            Z1_b = Z1 + self.b1
50            A1 = self.sigmoid(Z1_b)
51            Z2 = np.matmul(A1, self.W2)
52            Z2_b = Z2 + self.b2
53            A2 = self.sigmoid(Z2_b)
54
55            # Compute error term
56            dL_dA2 = -outputs/A2 + (1 - outputs)/(1 - A2)
57            dL_dZ2 = dL_dA2*A2*(1 - A2)
58            dL_dA1 = np.dot(dL_dZ2, self.W2.T)
59            dL_dZ1 = dL_dA1*A1*(1 - A1)
60
61            # Gradient descent update rules
62            self.W2 -= (1/m)*alpha*np.dot(A1.T, dL_dZ2)
63            self.W1 -= (1/m)*alpha*np.dot(inputs.T, dL_dZ1)
64            self.b2 -= (1/m)*alpha*np.sum(dL_dZ2, axis = 0, keepdims = True)
65            self.b1 -= (1/m)*alpha*np.sum(dL_dZ1, axis = 0, keepdims = True)
66
```

AdaGrad

Definition (AdaGrad optimizer [AdaGrad2011]):

AdaGrad is variation of gradient descent, which **adapts the learning rate for each parameter separately**, based on the historical gradient information for that parameter.

AdaGrad **scales down the learning rate for parameters that have received a large number of updates**, which can help prevent overfitting and improve generalization.

The gradient descent rule relies on variables G_W and G_b. Their values are defined as cumulated sum of the squared values of the gradients with respect to each parameter.

For instance, G_{W_2} starts at 0, and then, on each iteration:

$$G_{W_2} = G_{W_2} + \left(\frac{\partial L}{\partial W_2} \right)^2$$

This is then used to adjust the learning rate α_{W_2} for the W_2 update rule:

$$\alpha_{W_2} = \frac{\alpha}{\sqrt{G_{W_2} + \epsilon}}$$

Here, the value of ϵ is set to $1e^{-6}$ and is simply used to prevent divisions by zero.

The update rule for W_2 is then:

$$W_2 \leftarrow W_2 - \alpha_{W_2} \frac{\partial L}{\partial W_2}$$

AdaGrad implementation

Keep history of gradient in a list called G_list, which will be passed as parameters and returned on each iteration of backward
(Note that it could have also been an attribute for class)

Compute gradients and G_... Coefficients for each trainable parameter

Update rule uses adjusted learning rate for each parameter and epsilon value.

```python
def backward(self, inputs, outputs, G_list, alpha = 1e-5):
    # Get the number of samples in dataset
    m = inputs.shape[0]

    # Forward propagate
    Z1 = np.matmul(inputs, self.W1)
    Z1_b = Z1 + self.b1
    A1 = self.sigmoid(Z1_b)
    Z2 = np.matmul(A1, self.W2)
    Z2_b = Z2 + self.b2
    A2 = self.sigmoid(Z2_b)

    # Compute error term
    dL_dA2 = -outputs/A2 + (1 - outputs)/(1 - A2)
    dL_dZ2 = dL_dA2*A2*(1 - A2)
    dL_dA1 = np.dot(dL_dZ2, self.W2.T)
    dL_dZ1 = dL_dA1*A1*(1 - A1)

    # Gradient descent update rules
    grad_W2 = (-1/m)*np.dot(A1.T, dL_dZ2)
    grad_W1 = (-1/m)*np.dot(inputs.T, dL_dZ1)
    grad_b2 = (-1/m)*np.sum(dL_dZ2, axis = 0, keepdims = True)
    grad_b1 = (-1/m)*np.sum(dL_dZ1, axis = 0, keepdims = True)

    # Momentum and gradient decay/normalization for each parameter
    G_W2, G_W1, G_b2, G_b1 = G_list
    G_W2 += grad_W2**2
    G_W1 += grad_W1**2
    G_b2 += grad_b2**2
    G_b1 += grad_b1**2
    G_list = [G_W2, G_W1, G_b2, G_b1]

    # Gradient descent update rules
    eps = 1e-6
    self.W2 += alpha*grad_W2/(np.sqrt(G_W2 + eps))
    self.W1 += alpha*grad_W1/(np.sqrt(G_W1 + eps))
    self.b2 += alpha*grad_b2/(np.sqrt(G_b2 + eps))
    self.b1 += alpha*grad_b1/(np.sqrt(G_b1 + eps))

    # Update loss
    self.CE_loss(inputs, outputs)
    return G_list
```

AdaGrad implementation

RMSProp

Definition (RMSProp optimizer [RMSProp2012]):

RMSProp is a variation of AdaGrad, but it uses a **running average of the squared gradients to scale the learning rate** for each parameter, rather than using the sum of the squared gradients as in AdaGrad.

RMSProp can be less sensitive to the learning rate than AdaGrad, and it is often used in conjunction with other techniques such as momentum.

We will typically keep this idea in mind for later when designing the **Adam** optimizer.

The value G_{W_2} is no longer defined as the squared values of the gradients with respect to each parameter, but instead it use a mixed formula reusing the previous values of these parameters, mixed with the current squared values of the gradients.

$$G_{W_2} = (1 - \rho)\left(\frac{\partial L}{\partial W_2}\right)^2 + \rho \, G_{W_2}$$

This reuses the logic of the momentum formula we discussed earlier. Here, ρ is an hyperparameter, whose value is set between 0 and 1.

RMSProp

Definition (RMSProp optimizer [RMSProp2012]):

RMSProp is a variation of AdaGrad, but it uses a **running average of the squared gradients to scale the learning rate** for each parameter, rather than using the sum of the squared gradients as in AdaGrad.

RMSProp can be less sensitive to the learning rate than AdaGrad, and it is often used in conjunction with other techniques such as momentum.

We will typically keep this idea in mind for later when designing the **Adam** optimizer.

As before, this is then used to adjust the learning rate α_{W_2} for the W_2 update rule:

$$\alpha_{W_2} = \frac{\alpha}{\sqrt{G_{W_2} + \epsilon}}$$

Here, the value of ϵ is set to $1e^{-6}$ and is simply used to prevent divisions by zero.

The update rule for W_2 is then:

$$W_2 \leftarrow W_2 - \alpha_{W_2} \frac{\partial L}{\partial W_2}$$

RMSProp implementation

The only change compared to AdaGrad is in this formula and the addition of another hyperparameter ρ

```python
def backward(self, inputs, outputs, G_list, alpha = 1e-5, rho = 0.1):
    # Get the number of samples in dataset
    m = inputs.shape[0]

    # Forward propagate
    Z1 = np.matmul(inputs, self.W1)
    Z1_b = Z1 + self.b1
    A1 = self.sigmoid(Z1_b)
    Z2 = np.matmul(A1, self.W2)
    Z2_b = Z2 + self.b2
    A2 = self.sigmoid(Z2_b)

    # Compute error term
    dL_dA2 = -outputs/A2 + (1 - outputs)/(1 - A2)
    dL_dZ2 = dL_dA2*A2*(1 - A2)
    dL_dA1 = np.dot(dL_dZ2, self.W2.T)
    dL_dZ1 = dL_dA1*A1*(1 - A1)

    # Compute gradients
    grad_W2 = (-1/m)*np.dot(A1.T, dL_dZ2)
    grad_W1 = (-1/m)*np.dot(inputs.T, dL_dZ1)
    grad_b2 = (-1/m)*np.sum(dL_dZ2, axis = 0, keepdims = True)
    grad_b1 = (-1/m)*np.sum(dL_dZ1, axis = 0, keepdims = True)

    # Momentum and gradient decay/normalization for each parameter
    G_W2, G_W1, G_b2, G_b1 = G_list
    G_W2 = rho*G_W2 + (1 - rho)*grad_W2**2
    G_W1 = rho*G_W1 + (1 - rho)*grad_W1**2
    G_b2 = rho*G_b2 + (1 - rho)*grad_b2**2
    G_b1 = rho*G_b1 + (1 - rho)*grad_b1**2
    G_list = [G_W2, G_W1, G_b2, G_b1]

    # Gradient descent update rules
    eps = 1e-6
    self.W2 += alpha*grad_W2/(np.sqrt(G_W2 + eps))
    self.W1 += alpha*grad_W1/(np.sqrt(G_W1 + eps))
    self.b2 += alpha*grad_b2/(np.sqrt(G_b2 + eps))
    self.b1 += alpha*grad_b1/(np.sqrt(G_b1 + eps))

    # Update loss
    self.CE_loss(inputs, outputs)
    return G_list
```

RMSProp implementation

Adam

Definition (Adam optimizer [Adam2015]):

Adam is a popular variation of gradient descent that combines the ideas of momentum and RMSProp. Adam combines

- an **exponentially decaying averages of the past gradients to scale the learning rate** for each parameter,
- and an **exponentially decaying average of the past squared gradients to scale the learning rate**.

Adam is widely used as it can often achieve good performance with relatively little hyperparameter tuning.

As before, in AdaGrad and RMSProp, **Adam** relies on additional variables.

This time, it will be two of them for each trainable parameter.

For instance, for the weight W_2, we have two variables V_{W_2} and S_{W_2}.

The V and S values are defined following the logic in RMSProp, reusing the intuition of the momentum formula:

$$V_{W_2} = (1 - \beta_1)\frac{\partial L}{\partial W_2} + \beta_1 V_{W_2}$$

$$S_{W_2} = (1 - \beta_2)\left(\frac{\partial L}{\partial W_2}\right)^2 + \beta_2 S_{W_2}$$

Adam

Definition (Adam optimizer [Adam2015]):

Adam is a popular variation of gradient descent that combines the ideas of momentum and RMSProp. Adam combines

- an **exponentially decaying averages of the past gradients to scale the learning rate** for each parameter,
- and an **exponentially decaying average of the past squared gradients to scale the learning rate**.

Adam is widely used as it can often achieve good performance with relatively little hyperparameter tuning.

Here, we have two hyperparameters

- First, β_1 is an hyperparameter, whose value is set between 0 and 1, often set to 0.

- Second, β_2 is an hyperparameter, whose value is set between 0 and 1, often set to 0.999.

The values are then used to calculate the change C_{W_2} below, used for the W_2 update rule:

$$C_{W_2} = \alpha \frac{V_{W_2}}{\sqrt{S_{W_2}} + \epsilon}$$

Here, the value of ϵ is set to $1e^{-6}$ and is simply used to prevent divisions by zero. Do note that ϵ is no longer in the square root. The update rule for W_2 is then:

$$W_2 \leftarrow W_2 - C_{W_2}$$

Adam implementation

Compute all eight S and V variables (we have four trainable parameters).
Use update formula as mentioned.

```python
# Compute gradients
grad_W2 = (-1/m)*np.dot(A1.T, dL_dZ2)
grad_W1 = (-1/m)*np.dot(inputs.T, dL_dZ1)
grad_b2 = (-1/m)*np.sum(dL_dZ2, axis = 0, keepdims = True)
grad_b1 = (-1/m)*np.sum(dL_dZ1, axis = 0, keepdims = True)

# Momentum and gradient decay/normalization for each parameter
V_W2, V_W1, V_b2, V_b1, S_W2, S_W1, S_b2, S_b1 = G_list
V_W2 = beta1*V_W2 + (1 - beta1)*grad_W2
V_W1 = beta1*V_W1 + (1 - beta1)*grad_W1
V_b2 = beta1*V_b2 + (1 - beta1)*grad_b2
V_b1 = beta1*V_b1 + (1 - beta1)*grad_b1
V_W2_norm = V_W2/(1 - beta1**iteration_number)
V_W1_norm = V_W1/(1 - beta1**iteration_number)
V_b2_norm = V_b2/(1 - beta1**iteration_number)
V_b1_norm = V_b1/(1 - beta1**iteration_number)
S_W2 = beta2*S_W2 + (1 - beta2)*grad_W2**2
S_W1 = beta2*S_W1 + (1 - beta2)*grad_W1**2
S_b2 = beta2*S_b2 + (1 - beta2)*grad_b2**2
S_b1 = beta2*S_b1 + (1 - beta2)*grad_b1**2
S_W2_norm = S_W2/(1 - beta2**iteration_number)
S_W1_norm = S_W1/(1 - beta2**iteration_number)
S_b2_norm = S_b2/(1 - beta2**iteration_number)
S_b1_norm = S_b1/(1 - beta2**iteration_number)
G_list = [V_W2, V_W1, V_b2, V_b1, S_W2, S_W1, S_b2, S_b1]

# Gradient descent update rules
eps = 1e-6
self.W2 += alpha*V_W2_norm/(np.sqrt(S_W2_norm) + eps)
self.W1 += alpha*V_W1_norm/(np.sqrt(S_W1_norm) + eps)
self.b2 += alpha*V_b2_norm/(np.sqrt(S_b2_norm) + eps)
self.b1 += alpha*V_b1_norm/(np.sqrt(S_b1_norm) + eps)

# Update loss
self.CE_loss(inputs, outputs)
return G_list
```
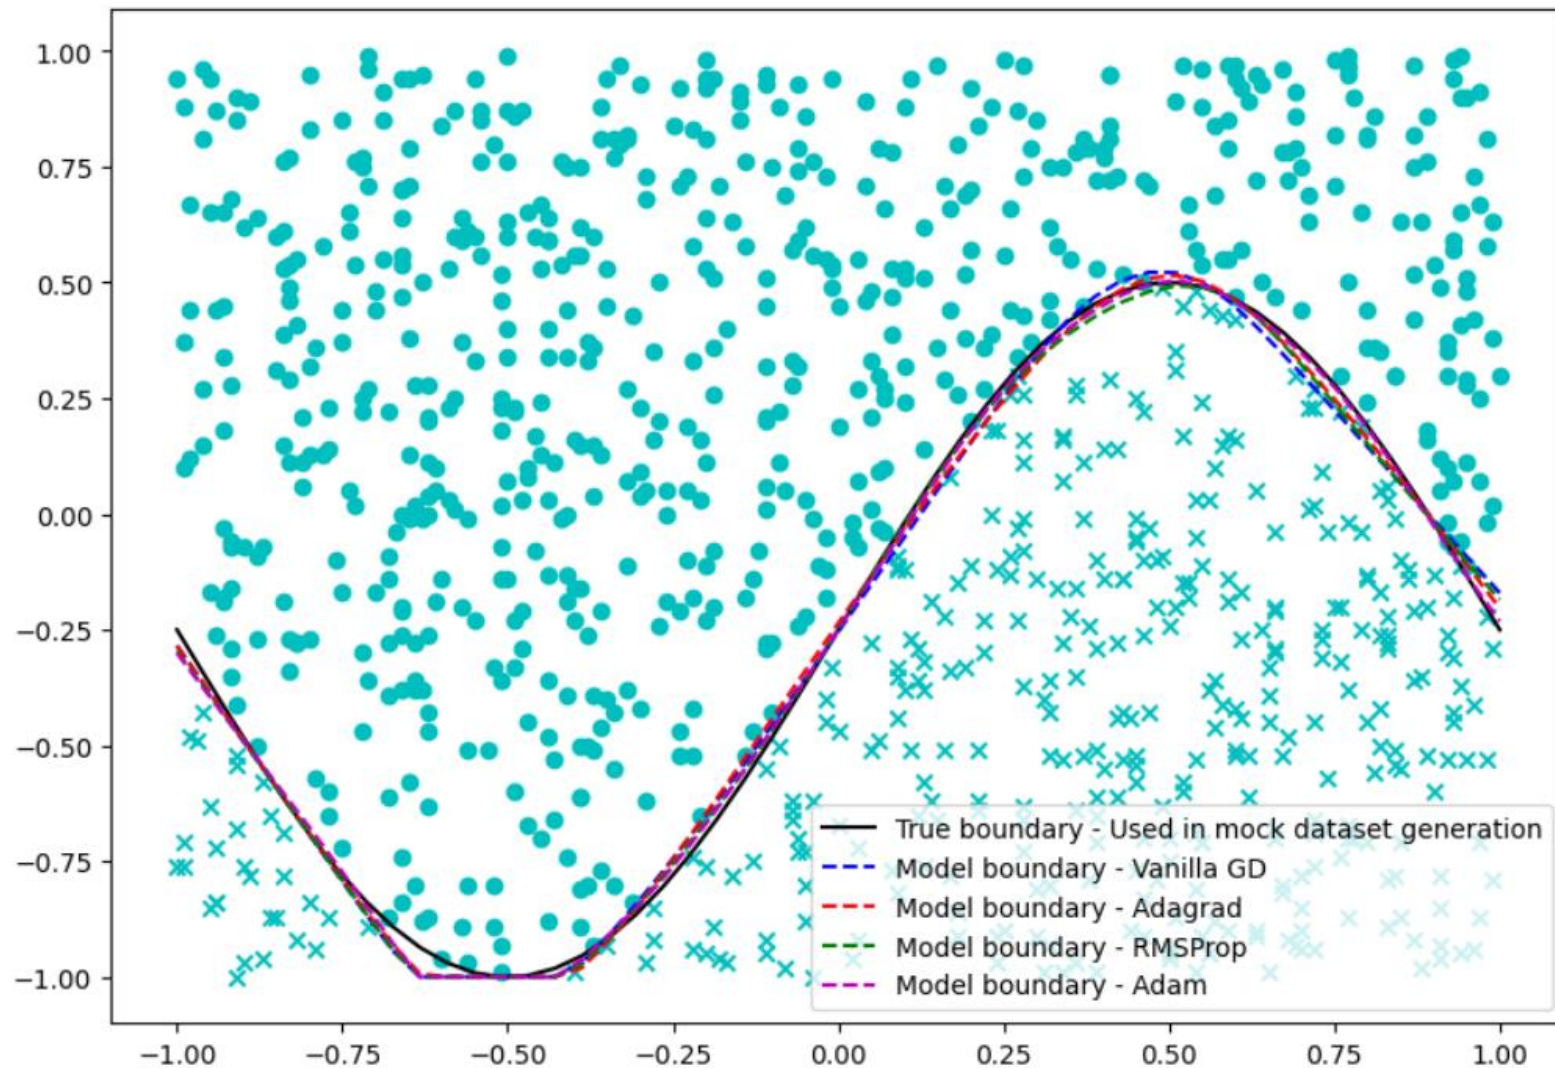
Adam implementation

Comparing all four optimizers

Comparing all four optimizers

PyTorch has many more optimizers

- These days, most Neural Networks seem to rely on Adam.

- There are a few more optimizers, typically implemented in PyTorch.

- Worth looking into: https://pytorch.org/docs/stable/optim.html#torch.optim.Optimizer

Algorithms

Adadelta	Implements Adadelta algorithm.
Adagrad	Implements Adagrad algorithm.
Adam	Implements Adam algorithm.
AdamW	Implements AdamW algorithm.
SparseAdam	Implements lazy version of Adam algorithm suitable for sparse tensors.
Adamax	Implements Adamax algorithm (a variant of Adam based on infinity norm).
ASGD	Implements Averaged Stochastic Gradient Descent.
LBFGS	Implements L-BFGS algorithm, heavily inspired by minFunc.
NAdam	Implements NAdam algorithm.
RAdam	Implements RAdam algorithm.

PyTorch has many more optimizers

There are several recent optimizers that are worth looking into beyond Adam, typically **AdaBound** , **AMSGrad**, and **Lookahead**.

- **AdaBound** is an improved version of Adam that combines the advantages of Adam and the AdaGrad algorithm.

- **AMSGrad**, which stands for Adaptive Moment Estimation with Simulated Annealing and Gradient Descent, is an improved version of Adam that utilizes a special bounding mechanism to prevent overshooting and divergence.

- Finally, **Lookahead** is an algorithm that combines the advantages of several optimization techniques, such as Adam and momentum, to speed up the optimization process.

These are out of scope, but worth mentioning for the curious reader.

Just a question

We have a dataset of N samples, used for training.

Our training procedure, so far consists of:

- Formulating predictions for all N samples, using current model parameters,

- Using all N predictions to compute the value of the loss function L for current model parameters,

- Backpropagation using the loss function we just calculated (parameters update rules calculations and parameters adjustment).

We would therefore perform 1 parameter update every N samples.

Just a question

We have a dataset of N samples, used for training. Our training procedure, so far consists of:

- Formulating predictions for all N samples, using current model parameters,

- Using all N predictions to compute the value of the loss function L for current model parameters,

- Backpropagation using the loss function we just calculated (parameters update rules calculations and parameters adjustment).

We would therefore perform **1 parameter update every N samples**.

Note: this is called a **(Full) Batch Gradient Descent**.

→ While this is a **slow and safe way to go**, could we perform **more regular updates** to reduce the computation for each parameter update?

More regular updates?   1 update every N samples

Stochastic Gradient Descent

Definition (Stochastic Gradient Descent [StoGD1951]):

Vanilla Gradient Descent would perform **1 parameter update every N samples**.

If we wanted to design the most frequent update scheme, it would probably consist of using of

- Formulating predictions **for a single sample, randomly drawn in dataset**, using current model parameters,

- Using this single prediction to compute the value of the loss function L for current model parameters,

- Backpropagation using the loss function we just calculated (parameters update rules calculations and parameters adjustment).

This procedure is called **Stochastic Gradient Descent**.

⟵——————————————⟶

One update per sample?!   1 update every N samples

Stochastic Gradient Descent implementation

Implementing the Stochastic Gradient Descent requires to slightly change the trainer method of our class. **It will now draw a random sample in the dataset and use this single sample to perform the backward operations.**

```python
 98      def train(self, inputs, outputs, N_max = 1000, alpha = 1e-5, beta1 = 0.9, beta2 = 0.999, delta = 1e-5, displa
 99          # Get number of samples
100          M = inputs.shape[0]
101          # List of losses, starts with the current loss
102          self.losses_list = [self.CE_loss(inputs, outputs)]
103          # Initialize G_list
104          G_list = [0*self.W2, 0*self.W1, 0*self.b2, 0*self.b1, \
105                    0*self.W2, 0*self.W1, 0*self.b2, 0*self.b1]
106          # Repeat iterations
107          for iteration_number in range(1, N_max + 1):
108
109              # Stochastic GD on one randomly chosen sample
110              indexes = np.random.randint(0, M)
111              inputs_sub = np.array([inputs[indexes, :]])
112              outputs_sub = np.array([outputs[indexes, :]])
113
114              # Backpropagate
115              G_list, loss = self.backward(inputs_sub, outputs_sub, G_list, iteration_number, alpha, beta1, beta2)
116
```

Stochastic Gradient Descent implementation

In general, slightly faster training, but a lot more erratic.

Stochastic Gradient Descent implementation

In general, slightly faster training, but a lot more erratic.

Reason: most loss function are mean/averaging error values over several samples, e.g. Mean Square Error.

- Using only one sample in the MSE formula will not lead to a good estimation of the Mean Square Error of the model on the dataset.

- Using all samples however leads to the best possible estimation but is slow to compute.

Stochastic GD: worst estimation of the MSE, but fast to compute

What does the optimal middle ground consists of?

Batch GD: best estimation of MSE loss,but slow to compute.

Intuition behind Mini-batch Stochastic GD

What if we used a subset of $N'(with\ N' < N)$ randomly drawn samples from the dataset to perform one iteration of parameter update instead of

- the entire dataset as in Batch GD,

- only one sample as in Stochastic GD?

Definition (samples mini-batch):

We call **samples mini-batch**, or simply **mini-batch**, a subset of $N'(with\ N' < N)$ randomly drawn samples from the dataset.

Stochastic Gradient Descent

Definition (mini-batch stochastic gradient descent [MiniGD1952]):

The **mini-batch stochastic gradient descent** uses the following steps:

- Formulate predictions **for a single mini-batch of N' samples, randomly drawn in dataset**, using current model parameters,

- Using these N' predictions to compute the value of the loss function L for current model parameters,

- Backpropagation using the loss function we just calculated.

It is therefore our **middle ground**.

Stochastic GD: worst estimation of the MSE, but fast to compute

What does the optimal middle ground consists of?
→ Mini-batch stochastic GD?

Batch GD: best estimation of MSE loss,but slow to compute.

Mini-batch Stochastic GD implementation

Our trainer **will now draw N' random sample in the dataset and use these to perform the backward operations.**

```python
 98    def train(self, inputs, outputs, N_max = 1000, alpha = 1e-5, beta1 = 0.9, beta2 = 0.999, \
 99            delta = 1e-5, batch_size = 100, display = True):
100        # Get number of samples
101        M = inputs.shape[0]
102        # List of losses, starts with the current loss
103        self.losses_list = [self.CE_loss(inputs, outputs)]
104        # Initialize G_list
105        G_list = [0*self.W2, 0*self.W1, 0*self.b2, 0*self.b1, \
106                0*self.W2, 0*self.W1, 0*self.b2, 0*self.b1]
107
108        # Define RNG for stochastic minibatches
109        rng = default_rng()
110
111        # Repeat iterations
112        for iteration_number in range(1, N_max + 1):
113
114            # Select a subset of inputs and outputs with given batch size
115            shuffler = rng.choice(M, size = batch_size, replace = False)
116            inputs_sub = inputs[shuffler, :]
117            outputs_sub = outputs[shuffler, :]
118
119            # Backpropagate
120            G_list, loss = self.backward(inputs_sub, outputs_sub, G_list, iteration_number, alpha, beta1, beta2)
```

Mini-batch Stochastic GD implementation

In general, good performance.

Mini-batch Stochastic GD implementation

A quick note on batch sizes: In general,

- We will choose a batch size, defined as a power of 2, between 32 and 512, that is $batch\ size = \{32, 64, 128, 256, 512\}$.

- In general, a larger batch size means slower computation but better performance.

- As in all things so far, **it will be about finding the correct balance!**

Mini-batch Stochastic GD implementation

Mini-batch Stochastic GD implementation

Combining everything into a great optimizer

From now on, we will be combining:

- An **advanced optimizer** (like **Adam**), implementing **momentum** and **gradient-based learning rate control**,

- A **mini-batch stochastic gradient descent** procedure, using mini-batches to perform parameter updates instead of the entire dataset,

Later on, we will implement more controls on the learning rate and Gradient Descent, for instance:

- **Learning rate decay:** decrease the value of learning rate over iterations like before. (Feel free to try implementing it for practice?)

- **Early stopping:** stop iterations of GD, if network has already achieved a good performance (next lecture).

Learn more about these topics

Out of class, supporting papers, for those of you who are curious.

- [AdaGrad2011] J. Duchi, E. **Hazan**, and Y. Singer, "Adaptive Subgradient Methods for Online Learning and Stochastic Optimization", 2011.

- [RMSProp2012] G. **Hinton** "Neural Networks for Machine Learning", 2012.

- [Adam2015] D. **Kingma** and J. **Ba** "Adam: A Method for Stochastic Optimization", 2015

Learn more about these topics

Out of class, supporting papers, for those of you who are curious.

- [StoGD1951] H. Robbins and S. Monro, "A Stochastic Approximation Method", 1951.

- [MiniGD1952] Kiefer and Wolfowitz, "On the Use of Stochastic Approximation Methods in Optimization and Control Problems", 1952.

- Leon Bottou, Yoshua Bengio and Yann LeCun have also worked on implementations of the Mini-batch stochastic Gradient Descent for Neural Networks.
  https://machinelearningmastery.com/gentle-introduction-mini-batch-gradient-descent-configure-batch-size/

Learn more about these topics

Tracking important names (Track their works and follow them on Scholar, Twitter, or whatever works for you!)

- **Elad Hazan: Professor** at **Princeton University**, **Director** and **co-founder** of **Google AI Princeton**.
https://www.ehazan.com
https://scholar.google.com/citations?user=LnhCGNMAAAAJ&hl=fr

- **Diederik P. Kingma: Research Scientist** at **Google Brain.**
http://dpkingma.com/
https://scholar.google.nl/citations?user=yyIoQu4AAAAJ&hl=en

- **Jimmy Ba: Assistant Professor** at **University of Toronto.**
https://jimmylba.github.io/
https://scholar.google.ca/citations?user=ymzxRhAAAAAJ&hl=en

Learn more about these topics

Some extra (easy) reading and videos for those of you who are curious.

- [Analy_Vidhya_Optim] Learn more about optimizers. https://www.analyticsvidhya.com/blog/2021/10/a-comprehensive-guide-on-deep-learning-optimizers/

- [MLM_GD] Learn more about stochastic mini-batch gradient descent. https://machinelearningmastery.com/gentle-introduction-mini-batch-gradient-descent-configure-batch-size/

- Is gradient descent the only way to train parameters anyway?
An introduction to the new 2022 training algorithm, from Hinton, called the forward-forward method.
https://bdtechtalks.com/2022/12/19/forward-forward-algorithm-geoffrey-hinton/