

# 50.039 Theory and Practice of Deep Learning

## W3-S3 Introduction to Deep Learning using the PyTorch framework

Matthieu De Mari



# Introduction (Week 3)

1. What is the **PyTorch library** and its **benefits**?
2. What is a **PyTorch tensor object** and its typical **attributes**?
3. How to implement some typical **tensor operations**?
4. What is **broadcasting** on tensors?
5. What are **tensor locations** in terms of computation?
6. How to **transform our original NumPy shallow Neural Network** class so it uses **PyTorch** now instead?
7. How to implement a **forward**, **loss** and **accuracy** metric in PyTorch?
8. What are some measurable **performance benefits** of using **PyTorch** over NumPy and **GPUs** over CPUs?

# Introduction (Week 3)

9. What is the **autograd/backprop** module in PyTorch, and how does it use a **computational graph** to **compute all derivatives**?
10. How to use the **autograd** to implement **derivatives** and a **vanilla gradient descent**?
11. How to implement **backprop** in PyTorch for our **shallow Neural Network** class?
12. How to use **PyTorch** to implement **advanced optimizers**?
13. How to use **PyTorch** to implement **advanced initializers**?
14. How to use **PyTorch** to implement **regularization**?
15. How to finally revise our **trainer** function to obtain a minimal, yet complete Neural Network in PyTorch?

# Introduction (Week 3)

16. What are the **Dataset** and **Dataloader** objects in **PyTorch**?
17. How to implement a custom **Dataloader** and **Dataset** object in PyTorch?
18. How to move from binary classification to **multi-class classification**?
19. How to adjust output probabilities using the **softmax** function?
20. How to change the **cross-entropy loss** so it works in **multi-class classification**?
21. How to implement **building blocks** in PyTorch?
22. How to implement and train our first **Deep Neural Network**?
23. What are **additional good practices** in PyTorch?

# To summarize (last two sessions)

We now have a full Neural Network class, in PyTorch, with:

- 2 linear layers, sigmoid activation functions,
- Xavier uniform initialization on trainable parameters,
- Forward pass method,
- Autograd backpropagation and trainer method,
- Adam optimizer,
- Dataloader allowing for stochastic mini-batches,
- Cross entropy loss and accuracies,
- L1 regularization.

And it runs/trains at the speed of light (almost...) on GPU!

```

1 class ShallowNeuralNet_PT(torch.nn.Module):
2     def __init__(self, n_x, n_h, n_y, device):
3         super().__init__()
4         self.n_x, self.n_h, self.n_y = n_x, n_h, n_y
5         self.W1 = torch.nn.Parameter(torch.zeros(size = (n_x, n_h), requires_grad = True, \
6                                                     dtype = torch.float64, device = device))
7         torch.nn.init.xavier_uniform_(self.W1.data)
8         self.b1 = torch.nn.Parameter(torch.zeros(size = (1, n_h), requires_grad = True, \
9                                                     dtype = torch.float64, device = device))
10        torch.nn.init.xavier_uniform_(self.b1.data)
11        self.W2 = torch.nn.Parameter(torch.zeros(size = (n_h, n_y), requires_grad = True, \
12                                                    dtype = torch.float64, device = device))
13        torch.nn.init.xavier_uniform_(self.W2.data)
14        self.b2 = torch.nn.Parameter(torch.zeros(size = (1, n_y), requires_grad = True, \
15                                                    dtype = torch.float64, device = device))
16        torch.nn.init.xavier_uniform_(self.b2.data)
17        self.loss = torch.nn.BCELoss()
18        self.accuracy = BinaryAccuracy()
19    def forward(self, inputs):
20        return torch.sigmoid(torch.matmul(torch.sigmoid(torch.matmul(inputs, self.W1) + self.b1), self.W2) + self.b2)
21    def train(self, inputs, outputs, N_max = 1000, alpha = 1, beta1 = 0.9, beta2 = 0.999, \
22             batch_size = 32, lambda_val = 1e-3):
23        dataset = torch.utils.data.TensorDataset(inputs, outputs)
24        data_loader = torch.utils.data.DataLoader(dataset, batch_size = batch_size, shuffle = True)
25        optimizer = torch.optim.Adam(self.parameters(), lr = alpha, betas = (beta1, beta2), eps = 1e-08)
26        optimizer.zero_grad()
27        self.loss_history = []
28        for iteration_number in range(1, N_max + 1):
29            for batch in data_loader:
30                inputs_batch, outputs_batch = batch
31                total_loss = self.loss(self(inputs_batch), outputs_batch.to(torch.float64))\
32                    + lambda_val*sum(torch.abs(param).sum() for param in self.parameters()).item()
33                self.loss_history.append(total_loss)
34                total_loss.backward()
35                optimizer.step()
36                optimizer.zero_grad()
37            if(iteration_number % (N_max//20) == 1):
38                pred = self(inputs)
39                acc_val = self.accuracy(pred, outputs).item()
40                print("Iteration {} - Loss = {} - Accuracy = {}".format(iteration_number, total_loss, acc_val))

```

# Built-in datasets

- Pytorch has a few **built-in datasets**, ready to be downloaded and used on models: typically the most common ones that have been used to demonstrate concepts, such as **MNIST** or **CIFAR-10**.
- For more details on the available **built-in datasets**, have a look at the following page: <https://pytorch.org/vision/stable/datasets.html>.

# Built-in datasets

- Let us demonstrate using the **FashionMNIST** dataset.
- This dataset consists of 28 by 28 greyscale images.
- It is typically used to design image classification models (i.e. models that receive images as inputs), and attempt to predict what is in the image in question.
- The 10 classes (bag, shirt, etc.) are indexed with 0-9 values, corresponding to the 10 types of fashion objects found in the dataset.

```
1 # Dictionary of labels and their identification
2 labels_map = {0: "T-Shirt",
3               1: "Trouser",
4               2: "Pullover",
5               3: "Dress",
6               4: "Coat",
7               5: "Sandal",
8               6: "Shirt",
9               7: "Sneaker",
10              8: "Bag",
11              9: "Ankle Boot"}
```

```
1 # Define folder path as string
2 # Dataset will be downloaded and stored there.
3 folder_path = "./data"
4
5 # Download (download = True) training data (train = True)
6 # to folder specified in root parameter.
7 # The transform parameter specifies that the image samples will
8 # be converted to Tensors, ready to be used by models.
9 training_data = datasets.FashionMNIST(root = folder_path, \
10                                     train = True, \
11                                     download = True, \
12                                     transform = ToTensor())
```



# Built-in datasets

- Let us demonstrate using the **FashionMNIST** dataset.
- This dataset consists of 28 by 28 greyscale images.
- It is typically used to design image classification models (i.e. models that receive images as inputs), and attempt to predict what is in the image in question.
- The 10 classes (bag, shirt, etc.) are indexed with 0-9 values, corresponding to the 10 types of fashion objects found in the dataset.

```
1 # Dataset contains 60000 samples, that are greyscale
2 # images with size 28 by 28 pixels.
3 print(training_data.data.shape)
```

```
torch.Size([60000, 28, 28])
```

```
1 # We can then fetch a sample using the [] notations
2 sample_index = 894
3 img, label = training_data[sample_index]
4 print("Image: ", img.shape)
5 print("Label: ", label)
```

```
Image: torch.Size([1, 28, 28])
```

```
Label: 8
```

```
1 # Display sample
2 sample_index = 894
3 img, label = training_data[sample_index]
4 plt.figure(figsize = (3, 3))
5 plt.axis("off")
6 plt.imshow(img.squeeze(), cmap="gray")
7 plt.title("Image {} - Label = {} ({} {})".format(sample_index, \
8                                                  label, \
9                                                  labels_map[label]))
10 plt.show()
```

Image 894 - Label = 8 (Bag)



# Writing a custom Dataset object

- Most of the time, when demonstrating concepts, we will rely on a "simple" **built-in dataset**, available in the PyTorch library.
- In practice, however, you will often play with a **custom dataset**, fitting your project needs.
- Most datasets will be provided by your future employers, or can be found on dataset search engines, such as **Kaggle**, **Google Dataset Search**, etc.
- Today, we will play with a simplified version of the **Ames Housing Dataset**, which can be found online, here:  
<https://www.kaggle.com/datasets/prevek18/ames-housing-dataset?resource=download>

AutoSaveOn

AmesHousing.xlsx

Last Modified: 9 January

Search

Matthieu DE MARI

FileHomeInsertPage LayoutFormulasDataReviewViewAutomateHelpAcrobatTable Design

CutCopyFormat Painter

Clipboard

Calibri11

Font

Alignment

General

\$%&

0.000.00

Number

NormalBadGoodNeutralCalculationCheck Cell

Conditional FormattingFormat as Table

Styles

InsertDeleteFormat

Cells

AutoSumFillClear

Editing

Sort & FilterFind & Select

Analysis

Analyze DataSensitivity

Sensitivity

CommentsShare

G10

1616

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q
1	Lot Area	Overall Qual	Overall Cond	Year Built	Year Remod/Add	Total Bsmt SF	1st Flr SF	2nd Flr SF	Gr Liv Area	Full Bath	Half Bath	Bedroom AbvGr	Kitchen AbvGr	TotRms AbvGr	Garage Area	Yr Sold	SalePrice
2	31770	6	5	1960	1960	1080	1656	0	1656	1	0	3	1	7	528	2010	215
3	11622	5	6	1961	1961	882	896	0	896	1	0	2	1	5	730	2010	105
4	14267	6	6	1958	1958	1329	1329	0	1329	1	1	3	1	6	312	2010	172
5	11160	7	5	1968	1968	2110	2110	0	2110	2	1	3	1	8	522	2010	244
6	13830	5	5	1997	1998	928	928	701	1629	2	1	3	1	6	482	2010	189
7	9978	6	6	1998	1998	926	926	678	1604	2	1	3	1	7	470	2010	195
8	4920	8	5	2001	2001	1338	1338	0	1338	2	0	2	1	6	582	2010	213
9	5005	8	5	1992	1992	1280	1280	0	1280	2	0	2	1	5	506	2010	193
10	5389	8	5	1995	1996	1595	1616	0	1616	2	0	2	1	5	608	2010	230
11	7500	7	5	1999	1999	994	1028	776	1804	2	1	3	1	7	442	2010	189
12	10000	6	5	1993	1994	763	763	892	1655	2	1	3	1	7	440	2010	175
13	7980	6	7	1992	2007	1168	1187	0	1187	2	0	3	1	6	420	2010	183
14	8402	6	5	1998	1998	789	789	676	1465	2	1	3	1	7	393	2010	180
15	10176	7	5	1990	1990	1300	1341	0	1341	1	1	2	1	5	506	2010	173
16	6820	8	5	1985	1985	1488	1502	0	1502	1	1	1	1	4	528	2010	212
17	53504	8	5	2003	2003	1650	1690	1589	3279	3	1	4	1	12	841	2010	538
18	12134	8	7	1988	2005	559	1080	672	1752	2	0	4	1	8	492	2010	164
19	11394	9	2	2010	2010	1856	1856	0	1856	1	1	1	1	8	834	2010	394
20	19138	4	5	1951	1951	864	864	0	864	1	0	2	1	4	400	2010	143
21	13175	6	6	1978	1988	1542	2073	0	2073	2	0	3	1	7	500	2010	210
22	11751	6	6	1977	1977	1844	1844	0	1844	2	0	3	1	7	546	2010	190
23	10625	7	6	1974	1974	1053	1173	0	1173	2	0	3	1	6	528	2010	170
24	7500	7	5	2000	2000	814	814	860	1674	2	1	3	1	7	663	2010	210
25	11241	6	7	1970	1970	1004	1004	0	1004	1	0	2	1	5	480	2010	149
26	12537	5	6	1971	2008	1078	1078	0	1078	1	1	3	1	6	500	2010	149
27	8450	5	6	1968	1968	1056	1056	0	1056	1	0	3	1	6	304	2010	142
28	8400	4	5	1970	1970	882	882	0	882	1	0	2	1	4	525	2010	120
29	10500	4	5	1971	1971	864	864	0	864	1	0	3	1	5	0	2010	115
30	5858	7	5	1999	1999	1405	1337	0	1337	2	0	2	1	5	511	2010	184
31	1680	6	5	1971	1971	483	483	504	987	1	1	2	1	5	264	2010	90
32	1680	5	5	1971	1971	525	525	567	1092	1	1	3	1	6	320	2010	105
33	1680	6	5	1971	1971	525	525	567	1092	1	1	3	1	6	264	2010	88
34	4043	6	5	1977	1977	1069	1069	0	1069	2	0	2	1	4	440	2010	127
35	2280	6	6	1975	1975	855	855	601	1456	2	1	3	1	6	440	2010	149
36	2280	7	6	1975	1975	836	836	0	836	1	0	2	1	4	308	2010	120
37	2280	6	5	1978	1978	855	855	601	1456	2	1	3	1	7	440	2010	140
38	12858	9	5	2009	2010	1590	1627	707	2334	2	1	3	1	10	751	2010	370
39	11478	8	5	2007	2008	1704	1704	0	1704	2	0	3	1	7	772	2010	300
40	10159	9	5	2009	2010	1930	1940	0	1940	2	1	3	1	8	606	2010	395
41	12883	8	5	2009	2010	1544	1544	0	1544	2	0	3	1	7	868	2010	290
42	12182	7	5	2005	2005	1541	1541	0	1541	2	0	3	1	7	532	2010	220
43	11520	9	5	2005	2005	1698	1698	0	1698	2	0	3	1	7	730	2010	275
44	14122	8	5	2005	2006	1822	1822	0	1822	2	0	3	1	8	678	2010	259

AmesHousing

# Writing a custom Dataset object

- The Ames dataset includes a variety of features for approximately 2,800 houses in Ames, Iowa.
- Features include the size of the house (in square feet), the number of bedrooms and bathrooms, and many more. It also includes the sale price for each house.
- The Ames Housing Dataset is a popular choice for machine learning projects, and it has been used to build AI models for predicting house prices, based on various house features.
- It consists of an Excel file (AmesHousing.xlsx) stored in the ./ames/ folder. The original dataset can be found in the AmesHousing.csv file, but we have simplified it by removing some of its features.

# Writing a custom Dataset object

The features we are interested in are:

- **Lot Area:** The area of the lot in square feet.
- **Overall Qual:** A rating of the overall material and finish of the house (1-10).
- **Overall Cond:** A rating of the overall condition of the house (1-10).
- **Year Built:** The year the house was built.
- **Year Remod/Add:** The year the house was remodeled or had an addition added.
- **Total Bsmt SF:** The total surface of the basement, in square feet.
- **1st Flr SF:** The first floor surface, in square feet.
- **2nd Flr SF:** The second floor surface, in square feet.
- **Gr Liv Area:** The above grade (ground) living area, in square feet.
- **Full Bath:** The number of full bathrooms.
- **Half Bath:** The number of half bathrooms.
- **Bedroom AbvGr:** The number of bedrooms.
- **Kitchen AbvGr:** The number of kitchens.
- **TotRms AbvGrd:** The total number of rooms (does not include bathrooms).
- **Garage Area:** The size of the garage, in square feet.
- **Yr Sold:** The year the property was sold.

These **16 features** will be used as **inputs**, and the **output** will consist of just **1 feature**, in the final column of the Excel file, which is:

- **SalePrice:** The sale price, in dollars.

# Writing a custom Dataset object

Let us start by loading the **Excel** file into a pandas **DataFrame** first.

*Note: Not familiar with the pandas library? Find 10 minutes to pick it up, it will definitely serve you in the long run!*

[https://pandas.pydata.org/docs/user\\_guide/10min.html](https://pandas.pydata.org/docs/user_guide/10min.html)

```
1 # Load dataset using pandas, and showing the first five entries
2 ames_dataset = pd.read_excel("./ames/AmesHousing.xlsx")
3 print(ames_dataset.head(5))
```

	Lot Area	Overall Qual	Overall Cond	Year Built	Year Remod/Add	\
0	31770	6	5	1960	1960	
1	11622	5	6	1961	1961	
2	14267	6	6	1958	1958	
3	11160	7	5	1968	1968	
4	13830	5	5	1997	1998	

	Total Bsmt SF	1st Flr SF	2nd Flr SF	Gr Liv Area	Full Bath	Half Bath	\
0	1080	1656	0	1656	1	0	
1	882	896	0	896	1	0	
2	1329	1329	0	1329	1	1	
3	2110	2110	0	2110	2	1	
4	928	928	701	1629	2	1	

	Bedroom AbvGr	Kitchen AbvGr	TotRms AbvGrd	Garage Area	Yr Sold	\
0	3	1	7	528	2010	
1	2	1	5	730	2010	
2	3	1	6	312	2010	
3	3	1	8	522	2010	
4	3	1	6	482	2010	

	SalePrice
0	215000
1	105000
2	172000
3	244000
4	189900

# Writing a custom Dataset object

To write a custom **dataset class** in PyTorch, we need to:

- Create a **class** that **subclasses** **torch.utils.data.Dataset**.
- Define a **\_\_init\_\_** that takes in the required arguments, and stores them as member variables.
- Define a **\_\_getitem\_\_** that takes **index** as input, and **returns** the **data** and **label** at that index as an array. This will allow to use the square bracket notation on our dataset object.
- Define a method **\_\_len\_\_** that returns the **number of samples in the dataset**.

```
class AmesHousingDataset(torch.utils.data.Dataset):

    #The init method will simply initialize attributes, which consist
    # of the details related to the dataset.
    def __init__(self, file_path = "./ames/AmesHousing.xlsx"):
        # Whole data as a pandas array
        self.data = pd.read_excel(file_path)
        self.dataset_length = len(self.data) #2928

        # Extract inputs
        self.input_features_number = 16
        self.input_features = self.data.iloc[:, :16]

        # Extract outputs
        self.output_features_number = 1
        self.output_feature = self.data.iloc[:, 16]

    # The getitem method returns the sample with given index
    # x will consist of the 16 input features for the given sample,
    # whereas y will consist of the 1 output feature for the given sample.
    def __getitem__(self, index):
        # Fetch inputs
        x = self.input_features.iloc[index].values
        # Fetch outputs
        y = self.output_feature.iloc[index]
        return x, y

    # Finally, the len special method should return the number of samples,
    # in the dataset. We could use self.dataset_length, but it is more
    # modular to use len(self.data).
    def __len__(self):
        return len(self.data)
```



# Writing a custom Dataset object

To write a custom **dataset class** in PyTorch, we need to:

- Create a **class** that **subclasses** **torch.utils.data.Dataset**.
- Define a an **\_\_init\_\_** that takes in the required arguments, and stores them as member variables.
- Define a **\_\_getitem\_\_** that takes **index** as input, and **returns** the **data** and **label** at that index as an array. This will allow to use the square bracket notation on our dataset object.
- Define a method **\_\_len\_\_** that returns the **number of samples in the dataset**.

```

1  # Instantiate the dataset
2  ames_dataset = AmesHousingDataset('./ames/AmesHousing.xlsx')
3
4  # Fetch sample with index 286
5  sample_input, sample_output = ames_dataset[286]
6  # Input is a (16,) numpy array, with the following values
7  print(type(sample_input), sample_input.shape)
8  print(sample_input)
9  # Output is a single value, of type numpy int64
10 print(type(sample_output), sample_output.shape)
11 print(sample_output)

```

```

<class 'numpy.ndarray'> (16,)
[6858    6    4 1915 1950  806  841  806 1647    1    1    4    1    6
 216 2010]
<class 'numpy.int64'> ()
128000

```



# Writing a custom Dataloader

- Before we can feed this dataset object to Neural Networks, we need to supplement it with a **Dataloader**.
- The Dataloader will **shuffle the samples randomly** and produce **mini-batches** of a given size.
- This Dataloader typically allows for **stochastic mini-batches**, as discussed in Week 2.
- The Dataloader will also **transform arrays into tensors**.

```
1 # Instantiate the dataset
2 ames_dataset = AmesHousingDataset('./ames/AmesHousing.xlsx')
3
4 # Create a DataLoader for the dataset
5 ames_dataloader = torch.utils.data.DataLoader(ames_dataset, \
6                                               batch_size = 32, \
7                                               shuffle = True)
```

# Writing a custom Dataloader

- We can then use this custom Dataloader to generate mini-batches using a **for** loop.
- Notice how this Dataloader generates **92** ( $= 2928/32$ , rounded up) **batches** of **32 samples**.
- With the exception of the **last batch** (with index 91), that only contains **16 samples** ( $= 2928 \% 32$ ).

Restricted

```
1 for batch_number, batch in enumerate(ames_dataloader):
2     inputs, outputs = batch
3     print("---")
4     print("Batch number: ", batch_number)
5     print(inputs.shape)
6     print(outputs.shape)
```

```
---
Batch number:  0
torch.Size([32, 16])
torch.Size([32])
```

```
---
Batch number:  1
torch.Size([32, 16])
torch.Size([32])
```

```
---
Batch number:  2
torch.Size([32, 16])
torch.Size([32])
```

```
---
Batch number:  3
torch.Size([32, 16])
torch.Size([32])
```

```
---
Batch number:  4
```

```
---
Batch number:  91
torch.Size([16, 16])
torch.Size([16])
```

Restricted

# Writing a custom Dataloader

- We can then use this custom Dataloader to generate mini-batches using a **for** loop.
- **Good practice:** the custom Dataset and custom Dataloader definition should be repeated to generate **training, testing** and **validation sets dataloaders**.
- This (or these) Dataloader(s) will then be fed to our **trainer** function later on.

```

1  for batch_number, batch in enumerate(ames_dataloader):
2      inputs, outputs = batch
3      print("---")
4      print("Batch number: ", batch_number)
5      print(inputs.shape)
6      print(outputs.shape)

```

```

---
Batch number:  0
torch.Size([32, 16])
torch.Size([32])

```

```

---
Batch number:  1
torch.Size([32, 16])
torch.Size([32])

```

```

---
Batch number:  2
torch.Size([32, 16])
torch.Size([32])

```

```

---
Batch number:  3
torch.Size([32, 16])
torch.Size([32])

```

```

---
Batch number:  4

```

```

---
Batch number:  91
torch.Size([16, 16])
torch.Size([16])

```

# Introducing the MNIST dataset

**MNIST** is a widely-used dataset for the benchmarking of machine learning and computer vision algorithms.

It consists of

- a **training set** of 60,000 examples,
- and a **testing set** of 10,000 examples.

All samples consist of 28x28 pixel grayscale images of handwritten digits (0 to 9).

MNIST is often used as a "Hello, World!" example, due to its simplicity and the availability of efficient implementations of various learning algorithms.



# Introducing the MNIST dataset

**MNIST** is a widely-used dataset for the benchmarking of machine learning and computer vision algorithms.

It consists of

- a **training set** of 60,000 examples,
- and a **testing set** of 10,000 examples.

All samples consist of 28x28 pixel grayscale images of handwritten digits (0 to 9).

**Careful however:** the MNIST dataset is often accused of having been **overused** (which is true) and of **being too simple**.



# Introducing the MNIST dataset

- It is a good dataset to use for testing and comparing the performance of different models, as well as for getting familiar with the basics of machine learning and deep learning.
- The images serve as inputs, and the task is therefore to predict which of the ten digits appears in the image.
- This is therefore a **classification task**, like before, except that it consists of **10 different classes** (corresponding to the 0-9 digits) **instead of just two like in binary classification**.



# Writing the MNIST Dataset and Dataloader

```
1 # Define transform to convert images to tensors and normalize them
2 transform_data = Compose([ToTensor(),
3                             Normalize((0.1307,), (0.3081,))])
4
5 # Load the data
6 batch_size = 64
7 train_dataset = MNIST(root='./mnist/', train = True, download = True, transform = transform_data)
8 test_dataset = MNIST(root='./mnist/', train = False, download = True, transform = transform_data)
9 train_loader = torch.utils.data.DataLoader(train_dataset, batch_size = batch_size, shuffle = True)
10 test_loader = torch.utils.data.DataLoader(test_dataset, batch_size = batch_size, shuffle = False)
```

```
1 # Try the dataloader
2 for batch_number, batch in enumerate(train_loader):
3     inputs, outputs = batch
4     print("---")
5     print("Batch number: ", batch_number)
6     print(inputs.shape)
7     print(outputs.shape)
8     break
```

```
---
Batch number: 0
torch.Size([64, 1, 28, 28])
torch.Size([64])
```



# Introducing the MNIST dataset

## Good practice: normalize your data!

In general, do the following:

- Scale the data (pixel values) to the  $[0,1]$  range.
- Normalize to have zero mean and unit standard deviation.

In MNIST, we will then subtract a mean of 0.1307 and divide by a standard deviation of 0.3081.

These values are the original mean and standard deviation of the dataset before normalization!

```
# Define transform to convert images to tensors and normalize them  
transform_data = Compose([ToTensor(),  
                           Normalize((0.1307,), (0.3081,))])
```



# Writing the MNIST Dataset and Dataloader

```
1 # Define transform to convert images to tensors and normalize them
2 transform_data = Compose([ToTensor(),
3                             Normalize((0.1307,), (0.3081,))])
4
5 # Load the data
6 batch_size = 64
7 train_dataset = MNIST(root='./mnist/', train = True, download = True, transform = transform_data)
8 test_dataset = MNIST(root='./mnist/', train = False, download = True, transform = transform_data)
9 train_loader = torch.utils.data.DataLoader(train_dataset, batch_size = batch_size, shuffle = True)
10 test_loader = torch.utils.data.DataLoader(test_dataset, batch_size = batch_size, shuffle = False)
```

```
1 # Try the dataloader
2 for batch_number, batch in enumerate(train_loader):
3     inputs, outputs = batch
4     print("---")
5     print("Batch number: ", batch_number)
6     print(inputs.shape)
7     print(outputs.shape)
8     break
```

```
---
Batch number: 0
torch.Size([64, 1, 28, 28])
torch.Size([64])
```

# From binary to multi-class classification

In **binary classification**, we would produce a **single value  $p$**  as **output**, with  $p$  between 0 and 1.

- This value  $p$  would correspond to the **probability of being of class 1**.
- The **probability of being of class 0** would then simply be  $1 - p$ .
- We would then use a **threshold 0.5** to decide if the sample is predicted of class 0 or 1.

# From binary to multi-class classification

Unfortunately, when we have more than 2 classes, **we can no longer rely on a single output value  $p$ .**

- Instead, it is often preferable to have the model output **10 values:**  
$$(p_0, p_1, p_2, \dots p_9)$$
- Where each  $p_i$  corresponds to the **probability of being of class  $i$ .**
- This could typically be done by asking for **the final layer to produce  $n_y = 10$  values** instead of just a single  $n_y = 1$  value.

# From binary to multi-class classification

Unfortunately, this is not good enough.

- The  $p_i$  are probabilities and **their sum should be equal to 1**, i.e.

$$\sum_i^9 p_i = 1$$

- A **fully connected layer** ( $WX + b$ ) is not smart enough to do that on its own: it might produce values that may not sum up to 1.

To normalize the outputs produced by the final fully connected layer, we will use the **softmax** operation, which is a special activation function.

- It will force the values of the  $p_i$  to fall in the range of  $[0, 1]$ .
- And force their sum to be equal to 1, that is  $\sum_i^9 p_i = 1$ .

# Softmax function

## Definition (the **softmax** function):

The **softmax** function transforms a vector of  $N$  values

$$Y = (y_0, y_1, y_2, \dots, y_K),$$

into another vector of  $N$  values

$$P = (p_0, p_1, p_2, \dots, p_K).$$

The new vector  **$P$**  is **guaranteed to be summing up to 1**, i.e.

$$\sum_{k=0}^K p_k = 1$$

The **softmax** operation

$$p_i = s(y_i, y_{-i})$$

is defined,  $\forall i$ , as:

$$p_i = s(y_i, y_{-i}) = \frac{\exp(y_i)}{\sum_k^K \exp(y_k)}$$

**Note:** the  $y_{-i}$  notation comes from game theory, and consists of every element in vector  $Y$  except  $y_i$ .

$$y_{-i} = (y_1, \dots, y_{i-1}, y_{i+1}, \dots, y_K)$$

# Softmax function


Softmax can be manually implemented as:

```
1 def softmax(x):
2     # Subtract the maximum value from each element of the input vector x
3     # to avoid numerical instability (this is optional, but equivalent)
4     x = x - np.max(x)
5     # Compute the exponent of each element
6     exp_x = np.exp(x)
7     # Normalize the exponentiated values by their sum
8     return exp_x/np.sum(exp_x)
```

```
1 # Ten values that do not sum up to 1
2 Y = np.array([-1, 4, 10, 5, 7, 1, 4, -2, -1, 2])/20
3 print(sum(Y))
4 P = softmax(Y)
5 print(P)
6 print(np.sum(P))
```

```
1.45
[0.08089815 0.10387528 0.14021696 0.10920108 0.12068586 0.08940628
 0.10387528 0.0769527  0.08089815 0.09399024]
0.9999999999999999
```

It will rescale values so that the trend is preserved, but the new vector sums to 1.



# Softmax function

In the case of multi-label classification, we will use the **softmax** operation **as the final activation after the last fully connected layer**.

This will produce a vector of 10 values,

$$P = (p_0, p_1, p_2, \dots, p_9),$$

corresponding to the **probability for sample of being of each class  $i$** .

The **predicted class  $pred$**  will be the index  $i$  of the highest probability value  $p_i$ , i.e.

$$pred = \arg \max_i [p_i].$$

# Softmax function

For instance, this simple neural network consists of **two fully-connected/linear layers**.

A **single ReLU activation** is used between both layers.

**No final activation yet.**

It also consists of a **flattening** operation, which will transform our input images (2D tensors, size 28 by 28), into a “flattened” 1D tensor with size 784 ( $= 28 \times 28$ ).

```
1 class ShallowNeuralNet(torch.nn.Module):
2     def __init__(self, n_x, n_h, n_y):
3         super().__init__()
4         # Define two layers using the nn.Linear()
5         self.fc1 = torch.nn.Linear(n_x, n_h)
6         self.fc2 = torch.nn.Linear(n_h, n_y)
7
8     def forward(self, x):
9         # Flatten images (transform them from 28x28 2D
10        # matrices to 784 1D vectors)
11        x = x.view(x.size(0), -1)
12        # First Wx + b operation
13        out1 = self.fc1(x)
14        # Using ReLU operation as activation after first layer
15        act1 = torch.relu(out1)
16        # Second Wx + b operation and return
17        out2 = self.fc2(act1)
18        return out2
```



# Softmax function

- Consider a neural network model, with **784 input size**, **128 hidden size**, and **10 output size**, eventually transferring the model to a device.
- Next, we get a **single sample** from the **train\_loader** iterator and extract sample info in the variables **data** and **target**.
- This can be simply done with the **next** and **iter** functions.

```

1 # Initialize model
2 model = ShallowNeuralNet(n_x = 784, \
3                           n_h = 128, \
4                           n_y = 10).to(device)
5
6 # Get a single sample
7 sample = next(iter(train_loader))
8 data, target = sample
9 print(data.shape)
10 print(target.shape)
11 data1 = data[0].to(device)
12 target1 = target[0].to(device)
13 print(data1.shape)
14 print(target1)
15
16 # Forward pass
17 out2 = model(data1)
18 act2 = torch.nn.functional.softmax(out2, dim = 1)
19 print(out2)
20 print(act2)
21 print(act2.sum())

```

torch.Size([64, 1, 28, 28])  
 torch.Size([64])  
 torch.Size([1, 28, 28])  
 tensor(4, device='cuda:0')  
 tensor([[[-0.2145, -0.0527, -0.1197, -0.0488, 0.2503, -0.1870, -0.1401, 0.1202,  
 0.1220, 0.1017]], device='cuda:0', grad\_fn=<AddmmBackward0>)  
 tensor([[0.0812, 0.0954, 0.0892, 0.0958, 0.1292, 0.0834, 0.0874, 0.1134, 0.1136,  
 0.1113]], device='cuda:0', grad\_fn=<SoftmaxBackward0>)  
 tensor(1., device='cuda:0', grad\_fn=<SumBackward0>)

# Softmax function

- Perform a forward pass through the model, storing the **output (no softmax)** in the variable **out2**.
- Apply the **softmax** operation on **out2**. PyTorch offers a functional implementation of the softmax:

*torch.nn.functional.softmax()*

- We can then verify that **softmax** will adjust the output of the neural network correctly.

```

1 # Initialize model
2 model = ShallowNeuralNet(n_x = 784, \
3                           n_h = 128, \
4                           n_y = 10).to(device)
5
6 # Get a single sample
7 sample = next(iter(train_loader))
8 data, target = sample
9 print(data.shape)
10 print(target.shape)
11 data1 = data[0].to(device)
12 target1 = target[0].to(device)
13 print(data1.shape)
14 print(target1)
15
16 # Forward pass
17 out2 = model(data1)
18 act2 = torch.nn.functional.softmax(out2, dim = 1)
19 print(out2)
20 print(act2)
21 print(act2.sum())

```

```

torch.Size([64, 1, 28, 28])
torch.Size([64])
torch.Size([1, 28, 28])
tensor(4, device='cuda:0')
tensor([[[-0.2145, -0.0527, -0.1197, -0.0488,  0.2503, -0.1870, -0.1401,  0.1202,
          0.1220,  0.1017]], device='cuda:0', grad_fn=<AddmmBackward0>])
tensor([[0.0812, 0.0954, 0.0892, 0.0958, 0.1292, 0.0834, 0.0874, 0.1134, 0.1136,
          0.1113]], device='cuda:0', grad_fn=<SoftmaxBackward0>])
tensor(1., device='cuda:0', grad_fn=<SumBackward0>)

```

# Forward implementation

- Our model is however the same as before.

We do not use the softmax operation in the forward.

- This is normal as **the softmax operation** will be used in the loss function, **cross\_entropy**, which will be used in the **trainer()** function later.

```
1 class ShallowNeuralNet(torch.nn.Module):
2     def __init__(self, n_x, n_h, n_y):
3         super().__init__()
4         # Define two layers using the nn.Linear()
5         self.fc1 = torch.nn.Linear(n_x, n_h)
6         self.fc2 = torch.nn.Linear(n_h, n_y)
7
8     def forward(self, x):
9         # Flatten images (transform them from 28x28 2D
10        # matrices to 784 1D vectors)
11        x = x.view(x.size(0), -1)
12        # First Wx + b operation
13        out1 = self.fc1(x)
14        # Using ReLU operation as activation after first layer
15        act1 = torch.relu(out1)
16        # Second Wx + b operation and return
17        out2 = self.fc2(act1)
18        return out2
```

# From binary cross entropy...

- Speaking of, in the case of the **binary classification**, we used the following loss function, namely the **log-likelihood function**.

$$L(x, y) = -\frac{1}{N} \sum_k^N y_k \ln(p(x_k)) + (1 - y_k) \ln(1 - p(x_k))$$

- But in the case of MNIST, **we have more than two classes...**

**How does the loss function change now that we have 10 classes?**

# ...To multi-class cross entropy!

The adjustment is actually quite simple, and the **multi-class cross-entropy loss function** simply rewrites as shown below:

$$L(x, y) = -\frac{1}{N} \sum_k^N \sum_{i=0}^9 y_k^i \ln(p_i(x_k)).$$

In the formula above,  $p_i(x_k)$  denotes the **probability for sample  $x_k$  of being of class  $i$** . In other words, it is the  $i$ -th value of the output vector produced by the model for sample  $x_k$ , **after softmax has been applied**.

...To multi-class cross entropy!

The adjustment is actually quite simple, and the **multi-class cross-entropy loss function** simply rewrites as shown below:

$$L(x, y) = -\frac{1}{N} \sum_k^N \sum_{i=0}^9 y_k^i \ln(p_i(x_k)).$$

The value  $y_k^i$  is the ground truth value for the probability of being of class  $i$  for sample  $x_k$ .

For instance, if the sample  $x_k$  is of class  $y_k = 2$ , we have:

$$Y_k = (y_k^0, y_k^1, y_k^2, y_k^3, \dots, y_k^9) = (0, 0, 1, 0, \dots, 0).$$

We say that  $Y_k$  is the **one-hot vector** for the given scalar value  $y_k = 2$ .

# Setting a model in train mode

**New good practice:** some operations (layers, activations, etc.) in forward will have **two different behaviors** depending on whether

- the model is currently **training**,
- or if we are using its trained version for **evaluation**.

*(Note: at the moment, we have not seen such operations.)*

But let us keep this in mind, and accept that is good practice to set the model to either **train()** or **eval()** mode.

```
1 # Initialize the model and optimizer
2 model = ShallowNeuralNet(n_x = 784, n_h = 64, n_y = 10).to(device)
3 optimizer = torch.optim.Adam(model.parameters(), lr = 1e-3)
4 # Set model in train mode!
5 model.train()
```

# Training our model with Adam GD, as before

```
1 # Training model
2 num_epochs = 5
3 for epoch in range(num_epochs):
4     # Go through all samples in train dataset
5     for i, (images, labels) in enumerate(train_loader):
6         # Get from dataloader and send to device
7         images = images.to(device)
8         labels = labels.to(device)
9         # Forward pass
10        outputs = model(images)
11        # Compute loss
12        loss = torch.nn.functional.cross_entropy(outputs, labels)
13        # Backward and optimize
14        optimizer.zero_grad()
15        loss.backward()
16        optimizer.step()
17        # Display
18        if (i+1) % 300 == 0:
19            print (f'Epoch [{epoch+1}/{num_epochs}], Step [{i+1}/{len(train_loader)}], Loss: {loss.item():.4f}')
```

```
Epoch [1/5], Step [300/938], Loss: 0.2895
Epoch [1/5], Step [600/938], Loss: 0.1588
Epoch [1/5], Step [900/938], Loss: 0.1753
Epoch [2/5], Step [300/938], Loss: 0.0461
Epoch [2/5], Step [600/938], Loss: 0.1251
Epoch [2/5], Step [900/938], Loss: 0.1592
Epoch [3/5], Step [300/938], Loss: 0.1241
Epoch [3/5], Step [600/938], Loss: 0.0511
Epoch [3/5], Step [900/938], Loss: 0.0553
Epoch [4/5], Step [300/938], Loss: 0.0514
Epoch [4/5], Step [600/938], Loss: 0.0339
```



# Eval mode and accuracy after training

After training, we will evaluate our trained model to check its accuracy on test set.

- Set the model in **eval()** mode (good practice for later).
- Predict on **test dataloader**.
- Calculate **accuracy** manually (we could have probably also used a torch function to do that).
- **97% = nicely trained model!**

```
1 # Evaluate model accuracy on test after training
2 # Set model in eval mode!
3 model.eval()
4 # Evaluate
5 with torch.no_grad():
6     correct = 0
7     total = 0
8     for images, labels in test_loader:
9         # Get images and labels from test loader
10        images = images.to(device)
11        labels = labels.to(device)
12        # Forward pass and predict class using max
13        outputs = model(images)
14        _, predicted = torch.max(outputs.data, 1)
15        # Check if predicted class matches label
16        # and count number of correct predictions
17        total += labels.size(0)
18        correct += (predicted == labels).sum().item()
19 # Compute final accuracy and display
20 accuracy = correct/total
21 print(f'Evaluation after training, Accuracy: {accuracy:.4f}')
```

Evaluation after training, Accuracy: 0.9712

# It is now time...

It is now time for us to define and train our first **Deep Neural Network**.

**By definition, a deep neural network consists of more than two layers.**

Here we will create a deep neural network with four layers:

- **three linear layers** with **ReLU** activation,
- followed by **one linear layer**, finished with a **softmax activation**.

# It is now time...

In general, **it is good practice to have the size decrease progressively by a factor of at least 2 from one layer to another.**

For instance, we decided here, to have

- the first layer receive inputs of size 784 and produce outputs of size 80,
- the second layer receive inputs of size 80 and produce outputs of size 40,
- the third layer receive inputs of size 40 and produce outputs of size 20,
- and the fourth layer receive inputs of size 20 and produce outputs of size 10.

# Creating blocks

**Good practice (another one):  
create building blocks for  
modularity.**

- The **DenseReLU** class is a custom PyTorch module that consists of a **linear layer** followed by a **ReLU activation** function.
- The **DenseNoReLU** class is similar, but it applies **no activation**.

**Important note:** no softmax as final activation, for the same reason as before.

```
1 class DenseReLU(torch.nn.Module):
2     def __init__(self, n_x, n_y):
3         super().__init__()
4         # Define Linear layer using the nn.Linear()
5         self.fc = torch.nn.Linear(n_x, n_y)
6
7     def forward(self, x):
8         # Wx + b operation
9         # Using ReLU operation as activation after
10        return torch.relu(self.fc(x))
```

```
1 class DenseNoReLU(torch.nn.Module):
2     def __init__(self, n_x, n_y):
3         super().__init__()
4         # Define Linear layer using the nn.Linear()
5         self.fc = torch.nn.Linear(n_x, n_y)
6
7     def forward(self, x):
8         # Wx + b operation
9         # No activation function
10        return self.fc(x)
```

# Creating blocks

The DeepNeuralNet class will here represent the overall deep neural network.

It starts by initializing four layers:

- Three **DenseReLU** layers,
- And one **DenseNoReLU**.

It then combines them into a single PyTorch sequential model using **torch.nn.Sequential()**.

```
class DeepNeuralNet(torch.nn.Module):
    def __init__(self, n_x, n_h, n_y):
        super().__init__()
        # Define three Dense + ReLU layers,
        # followed by one Dense + Softmax layer
        self.layer1 = DenseReLU(n_x, n_h[0])
        self.layer2 = DenseReLU(n_h[0], n_h[1])
        self.layer3 = DenseReLU(n_h[1], n_h[2])
        self.layer4 = DenseNoReLU(n_h[2], n_y)

        # Combine all four layers
        self.combined_layers = torch.nn.Sequential(self.layer1,
                                                    self.layer2,
                                                    self.layer3,
                                                    self.layer4)

    def forward(self, x):
        # Flatten images (transform them from 28x28
        # 2D matrices to 784 1D vectors)
        x = x.view(x.size(0), -1)
        # Pass through all four layers
        out = self.combined_layers(x)
        return out
```

```
# Initialize the model and optimizer
model = DeepNeuralNet(n_x = 784, n_h = [80, 40, 20], n_y = 10).to(device)
optimizer = torch.optim.Adam(model.parameters(), lr = 1e-3)
```

# Creating blocks

The forward pass of the network is then simply defined:

- The input image is first flattened from a 2D image to a 1D vector,
- It is then passed through the combined layers we have assembled in Sequential().

```
class DeepNeuralNet(torch.nn.Module):
    def __init__(self, n_x, n_h, n_y):
        super().__init__()
        # Define three Dense + ReLU layers,
        # followed by one Dense + Softmax layer
        self.layer1 = DenseReLU(n_x, n_h[0])
        self.layer2 = DenseReLU(n_h[0], n_h[1])
        self.layer3 = DenseReLU(n_h[1], n_h[2])
        self.layer4 = DenseNoReLU(n_h[2], n_y)

        # Combine all four layers
        self.combined_layers = torch.nn.Sequential(self.layer1,
                                                    self.layer2,
                                                    self.layer3,
                                                    self.layer4)

    def forward(self, x):
        # Flatten images (transform them from 28x28
        # 2D matrices to 784 1D vectors)
        x = x.view(x.size(0), -1)
        # Pass through all four layers
        out = self.combined_layers(x)
        return out
```

```
# Initialize the model and optimizer
model = DeepNeuralNet(n_x = 784, n_h = [80, 40, 20], n_y = 10).to(device)
optimizer = torch.optim.Adam(model.parameters(), lr = 1e-3)
```

# Creating blocks

**Good practice:** This modular approach,

- **defining blocks of layers,**
- **and eventually assembling them in a larger Deep Neural Network network,**

is very common in deep neural networks, especially when the architectures are very heavy and include many layers.

Organize the mess!

```
class DeepNeuralNet(torch.nn.Module):
    def __init__(self, n_x, n_h, n_y):
        super().__init__()
        # Define three Dense + ReLU layers,
        # followed by one Dense + Softmax layer
        self.layer1 = DenseReLU(n_x, n_h[0])
        self.layer2 = DenseReLU(n_h[0], n_h[1])
        self.layer3 = DenseReLU(n_h[1], n_h[2])
        self.layer4 = DenseNoReLU(n_h[2], n_y)

        # Combine all four layers
        self.combined_layers = torch.nn.Sequential(self.layer1,
                                                    self.layer2,
                                                    self.layer3,
                                                    self.layer4)

    def forward(self, x):
        # Flatten images (transform them from 28x28
        # 2D matrices to 784 1D vectors)
        x = x.view(x.size(0), -1)
        # Pass through all four layers
        out = self.combined_layers(x)
        return out
```

```
# Initialize the model and optimizer
model = DeepNeuralNet(n_x = 784, n_h = [80, 40, 20], n_y = 10).to(device)
optimizer = torch.optim.Adam(model.parameters(), lr = 1e-3)
```

# Training our model as before

```
1 # Training model
2 num_epochs = 10
3 for epoch in range(num_epochs):
4     # Go through all samples in train dataset
5     for i, (images, labels) in enumerate(train_loader):
6         # Get from dataloader and send to device
7         images = images.to(device)
8         labels = labels.to(device)
9         # Forward pass
10        outputs = model(images)
11        # Compute loss
12        loss = torch.nn.functional.cross_entropy(outputs, labels)
13        # Backward and optimize
14        optimizer.zero_grad()
15        loss.backward()
16        optimizer.step()
17        # Display
18        if (i+1) % 25 == 0:
19            print (f'Epoch [{epoch+1}/{num_epochs}], Step [{i+1}/{len(train_loader)}], Loss: {loss.item():.4f}')
```

```
Epoch [1/10], Step [25/235], Loss: 2.2469
Epoch [1/10], Step [50/235], Loss: 2.1274
Epoch [1/10], Step [75/235], Loss: 1.9379
Epoch [1/10], Step [100/235], Loss: 1.8415
Epoch [1/10], Step [125/235], Loss: 1.6548
Epoch [1/10], Step [150/235], Loss: 1.4660
Epoch [1/10], Step [175/235], Loss: 1.2895
Epoch [1/10], Step [200/235], Loss: 1.1478
Epoch [1/10], Step [225/235], Loss: 0.9910
Epoch [2/10], Step [25/235], Loss: 0.8251
```



# On Deep Neural Networks complexity

We can therefore raise a fairly important question:

**What is the appropriate number of layers to use and how many neurons should we use on each layer?**

- We established in Week 2 – Notebook 8, that there is no fixed rule for how many layers should be used in a deep neural network.
- Instead, the number of layers, as well as the number of neurons in each layer, should be chosen based on the complexity of the task and the amount of available data.
- In general, deep neural networks with many layers (hundreds or even thousands) can learn very complex patterns in data, but they will require a large amount of data and computational resources to train.

# On Deep Neural Networks complexity

- More importantly, if the network is **too deep**, it may also be prone to **overfitting**, which can hinder its generalization performance on unseen data.
- On the other hand, **shallow networks** with fewer layers may be easier to train and require less data, but **they may not be able to learn as complex patterns**.
- **NFL**: Finding the optimal number of layers and the optimal architecture of a deep neural network is often a trade-off between model complexity, computational resources, and performance, and requires some experimentation and model selection.

# In fact, our DNN is overfitting at the moment!

**Shallow Neural Net: 96.5% test acc  
(not too bad)**

```

1 # Evaluate model accuracy on test after training
2 # Set model in eval mode!
3 model.eval()
4 # Evaluate
5 with torch.no_grad():
6     correct = 0
7     total = 0
8     for images, labels in test_loader:
9         # Get images and labels from test loader
10         images = images.to(device)
11         labels = labels.to(device)
12         # Forward pass and predict class using max
13         outputs = model(images)
14         _, predicted = torch.max(outputs.data, 1)
15         # Check if predicted class matches label
16         # and count number of correct predictions
17         total += labels.size(0)
18         correct += (predicted == labels).sum().item()
19 # Compute final accuracy and display
20 accuracy = correct/total
21 print(f'Evaluation after training, Accuracy: {accuracy:.4f}')
```

Evaluation after training, Accuracy: 0.9649

**Deep Neural Net: 93.7% test acc  
(lower, even though we had a lower loss!)**

```

1 # Evaluate model accuracy on test after training
2 # Set model in eval mode!
3 model.eval()
4 # Evaluate
5 with torch.no_grad():
6     correct = 0
7     total = 0
8     for images, labels in test_loader:
9         # Get images and labels from test loader
10         images = images.to(device)
11         labels = labels.to(device)
12         # Forward pass and predict class using max
13         outputs = model(images)
14         _, predicted = torch.max(outputs.data, 1)
15         # Check if predicted class matches label
16         # and count number of correct predictions
17         total += labels.size(0)
18         correct += (predicted == labels).sum().item()
19 # Compute final accuracy and display
20 accuracy = correct/total
21 print(f'Evaluation after training, test accuracy: {accuracy:.4f}')
```

Evaluation after training, test accuracy: 0.9367

# Experimenting on layers numbers and sizes

In fact, in Notebook 7, we trained three DNN models:

- **Model 1:** 6 layers (probably too many layers)

$$n_h = [320, 160, 80, 40, 20],$$


- **Model 2:** 3 layers (layers probably too large)

$$n_h = [400, 200],$$

- **Model 3:** 3 layers (just fine?)

$$n_h = [40, 20].$$

While simpler, model 3 has highest test accuracy!



# Introduction (Week 3)

1. What is the **PyTorch library** and its **benefits**?
2. What is a **PyTorch tensor object** and its typical **attributes**?
3. How to implement some typical **tensor operations**?
4. What is **broadcasting** on tensors?
5. What are **tensor locations** in terms of computation?
6. How to **transform our original NumPy shallow Neural Network** class so it uses **PyTorch** now instead?
7. How to implement a **forward**, **loss** and **accuracy** metric in PyTorch?
8. What are some measurable **performance benefits** of using **PyTorch** over NumPy and **GPUs** over CPUs?

# Introduction (Week 3)

9. What is the **autograd/backprop** module in PyTorch, and how does it use a **computational graph** to **compute all derivatives**?
10. How to use the **autograd** to implement **derivatives** and a **vanilla gradient descent**?
11. How to implement **backprop** in PyTorch for our **shallow Neural Network** class?
12. How to use **PyTorch** to implement **advanced optimizers**?
13. How to use **PyTorch** to implement **advanced initializers**?
14. How to use **PyTorch** to implement **regularization**?
15. How to finally revise our **trainer** function to obtain a minimal, yet complete Neural Network in PyTorch?

# Introduction (Week 3)

16. What are the **Dataset** and **Dataloader** objects in **PyTorch**?
17. How to implement a custom **Dataloader** and **Dataset** object in PyTorch?
18. How to move from binary classification to **multi-class classification**?
19. How to adjust output probabilities using the **softmax** function?
20. How to change the **cross-entropy loss** so it works in **multi-class classification**?
21. How to implement **building blocks** in PyTorch?
22. How to implement and train our first **Deep Neural Network**?
23. What are **additional good practices** in PyTorch?

# Conclusion (Week 3)

- PyTorch library and its benefits
- Tensor objects, attributes and operations on tensors
- Converting our NumPy shallow neural network into PyTorch
- Parameter objects
- Forward method implementation
- Performance benefits of GPU acceleration
- Using autograd and computational graphs
- Advanced optimizers in PyTorch
- Initializers in PyTorch
- Regularization in PyTorch
- Dataset and Dataloader objects
- Multi-class classification
- Softmax function and multi-class cross entropy loss
- Building blocks in PyTorch
- Our first Deep Neural Network!
- Network size vs. overfitting tradeoff



# Project announcement!

Let us discuss it now.